

# Algebra Of Disagreement - Nonabelian Cohomology Consensus Failures

2025-07-12

## Table of Contents

- Part 1: An Algebraic-Topological Framework for Consensus Failures
  - Chapter 1.1: Modeling Distributed Executions as Topological Spaces
  - Chapter 1.2: The Torsor of Disagreement: First Cohomology as a Consensus Obstruction
  - Chapter 1.3: Non-Abelian Holonomy and Path-Dependent System States
  - Chapter 1.4: Cocycles as Measures of Local-to-Global Inconsistency
  - Chapter 1.5: Gerbes and Second Cohomology: Obstructions to Protocol Composition
- Part 2: The Directed Topology of Asynchronous Computations
  - Chapter 2.1: Beyond Reversibility: Directed Spaces for Asynchronous Systems
  - Chapter 2.2: The Fundamental Category ( $\text{Cat}_1(X)$ ) and Causal Paths
  - Chapter 2.3: Geometric Semantics: Concurrency, Deadlock, and Directed Holes
  - Chapter 2.4: Directed Homology: Quantifying Causal Anomalies
  - Chapter 2.5: The Topology of Impossibility: Analyzing k-Set Consensus Protocols
- Part 3: Nonabelian State Symmetries and Path-Dependent Updates
  - Chapter 3.1: The Algebra of State Updates: Modeling Operations with Non-Abelian Groups
  - Chapter 3.2: State Symmetries and the Action of the Update Groupoid
  - Chapter 3.3: Path-Dependence as Non-Trivial Holonomy in the State Space
  - Chapter 3.4: Crossed Homomorphisms: The Algebraic Signature of Inconsistent Views
  - Chapter 3.5: Case Study: Consensus Failure in Systems with Non-Commutative Operations

- Part 4: Cohomological Obstructions to Distributed Consensus: Torsors and Gerbes
  - Chapter 4.1: The Torsor of Disagreement:  $H^1(X, A)$  as the Formal Measure of Consensus Impossibility
  - Chapter 4.2: Interpreting Non-Abelian Cocycles as Local State Mismatches
  - Chapter 4.3: The Structure of Disagreement: How the Non-Abelian Group  $A$  Encodes Failure Modes
  - Chapter 4.4: Beyond State Disagreement: Gerbes and  $H^2$  as Obstructions to Protocol Coherence
  - Chapter 4.5: A Cohomological Classification of Faults: From Benign Asynchrony to Byzantine Torsors
- Part 5: Applications in Protocol Analysis and Fault-Tolerance
  - Chapter 5.1: Cohomological Analysis of Paxos: Proving Consensus through a Trivial Disagreement Torsor
  - Chapter 5.2: Raft’s Leader Election as Symmetry Breaking in the Protocol Groupoid
  - Chapter 5.3: Modeling Byzantine Faults: The Non-Abelian Torsor of Malicious Views
  - Chapter 5.4: CRDTs as a Strategy for Trivializing Cohomology: Enforcing Commutativity to Guarantee Convergence
  - Chapter 5.5: Blockchain Forks as Manifestations of Non-Trivial Holonomy in the State Space

## Part 1: An Algebraic-Topological Framework for Consensus Failures

### Chapter 1.1: Modeling Distributed Executions as Topological Spaces

#### Modeling Distributed Executions as Topological Spaces\*\*

The fundamental challenge in distributed computing lies in reasoning about the collective behavior of asynchronous, concurrent, and potentially faulty processes. A system’s evolution is not a single, linear narrative but a complex, branching tapestry of all possible interleavings of events. A simple log of events or a state-machine description for individual processes fails to capture the emergent, global properties and, most critically, the inherent impossibilities that arise from this combinatorial explosion. To tame this complexity, we require a more powerful abstraction: a geometric representation of the entire computational landscape. This chapter introduces the foundational step of our framework: the modeling of distributed executions as topological spaces. By translating the discrete, event-driven nature of computation into a geometric setting, we can leverage the powerful tools of algebraic topology to analyze and classify consensus failures. The “shape” of the resulting space—its connectivity, holes, and twists—encodes deep truths about what a distributed system can and cannot achieve.

## The Simplicial Model: Capturing Global Consistency

The pioneering work in applying topology to distributed computing, most notably by Herlihy, Shavit, and others, introduced the concept of the *protocol complex*. This model uses a simplicial complex—a collection of vertices, edges, triangles, and their higher-dimensional counterparts—to represent the global states of a protocol.

### Construction of the Protocol Complex:

Let a distributed system consist of a set of  $n+1$  processes,  $P = \{p_0, p_1, \dots, p_n\}$ . At any given time, each process  $p_i$  is in some local state  $s_i$ , which includes its internal variables and the value it might propose or decide upon.

- **Vertices:** A vertex in the protocol complex  $P(\Pi)$  represents a possible local state for a single process. It is a pair  $v = (p_i, s_i)$ , indicating “process  $p_i$  is in state  $s_i$ .”
- **Simplices:** A collection of  $k+1$  vertices,  $\sigma = \{v_0, v_1, \dots, v_k\}$  where  $v_j = (p_{i_j}, s_{i_j})$ , forms a  $k$ -*simplex* if and only if the processes  $p_{i_0}, \dots, p_{i_k}$  can *simultaneously* hold the respective states  $s_{i_0}, \dots, s_{i_k}$  in some valid, reachable execution of the protocol  $\Pi$ . A “valid execution” is a sequence of computation or communication steps allowed by the protocol’s rules, originating from a valid initial state.

The entire **protocol complex**  $P(\Pi)$  is the abstract simplicial complex formed by all such valid simplices. This complex represents a snapshot of all possible “consistent views” held by subsets of processes.

### Topological Properties and Computational Tasks:

The utility of this model becomes apparent when we relate the topological properties of  $P(\Pi)$  to the solvability of a consensus problem. For instance, in binary consensus, processes start with an input value (0 or 1) and must decide on a single common output value. The protocol complex for such a system has an *input complex*  $I$ , representing the initial configurations, and an *output complex*  $O$ , representing the terminal configurations. A protocol solves consensus if there is a continuous, decision-map from the “states-after-one-step” complex to the output complex.

The celebrated Asynchronous Computability Theorem states that a task is solvable if and only if the protocol complex representing the task’s allowable final states is not only connected but also simply connected (i.e., has no “holes”). The FLP impossibility result, which proves that deterministic consensus is impossible in an asynchronous system with even one crash fault, can be elegantly re-cast in this framework. The set of possible execution paths for a fault-tolerant protocol necessarily creates a “hole” in the state space, corresponding to an undecidable configuration. Any attempt to map this “punctured” space to a discrete set of

decision values (e.g.,  $\{0, 1\}$ ) must fail, just as one cannot continuously map a circle onto its two endpoints without tearing it.

The protocol complex provides a powerful static picture of global consistency. However, it is fundamentally an *undirected* structure. It tells us *which* states can co-exist, but it erases the crucial information about causality, temporal ordering, and the non-commutative nature of operations—the very factors that govern path-dependent failures.

### Directed Topology: Incorporating Causality and Path Dependence

Asynchronous executions are inherently directed. Information flows from one process to another; time moves forward. An event  $e_1$  that causally precedes  $e_2$  is fundamentally different from a scenario where  $e_2$  precedes  $e_1$ . The order of operations—such as updates to a shared resource—often matters, a property that is central to our investigation of systems with nonabelian state transformations. To capture this, we must refine our model from a simple topological space to a *directed topological space*.

A directed space, or *d-space*, is a topological space  $X$  equipped with a privileged class of paths, called *d-paths* or *execution paths*. These paths are models for the causal evolution of the system. Formally, we might model executions as continuous maps  $\gamma : [0, 1] \rightarrow X$  where the direction of the path from  $\gamma(0)$  to  $\gamma(1)$  represents the forward flow of time and information.

### From Simplicial Complexes to Directed Spaces:

We can imbue our simplicial models with direction in several ways:

1. **State-Transition Graphs:** A natural representation is a graph where nodes are global states and directed edges represent single process steps. The geometric realization of this graph gives a 1-dimensional directed space. Higher-dimensional structures can be built by “filling in” commutative squares and cubes, where the order of two independent operations does not matter.
2. **Ordered Simplices:** We can define an ordering on the vertices of each simplex in the protocol complex. An *ordered simplex*  $(v_0, v_1, \dots, v_k)$  can represent a totally ordered sequence of events. The space of all such causally-consistent ordered simplices naturally forms a directed space.
3. **The Space of Executions:** A more abstract and powerful model, proposed by Fajstrup, Goubault, and Raussen, considers the space of all possible execution traces directly. Points in this space are themselves executions. This approach elegantly captures the branching nature of computation.

### The Fundamental Category $\text{Cat}_1(X)$ :

In a standard topological space, the primary algebraic invariant capturing path-connectivity is the fundamental group  $\pi_1(X, x_0)$ . It consists of homotopy classes of loops based at a point  $x_0$ , with the group operation being path concatenation. This structure inherently assumes that paths are reversible, as the inverse of a loop is simply the loop traversed backward.

This assumption breaks down in directed spaces. An execution path cannot, in general, be run in reverse. The appropriate algebraic structure is not the fundamental group but the **fundamental category**  $\text{Cat}_1(X)$  (or fundamental groupoid  $\Pi_1(X)$  if all paths between two points are considered equivalent).

- **Objects:** The objects of  $\text{Cat}_1(X)$  are the points of the space  $X$  (the global states).
- **Morphisms:** A morphism from state  $S_1$  to state  $S_2$  is a homotopy class of directed paths from  $S_1$  to  $S_2$ .

Composition of morphisms corresponds to the concatenation of execution traces. The non-commutativity of  $\text{Cat}_1(X)$ —the fact that following path  $\gamma_1$  then  $\gamma_2$  may lead to a different state than  $\gamma_2$  then  $\gamma_1$ —is a direct topological reflection of the non-commutative nature of the operations within the distributed system. It is this structure that our nonabelian cohomological framework is designed to probe. For instance, if process  $p_1$  applies an update  $g_1$  and  $p_2$  applies  $g_2$ , the final state may depend on the order, i.e.,  $g_2 \circ g_1 \circ s_0 \neq g_1 \circ g_2 \circ s_0$ . This non-commutativity is encoded as two distinct directed paths between the same start and end points in  $\text{Cat}_1(X)$ .

### Orbifolds: Modeling Symmetries and Indistinguishability

The final layer of geometric refinement addresses symmetries and faults. In many systems, processes may be anonymous or symmetric. More critically, Byzantine faults introduce a malicious form of symmetry: a faulty process can equivocate, sending conflicting information to different observers, thereby creating states that are indistinguishable from valid (but different) states for the correct processes.

When a group  $G$  of symmetries acts on our state space  $X$ , the true space of observations is the quotient space  $X/G$ , where all states in the same orbit under  $G$  are identified. However, this quotienting process can create singularities. For example, if a state  $x$  is a fixed point for some symmetry  $g \in G$  (i.e.,  $g \cdot x = x$ ), the corresponding point  $[x]$  in the quotient  $X/G$  will have a non-trivial isotropy group. Such points behave differently from regular points; they are “cone points” or other singularities.

An **orbifold** is the precise mathematical structure for such “well-behaved” quotient spaces. Locally, an orbifold looks like a patch of Euclidean space  $\mathbb{R}^n$  quotiented by the action of a finite group.

### Applications in Distributed Computing:

1. **Byzantine Faults:** The actions of a Byzantine adversary can be modeled as a group  $G$  acting on the space of executions. This group permutes messages, corrupts data, and generates false reports. The set of states that correct processes can observe is not the full space  $X$  of ideal executions, but the orbifold  $X/G$ . The singularities of this orbifold represent points of extreme ambiguity—configurations where the actions of the adversary create a view that is consistent with multiple, contradictory realities. Reaching consensus requires navigating a path on this orbifold, and the singularities act as topological obstructions.
2. **k-Set Consensus:** In  $k$ -set consensus, processes must decide on at most  $k$  distinct values from their initial inputs. The set of valid terminal states has a natural symmetry under the permutation group  $S_k$ . The task is to find a protocol that maps the initial, highly symmetric state space to one of these valid terminal configurations. The problem’s difficulty is directly related to the topological complexity of the quotient space of configurations modulo these symmetries.

## Conclusion: A Geometric Stage for Algebraic Probes

This chapter has laid the geometric groundwork for our analysis. We have moved from a naive event-log view to a sophisticated topological one, capable of representing the global structure of distributed computation.

- The **simplicial complex** provides the basic framework, capturing states of global consistency and revealing fundamental impossibility results through its static topology (e.g., holes corresponding to  $H_1(X)$ ).
- The **directed space** enriches this model with causality, enabling us to study path-dependent phenomena. Its fundamental invariant, the category  $\text{Cat}_1(X)$ , directly encodes the non-commutative nature of concurrent operations.
- The **orbifold** provides the language to handle systems with intrinsic symmetries or adversarial behavior, where the observable state space is a quotient with singularities that act as computational chokepoints.

These spaces— $X$ —are the domains upon which our algebraic machinery will operate. We have constructed the stage. In the subsequent chapters, we will introduce the actors: nonabelian groups  $A$  representing state transformations and the tool of nonabelian cohomology  $H^n(X, A)$  to probe the structure of this stage. The topological features we have described here—holes, non-reversible paths, and singularities—will manifest as non-trivial cohomological invariants, providing a precise, algebraic classification of the spectrum of consensus failures.

## Chapter 1.2: The Torsor of Disagreement: First Cohomology as a Consensus Obstruction

The Torsor of Disagreement: First Cohomology as a Consensus Obstruction

Having established a topological model for distributed executions—the protocol complex  $\mathbf{X}$ —we now confront the central question of consensus. The problem of consensus, at its core, is the problem of constructing a global section of a sheaf of states over  $\mathbf{X}$ . That is, can we find a single, consistent state  $\mathbf{s}$  that is agreed upon by all processes across all possible executions? The famous FLP impossibility result and the challenges of Byzantine fault tolerance demonstrate that local communication and computation are often insufficient to guarantee such a global section. The failure is not one of local logic but of global topology.

This chapter formalizes this failure by introducing the first nonabelian cohomology set,  $H^1(\mathbf{X}, \mathbf{A})$ , as the primary tool for detecting and classifying obstructions to consensus. We will demonstrate that when consensus is impossible, the system is not merely in a state of disagreement, but is confined to a specific mathematical structure known as a **torsor**. This “Torsor of Disagreement” is a direct manifestation of a nontrivial cohomology class, providing a powerful, algebraic invariant for consensus failures.

**The State Space as a Nonabelian Group  $\mathbf{A}$**  To analyze disagreements algebraically, we must first model the *differences* between states. Let the set of possible data values or states be  $\mathbf{S}$ . We are interested in the transformations between these states. We therefore define a **group of disagreements**, denoted  $\mathbf{A}$ , which acts on  $\mathbf{S}$ . For any two states  $\mathbf{s}_i, \mathbf{s}_j \in \mathbf{S}$ , there exists some  $\mathbf{a} \in \mathbf{A}$  such that  $\mathbf{s}_i = \mathbf{a} \cdot \mathbf{s}_j$ . This group  $\mathbf{A}$  represents the set of all possible relative discrepancies between the views of any two processes.

Crucially,  $\mathbf{A}$  is, in general, a **nonabelian group**. This noncommutativity is not a mathematical contrivance; it is the natural language for describing path-dependent updates in distributed systems. Consider a system where updates correspond to geometric rotations (e.g., attitude control for a swarm of drones) or permutations of a list. The final state depends explicitly on the order in which updates are applied. If  $\mathbf{a}, \mathbf{b} \in \mathbf{A}$  are two updates,  $\mathbf{a} \cdot \mathbf{b} \neq \mathbf{b} \cdot \mathbf{a}$ . This property is the algebraic root of path-dependence in the protocol complex  $\mathbf{X}$ .

The action of the system’s dynamics, represented by the fundamental groupoid  $\mathbf{Cat}_1(\mathbf{X})$  or the fundamental group  $\mathbf{G} = \pi_1(\mathbf{X})$ , can also induce transformations on the disagreement group  $\mathbf{A}$  itself. An execution path  $\gamma \in \pi_1(\mathbf{X})$  might, for instance, represent a network partition that is later healed, altering the trust assumptions and thus the nature of disagreements. This is captured by a group action  $\mathbf{G} \times \mathbf{A} \rightarrow \mathbf{A}$ , written  $(g, \mathbf{a}) \mapsto g \cdot \mathbf{a}$ . For many standard protocols, this action is trivial ( $g \cdot \mathbf{a} = \mathbf{a}$ ), but our framework accommodates the general case.

**1-Cocycles: The Law of Local Consistency** Imagine our distributed system is running. At any point, we can query pairs of processes  $(p_i, p_j)$  that have recently communicated. This communication forms an edge  $\mathbf{e}_{\{ij\}}$  in the protocol complex  $\mathbf{X}$ . Along this edge, we can measure the discrepancy in their states, which we represent as an element  $\mathbf{g}_{\{ij\}} \in \mathbf{A}$ . This gives us a function

from the 1-simplices (edges) of  $X$  to the group  $A$ , which we can formalize as a **1-cochain**.

For the system to be locally consistent, these measured disagreements must satisfy a compatibility condition. Consider any three processes ( $p_i, p_j, p_k$ ) that form a 2-simplex (a triangle) in  $X$ , meaning they have all mutually communicated their states. The disagreement between  $p_i$  and  $p_k$  must be consistent with the disagreements relayed through  $p_j$ . Algebraically, this translates to the **1-cocycle condition**:

For every 2-simplex  $(i, j, k)$  in  $X$ , the assigned group elements must satisfy:

$$g_{\{ik\}} = g_{\{ij\}} \cdot g_{\{jk\}}$$

Here, for simplicity, we assume the action of  $_1(X)$  on  $A$  is trivial. In the more general case, the condition is  $g_{\{ik\}} = g_{\{ij\}} \cdot (i \cdot g_{\{jk\}})$ , where the action is indexed by the path. A 1-cochain satisfying this condition is called a **1-cocycle**.

The cocycle condition is a statement of local sanity. It ensures that there are no contradictions within any minimal, fully connected cluster of processes. If one process  $p_i$  sees  $p_j$ 's state as  $s_j$  and  $p_j$  sees  $p_k$ 's state as  $s_k$ , then  $p_i$ 's view of  $p_k$ 's state, inferred via  $p_j$ , must match its direct observation (if available).

**1-Coboundaries: Resolvable Disagreements** While the cocycle condition ensures local consistency, it does not guarantee global consensus. A global consensus state exists only if all these local discrepancies can be explained as artifacts of different local perspectives on a single, underlying global state.

Suppose there *is* a way to “correct” each process’s local state to align with a hypothetical global reference frame. We can model this correction as a function  $f$  from the 0-simplices (processes) of  $X$  to the group  $A$ ,  $f: V(X) \rightarrow A$ . The element  $f_i = f(p_i) \in A$  represents the total “error” or “offset” of process  $p_i$  from the ideal consensus state.

If such a global correction map  $f$  exists, it would *induce* a set of local disagreements. The disagreement between  $p_i$  and  $p_j$  would simply be the composition of their respective errors:

$$g_{\{ij\}} = f_i \cdot f_j^{-1}$$

A 1-cocycle  $g$  that can be written in this form is called a **1-coboundary**. A coboundary represents a “trivial” or “resolvable” disagreement. The existence of the function  $f$  (a **0-cochain**) proves that the entire web of disagreements  $g_{\{ij\}}$  is not due to a topological obstruction, but is merely a “gauge artifact”—a result of choosing different local reference points. To achieve consensus, one simply needs to find  $f$  and apply the corrections  $f_i^{-1}$  to each process  $p_i$ . The



problem of consensus is thus equivalent to the problem of determining if a given cocycle of disagreements is, in fact, a coboundary.

**$H^1(X, A)$ : The Set of Consensus Obstructions** Herein lies the central thesis. What if a system is locally consistent (its state is a 1-cocycle) but the disagreement is not resolvable (it is not a 1-coboundary)? This is the signature of a fundamental consensus failure.

We define an equivalence relation on the set of 1-cocycles  $Z^1(X, A)$ : two cocycles  $g$  and  $g'$  are equivalent if they differ by a coboundary, i.e.,  $g' = (f) \cdot g$  for some 0-cochain  $f$ . The set of equivalence classes is the **first nonabelian cohomology set of  $X$  with coefficients in  $A$** , denoted  $H^1(X, A)$ .

- **The Trivial Element:** The class of all 1-coboundaries forms a distinguished “trivial” element in  $H^1(X, A)$ . If the cocycle  $g$  describing the system’s disagreements belongs to this trivial class,  $[g] = [1]$ , then consensus is possible.
- **Nontrivial Elements:** If  $H^1(X, A)$  contains more than one element, then there exist 1-cocycles that are not coboundaries. If the system enters a state described by such a cocycle, it is trapped. No function  $f: V(X) \rightarrow A$  exists that can explain the local disagreements  $g_{\{ij\}}$ . The obstruction is global and topological.

This leads to our primary interpretation:

**$H^1(X, A) \neq \{1\}$  is the necessary and sufficient condition for the existence of topologically obstructed consensus failures.**

Each nontrivial element of  $H^1(X, A)$  corresponds to a distinct *mode* of irresolvable disagreement. The system is locally consistent everywhere, yet no global agreement can be reached.

**The Torsor of Disagreement** When a system is in a state corresponding to a nontrivial class  $[g] \in H^1(X, A)$ , what is the structure of the space of states? It is an  **$A$ -torsor**, which we call the **Torsor of Disagreement**.

An  $A$ -torsor  $T$  is a space on which the group  $A$  acts freely and transitively. This means: 1. For any two states  $t_1, t_2 \in T$ , there is a unique  $a \in A$  such that  $t_2 = a \cdot t_1$ . 2. If  $a \cdot t = t$  for some  $t \in T$ , then  $a$  must be the identity element of  $A$ .

Intuitively, a torsor is a group that has “forgotten its identity element”. It has the full geometric structure of the group  $A$ —all relative distances and transformations are well-defined—but it lacks a canonical origin.

When  $[g]$  is nontrivial, the set of possible global state configurations forms an  $A$ -torsor. We can transition between states of disagreement by applying transformations from  $A$ , but we can never reach the “state of perfect agreement” because no such state exists within the torsor. The torsor itself *is* the obstruction. It

is a globally consistent mathematical object that perfectly models a globally inconsistent distributed state. The system is not broken; it is correctly inhabiting a twisted global structure from which there is no escape via local operations.

**Holonomy: The View from  $\mathcal{X}$**  An equivalent and powerful perspective comes from the fundamental group  $G = \pi_1(\mathcal{X})$ . A 1-cocycle  $g$  on  $\mathcal{X}$  can be integrated along any path. If we integrate  $g$  along a loop starting and ending at a basepoint  $x_0$ , we obtain an element  $\gamma \in A$  known as the **holonomy** of the loop. This  $\gamma$  represents the net disagreement accumulated by an execution that traverses the cycle  $\gamma$ .

The map  $\gamma : \pi_1(\mathcal{X}) \rightarrow A$  derived from a cocycle  $g$  is a **crossed homomorphism**, satisfying  $\gamma(\gamma') = \gamma(\gamma) \cdot (\gamma' \cdot \gamma)$ . A cocycle  $g$  is a coboundary if and only if its corresponding crossed homomorphism is “trivial,” meaning it is of the form  $\gamma(\gamma) = a \cdot (\gamma \cdot a^{-1})$  for some fixed  $a \in A$ .

This gives us a profound re-interpretation of consensus failure:

Consensus is impossible if and only if there exists an execution loop in the protocol complex  $\mathcal{X}$  with a nontrivial holonomy  $\gamma \neq 1$ .

This means a process, or a set of processes, can execute a sequence of perfectly valid local steps that brings them back to their initial configuration, yet find that their view of the world has been permanently altered by a transformation  $\gamma$ . This discrepancy cannot be undone because it is woven into the topology of the protocol itself. The FLP impossibility can be seen as a statement about the existence of such loops in any asynchronous, fault-prone protocol complex.

**Case Study: A Byzantine Torsor** Let’s illustrate with a minimal example of Byzantine failure. \* **Processes:** Three correct processes  $P_1, P_2, P_3$  and one Byzantine process  $B$ . \* **State Group:** They are trying to agree on a value from a nonabelian group  $A$  (e.g.,  $A = S_3$ , the permutation group on 3 elements). \* **Protocol:**  $B$  sends a proposed value to  $P_1$  and  $P_2$ .  $P_1$  and  $P_2$  then forward the value they received to  $P_3$ . \* **Byzantine Behavior:**  $B$  sends value  $a_1$  to  $P_1$  and a different value  $a_2$  to  $P_2$ , where  $a_1, a_2 \in A$  and  $a_1 \neq a_2$ . \* **The Cocycle:** We define a cocycle  $g$  of relative states.  $g_{\{B,1\}} = a_1$ ,  $g_{\{B,2\}} = a_2$ .  $P_1$  reports to  $P_3$ , so  $P_3$ ’s view of  $B$ ’s value via  $P_1$  is  $a_1$ .  $P_2$  reports to  $P_3$ , so  $P_3$ ’s view of  $B$ ’s value via  $P_2$  is  $a_2$ . \* **The Obstruction:**  $P_3$  now has two irreconcilable paths to  $B$ ’s value. This defines a loop in the protocol complex:  $P_3 \rightarrow P_1 \rightarrow B \rightarrow P_2 \rightarrow P_3$ . The holonomy around this loop is  $g_{\{31\}} \cdot g_{\{1B\}} \cdot g_{\{B2\}} \cdot g_{\{23\}}$ . Assuming states are propagated ( $g_{\{ij\}}$  is the value at  $i$  relative to  $j$ ), this holonomy computes to  $a_1 \cdot a_2^{-1}$ . Since  $a_1 \neq a_2$ , the holonomy is nontrivial.

The cocycle  $g$  describing this system state is not a coboundary. There is no global assignment of values  $f_i \in A$  to each process that can resolve the conflict. The system of four processes is trapped in an  $A$ -torsor.  $P_3$  is locked in a state

of irresolvable uncertainty, a direct consequence of the nontrivial element  $[g]$   $H^1(X, A)$  created by the Byzantine fault.

In conclusion, the language of first nonabelian cohomology provides a precise and powerful framework for understanding consensus failures. It elevates the discussion from ad-hoc arguments about message passing to a systematic classification of obstructions via algebraic topology. The Torsor of Disagreement is not merely a metaphor; it is the mathematical reality of a system caught in a state of locally consistent, yet globally irresolvable, disagreement. This insight paves the way for analyzing more complex failure modes, such as those involving the rules of the protocol itself, which will be the subject of our investigation into second cohomology,  $H^2(X, A)$ .

### Chapter 1.3: Non-Abelian Holonomy and Path-Dependent System States

#### Non-Abelian Holonomy and Path-Dependent System States

In the preceding chapter, we established that the first cohomology group,  $H^1(X, \mathcal{A})$ , provides a powerful tool for classifying static obstructions to global consensus. A non-trivial class in  $H^1(X, \mathcal{A})$  corresponds to an “A-torsor of disagreement”—a situation where local states are consistent with all local transition rules, yet no single global state can be defined consistently across the entire system. This model, however, implicitly assumed a relatively simple structure for the state transitions, often one where the order of operations is irrelevant. This is tantamount to the coefficient sheaf  $\mathcal{A}$  being a sheaf of Abelian groups.

In a vast number of realistic distributed systems, this assumption fails. Operations such as database transactions, updates to replicated data structures, or even physical manipulations by robotic agents are inherently non-commutative. Applying update  $g_1$  followed by  $g_2$  can yield a starkly different result from applying  $g_2$  followed by  $g_1$ . This non-commutativity introduces a fundamentally new and dynamic layer of complexity: **path-dependence**. The final state of a process or the system as a whole is no longer just a function of the set of operations performed, but of the specific *sequence* in which they were executed.

This chapter delves into the consequences of non-Abelian state transitions. We will demonstrate that the algebraic concept of **holonomy**, borrowed from differential geometry and topology, provides the precise mathematical language to describe and quantify path-dependent failures. By modeling the group of state updates  $A$  as a non-Abelian group, we will see that topological loops in the protocol complex  $X$  induce non-trivial transformations on the state space. This dynamic manifestation of disagreement, where traversing an execution cycle permanently alters a process’s view of the system state, is the core signature of non-Abelian consensus failures.

**The Algebra of State Updates: Non-Abelian Groups** Let us formalize the notion of a non-commutative state space. We model the set of possible state *transformations* or *updates* as a group, which we will call  $\mathbf{A}$ . A process  $P_i$  maintains a local state  $s_i$ , which is an element of some set  $S$ . The group  $\mathbf{A}$  acts on this set  $S$ . An update, represented by an element  $g \in A$ , transforms the local state via this action:  $s_i \mapsto g \cdot s_i$ .

The critical feature we now introduce is that  $\mathbf{A}$  is a **non-Abelian group**, meaning there exist elements  $g_1, g_2 \in A$  such that  $g_1 \cdot g_2 \neq g_2 \cdot g_1$ .

**Example: Rotational State in a Multi-Agent System** Consider a system of drones tasked with collaboratively tracking the orientation of a satellite. Each drone  $P_i$  maintains its own estimate  $s_i$  of the satellite’s orientation. The state  $s_i$  can be represented as an element of the special orthogonal group  $SO(3)$ , the group of rotations in three-dimensional space. An update message might command a relative rotation, for instance, “rotate 90 degrees around the satellite’s current x-axis” ( $g_x$ ) or “rotate 90 degrees around its y-axis” ( $g_y$ ).  $SO(3)$  is famously non-Abelian. Applying  $g_x$  then  $g_y$  results in a different final orientation than applying  $g_y$  then  $g_x$ . If drone  $P_1$  receives these updates in the order  $(g_x, g_y)$  while drone  $P_2$ , due to network latency, receives them as  $(g_y, g_x)$ , their local state estimates will diverge, even if they started identically and received the same set of updates.

This non-commutativity is the algebraic engine of path-dependence in distributed computations. Asynchrony, a defining characteristic of distributed systems, creates a multitude of possible execution paths corresponding to different interleavings of messages. When the update group  $\mathbf{A}$  is non-Abelian, these distinct topological paths can lead to distinct and irreconcilable final states.

**Holonomy: The Measure of Path-Dependence** To capture the effect of an entire execution history, we introduce the concept of holonomy. In our context, an execution history corresponds to a path  $\gamma$  in the protocol complex  $X$ . Such a path is a sequence of directed 1-simplices,  $\gamma = e_k \circ \dots \circ e_2 \circ e_1$ , where each edge  $e_i = (v_{i-1}, v_i)$  represents a computational step or a message delivery. Associated with each such elementary step is a transformation from our group  $\mathbf{A}$ .

Let the function  $\phi : E(X) \rightarrow A$  map each directed edge  $e$  of the complex to a group element  $\phi(e) \in A$ . This map represents the protocol’s transition rules.

**Definition (Holonomy):** The **holonomy** of a path  $\gamma = e_k \circ \dots \circ e_1$  in the protocol complex  $X$  is the ordered product of the transformations along the path, defined as:

$$\text{Hol}(\gamma) = \phi(e_k) \cdot \dots \cdot \phi(e_2) \cdot \phi(e_1)$$

The holonomy  $\text{Hol}(\gamma)$  is an element of  $\mathbf{A}$  that represents the net transformation accumulated by traversing the execution path  $\gamma$ . If a process starts in state

$s_{start}$  and experiences the sequence of events described by  $\gamma$ , its final state will be  $\text{Hol}(\gamma) \cdot s_{start}$ .

Path-dependence arises when two distinct paths,  $\gamma_1$  and  $\gamma_2$ , connect the same start and end configurations in  $X$  but yield different holonomies:  $\text{Hol}(\gamma_1) \neq \text{Hol}(\gamma_2)$ . In such a scenario, the final state of a process depends entirely on which of the two possible execution histories occurred.

The most revealing instances of path-dependence occur with **loops**. Consider a path  $\gamma$  that is a loop, starting and ending at the same vertex  $v \in X$ . This represents a sequence of operations after which the system's configuration (e.g., which processes have communicated, what messages are in flight) returns to its initial state. In an Abelian system, one would expect the net transformation to be trivial (the identity element). However, in a non-Abelian system, the holonomy of a loop,  $\text{Hol}(\gamma)$ , can be a non-identity element of  $A$ . This means a process can participate in a cycle of interactions and return to its starting point in the protocol, only to find that its internal state has been irrevocably altered. This lingering discrepancy,  $\text{Hol}(\gamma) \in A$ , is a direct measure of the inconsistency generated by the system's topology and non-commutative dynamics.

**Holonomy Representations and the Fundamental Groupoid** The concept of holonomy provides a bridge between the topology of the protocol complex and the algebra of the update group. This connection is most naturally expressed using the **fundamental groupoid** of the space, denoted  $\Pi_1(X)$ .

- The **objects** of  $\Pi_1(X)$  are the vertices (configurations) of the complex  $X$ .
- The **morphisms** between two vertices  $v_1$  and  $v_2$  are the homotopy classes of directed paths from  $v_1$  to  $v_2$ .

Our holonomy mapping can be elevated to a **functor**  $H : \Pi_1(X) \rightarrow A$ , where  $A$  is viewed as a category with a single object and whose morphisms are the elements of the group. This functor maps each path class  $[\gamma]$  to its holonomy  $\text{Hol}(\gamma)$ . The functoriality condition,  $\text{Hol}(\gamma_2 \circ \gamma_1) = \text{Hol}(\gamma_2) \cdot \text{Hol}(\gamma_1)$ , ensures that the composition of paths maps to the product of transformations.

For a fixed basepoint (initial configuration)  $v \in X$ , the morphisms from  $v$  to itself form the **fundamental group**,  $\pi_1(X, v)$ . The holonomy functor restricts to a group homomorphism, often called the **holonomy representation**:

$$h_v : \pi_1(X, v) \rightarrow A$$

This homomorphism maps each topological loop class in the protocol space to an algebraic discrepancy in the state transformation group. A non-trivial homomorphism, where some loop  $\gamma$  maps to  $h_v(\gamma) \neq \mathbf{1}_A$ , is a definitive signature of a systemic, path-dependent failure mode. It signifies a “topological gear” in the machine of the protocol that “slips,” inducing a permanent state error.

**Connection to Non-Abelian Cohomology**  $H^1(X, A)$  The framework of holonomy representations connects directly to the first non-Abelian cohomol-

ogy set,  $H^1(X, A)$ . When the fundamental group  $\pi_1(X)$  acts trivially on the coefficient group  $A$ , the set  $H^1(X, A)$  is defined as the set of conjugacy classes of homomorphisms from  $\pi_1(X)$  to  $A$ :

$$H^1(X, A) \cong \text{Hom}(\pi_1(X, v), A)/A$$

where the group  $A$  acts on the set of homomorphisms by conjugation:  $(\psi^a)(\gamma) = a \cdot \psi(\gamma) \cdot a^{-1}$  for  $a \in A$ .

Each non-trivial element of  $H^1(X, A)$  corresponds to a class of holonomy representations that cannot be “gauged away” or eliminated by a simple, global re-coordination of states (which corresponds to conjugation by an element  $a \in A$ ). A non-trivial class  $[h] \in H^1(X, A)$  implies the existence of a fundamental, topologically-enforced path-dependence that no mere change of “local coordinates” can resolve.

The torsor of disagreement, introduced in the Abelian context, now acquires a rich, dynamic interpretation. The different elements in the fiber of the torsor over a point are not just abstractly possible states; they are concrete states reached via different execution histories. The “distance” between two states  $s_1$  and  $s_2$  in the same fiber can be identified with the holonomy  $\text{Hol}(\gamma)$  of a loop  $\gamma$  such that  $s_2 = \text{Hol}(\gamma) \cdot s_1$ .

**A Case Study: Path-Dependence with Permutation Updates** Let’s illustrate these ideas with a concrete distributed computing scenario.

- **System:** Three processes,  $P_1, P_2, P_3$ , manage a replicated ordered list of three unique items,  $\{X, Y, Z\}$ . The initial state is  $(X, Y, Z)$  at all processes.
- **Update Algebra:** The allowed operations are “swap adjacent items.” The state space can be identified with the symmetric group  $S_3$ , the group of permutations on three elements. This group is non-Abelian. Let’s define two operations:
  - **g\_a** = (1 2): Swap the items at positions 1 and 2.
  - **g\_b** = (2 3): Swap the items at positions 2 and 3. Note that  $g_b \cdot g_a = (23)(12) = (132)$ , a 3-cycle, while  $g_a \cdot g_b = (12)(23) = (123)$ , a different 3-cycle.
- **Protocol and Execution Space:** Imagine a simple protocol where a client can issue swap requests to any process. Consider an asynchronous execution where a client issues **g\_a** and **g\_b** nearly simultaneously. Due to network delays, the processes receive these requests in different orders.
  - **Process  $P_1$**  receives the requests in the order  $(g_a, g_b)$ . Its local state evolves:  $(X, Y, Z) \xrightarrow{g_a} (Y, X, Z) \xrightarrow{g_b} (Y, Z, X)$ . The net transformation is  $g_a \cdot g_b = (123)$ .
  - **Process  $P_2$**  receives the requests in the order  $(g_b, g_a)$ . Its local state evolves:  $(X, Y, Z) \xrightarrow{g_b} (X, Z, Y) \xrightarrow{g_a} (Z, X, Y)$ . The net transformation is  $g_b \cdot g_a = (132)$ .

At the end of this execution,  $P_1$  believes the state is  $(Y, Z, X)$ , while  $P_2$  believes it is  $(Z, X, Y)$ . They have processed the exact same set of updates, but their final states irreconcilably diverge.

- **Topological Interpretation:** This scenario corresponds to a loop in the protocol’s global state space. Let the initial configuration be  $c_0$ . Let  $c_a$  be the configuration after the  $\mathbf{g\_a}$  request is delivered and  $c_b$  be after the  $\mathbf{g\_b}$  request is delivered. Let  $c_{ab}$  be the configuration after both are delivered. There are two paths from  $c_0$  to  $c_{ab}$ :
  - $\gamma_1$ :  $c_0 \rightarrow c_a \rightarrow c_{ab}$  (corresponds to  $P_1$ ’s observation)
  - $\gamma_2$ :  $c_0 \rightarrow c_b \rightarrow c_{ab}$  (corresponds to  $P_2$ ’s observation)

The loop  $\gamma = \gamma_1 \circ \gamma_2^{-1}$  has a non-trivial holonomy. The discrepancy between the two resulting states is precisely the holonomy of this loop:

$$\text{Hol}(\gamma_1) \cdot \text{Hol}(\gamma_2)^{-1} = (g_a \cdot g_b) \cdot (g_b \cdot g_a)^{-1} = (123) \cdot (123) = (132) \neq \text{identity}$$

This non-trivial holonomy, an element of  $S_3$ , is the algebraic witness to the consensus failure. It represents a non-trivial homomorphism from  $\pi_1(X)$  to  $S_3$ , and thus a non-trivial element in  $H^1(X, S_3)$ . The system is plagued by a topological obstruction that manifests as dynamic, path-dependent disagreement.

**Conclusion: Holonomy as the Engine of Dynamic Disagreement** The introduction of non-Abelian groups of transformations moves our analysis from static consensus obstructions to dynamic ones. Path-dependence, a direct consequence of asynchrony and non-commutative operations, is no longer an amorphous problem but a structured phenomenon that can be precisely quantified.

- **Non-Abelian update groups  $\mathbf{A}$**  are the algebraic source of path-dependent states.
- **Holonomy  $\text{Hol}(\cdot)$**  is the mathematical tool to compute the net discrepancy accumulated along any execution path.
- The **holonomy representation  $\mathbf{h}: \pi_1(X) \rightarrow \mathbf{A}$**  captures the system’s global, structural response to race conditions and other cyclic execution patterns.
- Non-trivial elements in the **first non-Abelian cohomology set  $H^1(X, \mathbf{A})$**  classify these path-dependent failure modes, providing a rigorous framework for understanding why certain systems can never guarantee a consistent global state.

By embracing the non-Abelian nature of state, we have uncovered holonomy as the engine driving a dynamic form of disagreement. The torsor of disagreement is no longer a static puzzle but a landscape of possibilities shaped and navigated by execution paths. This deepens our understanding of consensus failures rooted in state ambiguity. However, this is not the final layer of complexity. What if the very rules of state transition—the “gauge transformations” themselves—are ambiguous or inconsistent? This question forces us to ascend the cohomological ladder, to investigate obstructions to the protocol’s structure itself, a domain

governed by the second non-Abelian cohomology  $H^2(X, \mathcal{A})$  and the theory of gerbes.

## Chapter 1.4: Cocycles as Measures of Local-to-Global Inconsistency

### Cocycles as Measures of Local-to-Global Inconsistency

In the preceding chapters, we established a high-level correspondence between the failure to achieve consensus and the non-triviality of the first nonabelian cohomology set,  $H^1(X, \mathcal{A})$ . A non-trivial element of this set corresponds to an  $\mathcal{A}$ -torsor over the protocol complex  $X$ —a geometric object that admits local sections but lacks a global one. The absence of a global section is precisely the impossibility of finding a single state value agreed upon by all processes. While the torsor provides a powerful and elegant geometric picture of disagreement, it remains an abstract concept. For practical analysis and formal verification, we must descend from this abstraction to a more concrete, computable object that captures the same information.

This chapter bridges the conceptual gap between the geometric obstruction (the torsor) and the algebraic data that defines it. This data is encapsulated in the **cocycle**. A cocycle is a function defined on the local components of the protocol complex—its simplices—that measures the “discrepancy” or “transformation” between the views of adjacent processes. It represents the raw, local data of disagreement. The central theme is the classic local-to-global problem: a system can appear perfectly consistent from every local vantage point, yet these local consistencies can fail to cohere into a single, globally consistent state. The cocycle is the mathematical tool that quantifies this failure of integration, making the abstract notion of a topological obstruction manifest and measurable.

### From Torsors to Cocycles: The Local Data of Disagreement

To understand how global disagreement arises from local data, we must formalize the notion of a cocycle within our simplicial framework. Let  $X$  be the protocol complex, whose vertices  $v_i$  represent local states of processes and whose simplices represent sets of compatible, concurrently held local states. Let  $\mathcal{A}$  be the sheaf of nonabelian groups over  $X$  representing the state space; for simplicity, we can often consider a constant sheaf where the group is a single nonabelian group  $A$ .

A **1-cochain** is a function  $g$  that assigns an element of the group  $A$  to each 1-simplex (edge) of  $X$ . If an edge  $\sigma = (v_i, v_j)$  connects the local state of process  $i$  to that of process  $j$ , the value  $g_{ij} = g(\sigma) \in A$  can be interpreted as the transformation required to align the reference frame of  $v_i$  with that of  $v_j$ . It is the “relative state” or “perceived update” between these two points in the execution space.

A 1-cochain does not, by itself, imply any inconsistency. It is merely a collection of local relational data. The crucial structure emerges when we demand local consistency. A 1-cochain  $g$  is a **1-cocycle** if it satisfies the *cocycle condition* on



every 2-simplex (triangle) in  $X$ . For a 2-simplex  $\sigma = (v_i, v_j, v_k)$ , the condition dictates how the transformations along its edges must relate. In the nonabelian, multiplicative notation, this condition is:

$$g_{ij} \cdot g_{jk} = g_{ik}$$

This equation should be read as: the transformation from  $v_i$  to  $v_j$  followed by the transformation from  $v_j$  to  $v_k$  must yield the same result as the direct transformation from  $v_i$  to  $v_k$ .

**Interpretation of the Cocycle Condition:** \* **Local Coherence:** The cocycle condition ensures that there are no contradictions within any elementary “triangle” of observations. If process  $i$ ,  $j$ , and  $k$  can form a consistent collective state (i.e.,  $(v_i, v_j, v_k)$  is a 2-simplex in  $X$ ), then their relative views must be reconcilable. The information gained by composing the path  $i \rightarrow j \rightarrow k$  is identical to the information from the direct path  $i \rightarrow k$ . \* **Gauge Freedom Analogy:** In physics, this is analogous to the consistency of gauge transformations. The value  $g_{ij}$  acts like a gauge transformation connecting the “fiber” (the state space  $A$ ) over point  $v_i$  to the fiber over  $v_j$ . The cocycle condition ensures that these transformations compose associatively over any simply-connected region (a 2-simplex).

The set of all 1-cocycles is denoted  $Z^1(X, A)$ . It is this set that contains all possible configurations of *locally consistent* but potentially *globally inconsistent* states of disagreement.

### The Failure of Global Integration: Cocycles vs. Coboundaries

The satisfaction of the cocycle condition across all 2-simplices is a powerful constraint, but it is deceptively weak. It guarantees coherence over infinitesimal patches of the execution space but provides no guarantee of global coherence. The failure to integrate these local relationships into a global whole is the essence of a consensus failure.

This distinction is formalized by the concept of a **coboundary**. A 1-cocycle  $g$  is called a **1-coboundary** if there exists a **0-cochain**—a function  $s : V(X) \rightarrow A$  assigning a group element  $s_i = s(v_i)$  to each vertex  $v_i$  of the complex—such that for every edge  $(v_i, v_j)$ , the cocycle value is given by:

$$g_{ij} = s_i \cdot s_j^{-1}$$

(Note: The convention could also be  $s_j \cdot s_i^{-1}$ ; what matters is consistency. We choose this convention for its connection to torsor transition functions.)

The existence of such a 0-cochain  $s$  is of profound significance for consensus: \* **Global Section:** The 0-cochain  $s$  is precisely the global section of the state space we have been seeking. The element  $s_i$  can be interpreted as the absolute,

globally-defined state “at” the local process state  $v_i$ . \* **Consensus Achieved:** If such an  $s$  exists, then a global state assignment is possible. The system has reached consensus (or can reach it), and the “discrepancies” measured by the cocycle  $g_{ij}$  are not fundamental disagreements but merely the relative differences between well-defined global states. They are artifacts of the chosen “gauge” or coordinate system, which can be eliminated by choosing the “correct” global frame of reference defined by  $s$ .

A cocycle that is *not* a coboundary is called a **non-trivial** or **non-cobounding cocycle**. The existence of such a cocycle is the definitive signature of a persistent, structural disagreement. It represents a set of locally consistent relative states ( $g_{ij}$ ) that cannot be explained away as mere differences between elements of a single global state.

The first cohomology set  $H^1(X, A)$  is formally defined as the set of 1-cocycles modulo an equivalence relation, where two cocycles  $g$  and  $g'$  are equivalent if they differ by a coboundary. In the nonabelian setting, this is  $g' = s^{-1}g(\partial s)$ , which simplifies to  $g'_{ij} = s_i^{-1}g_{ij}s_j$  for a 0-cochain  $s$ . The distinguished “trivial” element of  $H^1(X, A)$  is the class of all coboundaries. If  $H^1(X, A)$  contains any other element, a global section does not exist, and consensus fails.

The link to the holonomy discussed in the previous chapter is now direct and computational. If we take any closed loop of edges in the protocol complex, say  $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_n \rightarrow v_1$ , the holonomy around this loop is the product of the cocycle values along it:

$$h = g_{12} \cdot g_{23} \cdot \dots \cdot g_{n1}$$

- If the cocycle  $g$  is a coboundary ( $g_{ij} = s_i s_j^{-1}$ ), this product becomes a telescoping series:

$$h = (s_1 s_2^{-1})(s_2 s_3^{-1}) \dots (s_n s_1^{-1}) = s_1 (s_2^{-1} s_2) \dots (s_n^{-1} s_n) s_1^{-1} = s_1 s_1^{-1} = e$$

where  $e$  is the identity element of  $A$ . The holonomy around any loop is trivial.

- If the cocycle is non-trivial, there must exist at least one loop in  $X$  for which this product is a non-identity element of  $A$ . This loop necessarily encloses a “hole” in the protocol complex—a region of impossible executions. The resulting non-identity holonomy is the accumulated, path-dependent disagreement, a concrete measurement of the system’s inability to reconcile its state upon returning to a logical starting point.

### Illustrative Example: Path-Dependence as a Non-Trivial Cocycle

Let us consider a simplified system with three processes,  $P_1, P_2, P_3$ , managing a resource whose state is described by its orientation, an element of the dihedral group  $D_4$  (the symmetries of a square).  $D_4$  is a nonabelian group of order

8. Imagine the “resource” is a shared data structure, and the operations are rotations  $(r_{90}, r_{180}, r_{270})$  and flips.

The protocol complex  $X$  for this system contains a topological hole, representing an execution race. For instance,  $P_1$  wants to apply a 90-degree rotation  $(r_{90})$ , and  $P_2$  wants to apply a horizontal flip  $(f_h)$ . The asynchrony of the system creates a non-contractible loop in the state space corresponding to the two possible orderings of these operations.

Let’s define a 1-cocycle  $g$  on the edges of this loop. The vertices of the loop represent local system views at different stages. 1.  $v_0$ : Initial state. 2.  $v_1$ : View after  $P_1$  has broadcast its intent to apply  $r_{90}$ . 3.  $v_2$ : View after  $P_2$  has broadcast its intent to apply  $f_h$ . 4.  $v_{12}$ : View after both intents are known, but observed in the order  $(P_1, \text{ then } P_2)$ . 5.  $v_{21}$ : View after both intents are known, but observed in the order  $(P_2, \text{ then } P_1)$ .

We can trace a loop in the protocol complex, for example, from a base state  $v_0$ , through states where one or the other operation has been seen, to states where both have been seen, and back. A simplified loop could be visualized as a square whose corners represent states. Let’s define the cocycle as the state transformation along the edges. \* Path A  $(v_0 \rightarrow v_a \rightarrow v_{ab})$ : \*  $g_{0a}$ : The effect of applying  $P_1$ ’s operation,  $r_{90}$ . \*  $g_{ab}$ : The effect of applying  $P_2$ ’s operation,  $f_h$ . \* Path B  $(v_0 \rightarrow v_b \rightarrow v_{ab})$ : \*  $g_{0b}$ : The effect of applying  $P_2$ ’s operation,  $f_h$ . \*  $g_{ba}$ : The effect of applying  $P_1$ ’s operation,  $r_{90}$ .

The cocycle condition,  $g_{ik} = g_{ij}g_{jk}$ , holds on the triangles that make up the surface of the execution space. However, traversing the entire hole reveals the inconsistency. The total transformation along Path A is  $g_{ab} \cdot g_{0a} = f_h \cdot r_{90}$ . The total transformation along Path B is  $g_{ba} \cdot g_{0b} = r_{90} \cdot f_h$ .

Since  $D_4$  is nonabelian,  $f_h \cdot r_{90} \neq r_{90} \cdot f_h$ . The discrepancy between these two paths means that the cocycle cannot be a coboundary. If it were, both paths from  $v_0$  to  $v_{ab}$  would have to evaluate to  $s_0 s_{ab}^{-1}$ , which is impossible as they yield different results. The product of cocycle values around the loop representing this race condition is non-trivial:  $(r_{90} \cdot f_h) \cdot (f_h \cdot r_{90})^{-1} \neq e$ . This holonomy is a direct measure of the path-dependent nature of the system and a concrete manifestation of the non-trivial cocycle classifying the consensus failure.

### Higher-Order Inconsistencies and 2-Cocycles

The framework of cohomology extends naturally to higher dimensions, describing more subtle and complex failures of consensus. We have focused on  $H^1(X, A)$ , which classifies obstructions to finding a global state. This corresponds to a failure of *consensus*.

The next level, governed by the second cohomology group  $H^2(X, A)$ , describes obstructions to the consistency of the cocycles themselves. A **2-cocycle**  $\alpha$  is a function on the 2-simplices of  $X$ . It measures the failure of the 1-cocycle condition. For a 2-simplex  $(v_i, v_j, v_k)$ , we can define:

$$\alpha_{ijk} = g_{ij} \cdot g_{jk} \cdot g_{ik}^{-1}$$

If  $g$  is a 1-cocycle, then  $\alpha_{ijk} = e$  for all triangles. What if this is not the case? This would imply that even the *rules for relating local states* are themselves ambiguous or context-dependent. The transformation from  $v_i$  to  $v_k$  is not just different from the path through  $v_j$ ; it's ill-defined.

This corresponds to an obstruction to what might be called **meta-consensus**: the ability of processes to agree on a consistent model of their disagreements. In such a system, processes might not only fail to agree on a value, but they might also fail to agree on a protocol for resolving their disagreements. These higher obstructions are classified by gerbes, the geometric objects corresponding to non-trivial elements in  $H^2(X, A)$ . While more abstract, they are relevant for analyzing complex, multi-layered protocols where the rules of interaction can themselves be subject to inconsistency.

In conclusion, the cocycle provides the essential link between the abstract topology of distributed executions and the concrete, measurable realities of a running system. It translates the geometric notion of a torsor into a set of local algebraic relations,  $g_{ij}$ . The failure of these relations to be integrable into a global state function  $s_i$  is the essence of consensus failure. By computing the holonomy—the product of cocycle values around loops in the protocol complex—we can directly measure the path-dependent disagreement that plagues non-commuting, asynchronous systems. This provides a powerful, quantitative tool for diagnosing, classifying, and ultimately designing protocols that can navigate the intricate algebra of disagreement.

## Chapter 1.5: Gerbes and Second Cohomology: Obstructions to Protocol Composition

### Gerbes and Second Cohomology: Obstructions to Protocol Composition

**From State Disagreement to Protocol Incoherence** In the preceding chapters, we established a foundational link between the first non-abelian cohomology group,  $H^1(X, \mathcal{A})$ , and the failure of a distributed system to achieve consensus. A non-trivial class in  $H^1(X, \mathcal{A})$  corresponds to an  $\mathcal{A}$ -torsor over the protocol complex  $X$ —the “torsor of disagreement.” This object elegantly captures a situation where local views of the system state are consistent within their own neighborhoods, yet cannot be reconciled into a single, globally defined state. The 1-cocycles  $\{g_{ij}\}$  represent the necessary, but ultimately incompatible, “gauge transformations” required to align the local state descriptions between overlapping execution paths  $U_i$  and  $U_j$ . The obstruction is fundamental: no matter how one redefines local state representations, the global inconsistency, encoded by the topology of  $X$  and the non-abelian nature of the state-update group  $A$ , persists.

This analysis, however, presumes a coherent underlying logic. It assumes that while we may fail to agree on a *state*, we have a consistent set of rules for *transforming* between local viewpoints. The transformations  $g_{ij}$  themselves are assumed to compose in a natural way. But what if this assumption breaks down? What if the protocol's logic is itself fractured, leading to inconsistencies not in the data, but in the rules governing the data? This elevates the problem from a failure of consensus to a failure of meta-consensus—an inability to agree on the protocol itself. To diagnose such deep-seated architectural flaws, we must ascend the cohomological ladder to the second cohomology group,  $H^2(X, \mathcal{A})$ , and its geometric avatar: the gerbe.

### The Second Cohomology Group and the Measure of Inconsistency

Just as  $H^1$  is built from 1-cocycles defined on the edges (1-simplices) of the protocol complex,  $H^2$  is constructed from 2-cocycles defined on its triangles (2-simplices). Let  $X$  be our protocol complex with a suitable open cover  $\{U_i\}$  corresponding to distinct local execution contexts (e.g., processor views, network partitions, or protocol phases).

A **1-cochain** is a collection of maps  $\{g_{ij} : U_i \cap U_j \rightarrow A\}$ , representing the transition logic between contexts  $i$  and  $j$ . This collection is a **1-cocycle** if on every triple overlap  $U_i \cap U_j \cap U_k$ , the composition rule holds:

$$g_{ij}(p) \cdot g_{jk}(p) = g_{ik}(p) \quad \forall p \in U_i \cap U_j \cap U_k$$

This condition ensures that the local state transformations form a coherent system, defining an  $\mathcal{A}$ -torsor. The existence of a global state is equivalent to this 1-cocycle being a **1-coboundary**, i.e.,  $g_{ij} = h_i h_j^{-1}$  for some collection of local functions  $\{h_i\}$ .

Now, let us consider a more complex scenario where the transition rules themselves are not perfectly aligned. We might have a set of local  $\mathcal{A}$ -torsors, one for each  $U_i$ , but find it impossible to glue them together. This means the 1-cocycle condition fails. We can measure this failure with a **2-cochain**, a collection of maps  $\{g_{ijk} : U_i \cap U_j \cap U_k \rightarrow A\}$  defined as:

$$g_{ijk} := g_{ij} \cdot g_{jk} \cdot g_{ki}$$

where  $g_{ki} = g_{ik}^{-1}$ . If the 1-cocycle condition holds, then  $g_{ijk} = 1$  for all  $i, j, k$ . A non-trivial  $g_{ijk}$  is an element of  $A$  that quantifies the “error” or “defect” in composing the transition logic around the boundary of the 2-simplex  $(i, j, k)$ . It is the holonomy of the transition logic itself.

For this collection of defects  $\{g_{ijk}\}$  to represent a stable, global obstruction, it must satisfy its own consistency condition. This condition, the **2-cocycle condition**, is defined on tetrahedra (3-simplices) in the protocol complex. For any four overlapping contexts  $U_i, U_j, U_k, U_l$ , it states:

$$(\partial g)_{ijkl} = g_{jkl} \cdot g_{ikl}^{-1} \cdot g_{ijl} \cdot g_{ijk}^{-1} = 1$$

(Note: The precise formula depends on group-theoretic and categorical conventions, but the concept remains the same). This equation ensures that the “defects” measured on the four faces of the tetrahedron  $(i, j, k, l)$  are mutually consistent. A 2-cochain satisfying this condition is a **2-cocycle**.

The **second non-abelian cohomology set**, denoted  $H^2(X, \mathcal{A})$ , is the set of equivalence classes of these 2-cocycles, where two 2-cocycles are equivalent if they differ by a 2-coboundary. A non-trivial element in  $H^2(X, \mathcal{A})$  signifies a fundamental, irremovable inconsistency in the protocol’s transition logic.

**Gerbes: The Geometry of Inconsistent Logics** What kind of object does a non-trivial class in  $H^2(X, \mathcal{A})$  represent? The answer is a **gerbe**. While a torsor is a “bundle of points” that locally looks like the group  $A$  but may lack a global point (section), a gerbe is a “stack” or “2-bundle” whose local structure is that of a torsor.

- **Torsor ( $H^1$  obstruction):** There is no globally consistent *state*. It is locally like asking, “What is the state?” and getting a valid answer, but the answers from different locations can’t be reconciled.
- **Gerbe ( $H^2$  obstruction):** There is no globally consistent *torsor of disagreement*. It is locally like asking, “What is the set of allowed state transformations?” and getting a consistent answer (a local torsor), but the transformation rules between these local sets of rules are themselves inconsistent.

A gerbe represents a system where the very notion of “disagreement” is ill-defined globally. On each patch  $U_i$ , the system’s logic is sound (it is described by a torsor, possibly trivial). On overlaps  $U_i \cap U_j$ , we have a way to translate between these logical frameworks, given by  $g_{ij}$ . However, on triple overlaps  $U_i \cap U_j \cap U_k$ , the translation from  $i \rightarrow j \rightarrow k$  does not match the translation from  $i \rightarrow k$ . The discrepancy is precisely the 2-cocycle  $g_{ijk}$ . This is an obstruction to “lifting” the local torsors into a single, global torsor over  $X$ . The system is afflicted not by a simple disagreement over state, but by a fundamental schizophrenia in its logic.

**Interpretation: Obstructions to Protocol Composition** The most natural interpretation of a gerbe-like obstruction in distributed systems arises in the context of **protocol composition**. Modern distributed systems are rarely monolithic; they are composed of multiple sub-protocols, each governing a different aspect of the system’s behavior. Examples include:

1. A core consensus protocol (e.g., Raft, Paxos).
2. A membership reconfiguration protocol for adding/removing nodes.
3. A failure detection and recovery protocol.
4. A data sharding and load balancing protocol.

The system transitions between these operational modes based on internal or external events. The protocol complex  $X$  can be seen as a model of these

transitions, where each patch  $U_i$  corresponds to the system operating under a specific sub-protocol  $P_i$ .

- **1-Cochains ( $g_{ij}$ ):** These represent the **interface protocol** or **translation logic** for switching from protocol  $P_i$  to  $P_j$ . For example, when a node failure is detected (transitioning from “normal operation”  $P_i$  to “recovery mode”  $P_j$ ),  $g_{ij}$  would encode the rules for freezing logs, transferring leadership roles, and initializing the recovery state.
- **2-Cocycles ( $g_{ijk}$ ):** A non-trivial 2-cocycle  $g_{ijk} \neq 1$  represents a **composition anomaly**. It signifies that the sequence of transitions  $P_i \rightarrow P_j \rightarrow P_k$  results in a system state or configuration that is incompatible with the state that would result from a direct transition  $P_i \rightarrow P_k$  (or any other path). The element  $g_{ijk} \in A$  is the algebraic “error term” that must be applied to reconcile the outcomes.

### Example: Composing Consensus and Reconfiguration

Consider a Raft-based system. \* Let  $P_i$  be “Normal Operation” in configuration  $C_1 = \{N_1, N_2, N_3\}$ . \* Let  $P_j$  be the “Reconfiguration Protocol” to transition to  $C_2 = \{N_1, N_2, N_4\}$ . \* Let  $P_k$  be “Normal Operation” in configuration  $C_2$ .

The transition  $P_i \rightarrow P_j$  involves the current leader proposing a new configuration entry in its log. The transition  $P_j \rightarrow P_k$  occurs once this entry is committed. The combined path  $i \rightarrow j \rightarrow k$  defines the standard Raft procedure for changing cluster membership.

Now, imagine a more complex scenario involving a third protocol,  $P_l$ , such as an emergency “snapshot-based recovery” protocol triggered by suspected data corruption. \*  $U_i$ : Normal operation. \*  $U_j$ : Standard reconfiguration. \*  $U_k$ : Snapshot recovery.

Suppose a sequence of events leads to a transition path  $i \rightarrow j \rightarrow k$ : the system begins a standard reconfiguration, but during this delicate phase, a corruption alert triggers the snapshot recovery. The interface logic  $g_{jk}$  must define how to safely abort the reconfiguration and apply the snapshot.

A non-trivial 2-cocycle  $g_{ijk}$  would emerge if, for instance, the state resulting from the aborted reconfiguration followed by recovery (**path i→j→k**) has a different leader epoch, set of committed logs, or voter configuration than if the system had transitioned directly from normal operation to snapshot recovery (**path i→k**). The element  $g_{ijk}$  represents this discrepancy—perhaps an off-by-one error in an epoch number, or an ambiguity in which nodes have voting rights.

A non-trivial class in  $H^2(X, \mathcal{A})$  implies this is not a simple bug that can be patched at one interface. It is a systemic flaw. Fixing the  $g_{ijk}$  inconsistency by changing the  $g_{jk}$  interface might simply shift the anomaly, creating a new non-trivial 2-cocycle  $g_{jkl}$  at another protocol intersection. The gerbe is the global manifestation of this inescapable web of local inconsistencies. It is an obstruction baked into the architecture, a proof that the sub-protocols cannot be composed into a coherent whole.

**Conclusion: Higher Obstructions and the Limits of Local Fixes** The progression from  $H^1$  to  $H^2$  provides a powerful diagnostic framework for classifying distributed system failures:

- $H^0$  (**non-trivial center**): Obstruction to reaching a *unique* global state. Multiple consensus values are possible (**k-set consensus**).
- $H^1$  (**non-trivial**): Obstruction to reaching *any* global state. A “torsor of disagreement” exists, defined by incompatible local views. The protocol logic is sound, but the state is fractured.
- $H^2$  (**non-trivial**): Obstruction to having a *coherent protocol logic*. A “gerbe of incoherence” exists, defined by incompatible rules for transitioning between local operational contexts. The very definition of state and disagreement becomes locally defined but globally inconsistent.

Gerbes and second cohomology formalize the intuition that the most pernicious bugs in complex systems are not isolated errors but emergent properties of component interactions. They arise from the failure to ensure that local interface agreements compose globally. By identifying the topological and algebraic sources of these higher-order failures, this framework provides not only a new language for describing them but also a guide for designing future protocols that are provably robust against such deep architectural flaws. The presence of a non-trivial gerbe is a mathematical certificate of a system’s inevitable descent into logical chaos.

## Part 2: The Directed Topology of Asynchronous Computations

### Chapter 2.1: Beyond Reversibility: Directed Spaces for Asynchronous Systems

Beyond Reversibility: Directed Spaces for Asynchronous Systems

The topological modeling of distributed systems, particularly through the lens of simplicial complexes, has provided profound insights into the nature of computability. By representing the global states of a protocol as simplices, we transform questions of task solvability into questions of topological connectivity and contractibility. The seminal proof of the impossibility of consensus in the presence of a single crash fault (the FLP result) and its topological counterparts have demonstrated the power of this abstraction. They reveal that consensus is obstructed by non-trivial “holes” in the protocol complex, which prevent processes from converging to a common decision value.

However, this classical framework, for all its power, rests on a fundamental simplification that elides a crucial aspect of computation: **causality**. Standard algebraic topology is built upon the notion of a path, a continuous map from the unit interval  $[0,1]$  into a space  $\mathbf{X}$ . Crucially, if a path from point  $\mathbf{a}$  to  $\mathbf{b}$  exists, so does its inverse  $^{-1}$  from  $\mathbf{b}$  to  $\mathbf{a}$ . The fundamental group  $(\mathbf{X}, \mathbf{x})$  and its associated homology groups treat paths as reversible entities. This



assumption of reversibility is deeply at odds with the physical and logical reality of asynchronous systems.

In a distributed computation, events—sending a message, receiving a message, updating a local state—are causally ordered. Time, as experienced by the system, flows inexorably forward. An execution is a sequence of events, and this sequence cannot simply be run in reverse. While one might be able to define an operation that “undoes” a previous one, this is itself a new, distinct forward-moving computation, not a reversal of the original path through the state space. The very essence of asynchrony is the uncertainty of the ordering of concurrent events, leading to a partial order of states, not a symmetric graph of transitions. Modeling executions with tools that assume reversibility thus erases the arrow of time, conflating what *is reachable* with what *has happened*. To capture the true geometry of asynchronous computation, we must move beyond reversible spaces to the more nuanced framework of directed topology.

**From Paths to Dipaths: The Structure of Directed Spaces** A **directed space**, or **d-space**, is the formal structure that reintroduces causality into topology. It consists of a topological space  $X$  endowed with a privileged collection of paths, called **directed paths** or **dipaths**, which represent the “allowed” trajectories or executions. These dipaths must satisfy two intuitive axioms: 1. **Constants:** All constant paths are dipaths. A system can remain in a state. 2. **Concatenation:** If  $\gamma$  is a dipath from  $x$  to  $y$  and  $\delta$  is a dipath from  $y$  to  $z$ , then their concatenation  $\gamma * \delta$  is a dipath from  $x$  to  $z$ . Executions can be composed sequentially.

Crucially, the inverse of a non-constant dipath is typically *not* a dipath. For instance, the directed interval  $\uparrow[0,1]$  is the standard unit interval where the dipaths are only the non-decreasing functions. One can travel from 0 to 1, but not from 1 to 0 along an allowed path.

This simple modification has profound consequences. In the context of a protocol complex  $X$ , the direction is naturally induced by the progression of the computation. An  $n$ -simplex  $\sigma = (p : v, \dots, p : v)$  represents a global state where  $n+1$  processes have participated. Its faces represent prior states. For instance, the face  $\sigma_{\bar{p}}$  is the state before process  $p$  contributed its value  $v$ . Thus, there is a natural directed path from the barycenter of  $\sigma_{\bar{p}}$  to the barycenter of  $\sigma$ , representing the execution of a single computational step. The entire protocol complex becomes a d-space where executions are chains of such elementary dipaths, always moving from lower-dimensional simplices (fewer events) to higher-dimensional ones (more events).

**The Fundamental Category:  $\text{Cat}(X)$**  The most immediate casualty of abandoning reversibility is the algebraic structure used to classify paths. In classical topology, the set of homotopy classes of loops based at a point  $x$  forms the fundamental group  $\pi_1(X, x)$ . The invertibility of paths ensures that every element has an inverse, satisfying the group axiom.

In a directed space, this is no longer true. A **dihomotopy** is a homotopy that proceeds only along dipaths. Two dipaths are dihomotopic if one can be continuously deformed into the other through a family of intermediate dipaths. The resulting equivalence classes of dipaths do not form a group. A directed loop from  $x$  to  $x$  may not be dihomotopic to the constant path, and it may not have an inverse.

The natural algebraic object that emerges is not a group, but a **category**. The **fundamental category** of a d-space  $X$ , denoted  $\text{Cat}(X)$ , has the points of  $X$  as its objects. A morphism from  $x$  to  $y$  is a dihomotopy class of dipaths from  $x$  to  $y$ . Composition of morphisms is defined by concatenation of representative dipaths.

Unlike the fundamental groupoid  $\Pi(X)$  of a standard space, where every morphism is an isomorphism,  $\text{Cat}(X)$  contains non-invertible morphisms. The existence of a morphism from  $x$  to  $y$  means that state  $y$  is computationally reachable from state  $x$ . The absence of a morphism from  $y$  to  $x$  captures the irreversibility of the process. This categorical perspective allows for a much richer description of a system's dynamics: - **Reachability**:  $\text{Hom}(x, y)$  means  $y$  is reachable from  $x$ . - **Deadlocks**: A state  $x$  is a deadlock if  $\text{Hom}(x, y) = \emptyset$  for all  $y \neq x$ . It is a terminal object in its component of the category. - **Concurrency**: If two dipaths from  $x$  to  $y$  are not dihomotopic, it indicates two genuinely different, non-interchangeable computational pathways between the same initial and final states. This is a geometric manifestation of non-commutativity.

**Directed Homology and Cohomological Obstructions** Just as the fundamental group gives way to the fundamental category, standard homology theories can be refined into directed variants. While a full exposition is beyond our scope, the intuition is that directed homology measures “features” of the state space that respect causality. It can detect phenomena invisible to standard homology, such as states that are “sources” or “sinks” of execution flow, or cycles that can only be traversed in one direction.

This directed framework provides the natural substrate for the nonabelian cohomology theory of consensus failures. Let  $X$  be the protocol complex, viewed as a d-space, and let  $A$  be the nonabelian group of possible state transformations (e.g., updates to a shared object). A **local section** of a potential global state can be represented as a function  $s$  on the vertices (processes) of  $X$ , with  $s(p) \in A$  being the local state of process  $p$ .

Now, consider an edge  $e = (p, q)$  in the complex, representing a communication or ordering of events. The system's specification may demand a transition function  $g_e : A \rightarrow A$  that relates the states across this edge. For a global state  $s$  to be consistent, we would require  $s(q) = s(p) \cdot g_e$ . This is the familiar setup.

The directed nature of the space now becomes critical. An execution is a dipath in  $X$ . As we traverse this path, the local views are transformed. The discrepancy accumulated along a closed loop is its **holonomy**—the group element

in  $A$  obtained by composing the transformations along the loop. In a standard topological setting, one could traverse the loop in reverse to “undo” this holonomy. In a d-space, this is forbidden. The holonomy represents an *irreversible* desynchronization.

The 1-cocycles of nonabelian cohomology,  $f: G \rightarrow A$  where  $G = \text{Cat}(X)$  and  $f$  satisfies  $f(gh) = f(g) \cdot g \cdot f(h)$ , are precisely the functions that measure this path-dependent discrepancy. The cocycle condition is a statement about composing *dipaths*. A nontrivial cohomology class in  $H^1(X, A)$  corresponds to the existence of a “twisted” state assignment—an  $A$ -torsor—that is locally consistent but globally irreconcilable. The directed topology reveals that this is not just a static “hole” but a dynamic “vortex” in the state space. Any execution path that winds around this obstruction accumulates a permanent, irreversible phase error in the system’s state.

**Case Study Revisited: k-Set Consensus** Let us reconsider the problem of  $k$ -set consensus, where processes must agree on at most  $k$  different values. The protocol complex for this problem is known to be  $(n-k-1)$ -connected, where  $n$  is the number of processes. For  $(k+1)$ -process executions, this implies the existence of topological spheres of dimension  $n-(k+1)$ .

In a classical view, these spheres are simply obstructions. In the directed view, they become causal traps. Consider the case of 2-set consensus among 4 processes ( $n=4$ ,  $k=2$ ). The protocol complex contains 1-spheres (loops). A dipath corresponding to an execution might enter a region corresponding to a 3-valent state (e.g., processes  $p$ ,  $p$ ,  $p$  have proposed distinct values  $v$ ,  $v$ ,  $v$ ). From here, the system is in a precarious position. The directed topology of the complex shows that all available dipaths might lead to states where a fourth process  $p$  observes one of these three values, perpetuating the disagreement. There may be no directed path from this region of “high contention” to a 2-valent (decided) region that does not violate the protocol’s rules (e.g., by dropping a message or ignoring a value, which are not modeled as valid steps). The impossibility arises not just because a “hole” exists, but because executions are “sucked” into it by the forward flow of time, with no allowed path leading out to a valid, decided state for all participants. The obstruction is a “waterfall,” not a “donut.”

**Conclusion: The Geometry of Irreversibility** The transition from classical to directed topology is not a mere technicality; it is a fundamental paradigm shift required to faithfully model the physics of information in asynchronous systems. By respecting causality, directed topology provides a geometry of irreversibility.

- It replaces the symmetric notion of path-connectedness with the asymmetric notion of **reachability**.
- It elevates the algebraic model from the fundamental group  $\pi_1(X)$  to the richer **fundamental category**  $\text{Cat}(X)$ , where non-invertible morphisms encode irreversible computational steps.

- It provides the natural setting for understanding **holonomy** as an accumulated, path-dependent, and permanent discrepancy in system state.

This framework demonstrates *why* nonabelian algebraic structures are the correct tools for analyzing distributed consensus. The non-commutativity of state updates and the irreversibility of time are two sides of the same coin, and they find their unified geometric expression in the directed structure of the execution space. The cohomological obstructions, the **A**-torsors and gerbes that we analyze in subsequent chapters, are not abstract algebraic artifacts. They are the precise mathematical shadows cast by the causal, irreversible, and often paradoxical geometry of asynchronous computation. Directed topology provides the light source that makes these shadows intelligible.

## Chapter 2.2: The Fundamental Category ( $\text{Cat}_1(X)$ ) and Causal Paths

### The Fundamental Category ( $\text{Cat}_1(X)$ ) and Causal Paths

The transition from standard topological spaces to directed spaces, as motivated in the previous chapter, necessitates a corresponding evolution in our algebraic-topological tools. The classical fundamental group,  $\pi_1(X)$ , while powerful, is fundamentally anchored in the notion of invertible paths. In the context of asynchronous systems, where computation is irreversible and causality imposes a strict temporal ordering, this assumption is untenable. An event cannot be “un-happened,” and a message, once received, cannot be causally “un-received.” To capture this intrinsic directionality, we replace the fundamental group with a richer algebraic structure: the **fundamental category**, denoted  $\text{Cat}_1(X)$  or sometimes  $\Pi_1(X)$ .

**From Reversible Loops to Directed Morphisms** The fundamental group  $\pi_1(X, x_0)$  of a space  $X$  with basepoint  $x_0$  is the group of homotopy classes of loops starting and ending at  $x_0$ . The group structure relies on two key operations: path concatenation and path inversion. Path inversion, the ability to traverse any path  $\gamma$  backwards as  $\gamma^{-1}$ , is precisely what breaks down in a directed setting. A valid sequence of computational events from state  $A$  to state  $B$  does not guarantee the existence of a valid sequence from  $B$  to  $A$ .

The fundamental category  $\text{Cat}_1(X)$  resolves this by dispensing with the requirement of a single basepoint and the universal invertibility of paths. It instead provides a global picture of all possible causal evolutions between *all* pairs of states in the system.

**Definition of the Fundamental Category** For a given directed space  $X$ , its fundamental category  $\text{Cat}_1(X)$  is defined as follows:

- **Objects:** The objects of  $\text{Cat}_1(X)$  are the points of the space  $X$ .  $\text{Obj}(\text{Cat}_1(X)) = X$ . In our computational model, each object represents a distinct global state of the distributed system.

- **Morphisms:** The morphisms of  $\text{Cat}_1(X)$  are the *dihomotopy classes* of directed paths.  $\text{Mor}(\text{Cat}_1(X)) = \{[\gamma] \mid \gamma \text{ is a dipath in } X\} / \sim_{di}$ . A morphism from an object  $x$  to an object  $y$  is an equivalence class of directed paths starting at  $x$  and ending at  $y$ . We denote the set of all such morphisms as  $\text{Hom}_{\text{Cat}_1(X)}(x, y)$ .

To make this definition precise, we must clarify the concepts of directed paths and directed homotopy.

**Directed Paths (Dipaths):** A dipath in  $X$  is a continuous map  $\gamma : [0, 1] \rightarrow X$  that respects the directed structure of  $X$ . This means that for any  $s_1 \leq s_2$  in  $[0, 1]$ , the point  $\gamma(s_2)$  is in the causal future of  $\gamma(s_1)$ . For a protocol complex, this means the path only traverses simplices in the direction of increasing dimension or along allowed “edges.” A dipath is a model of a valid, causally consistent execution trace of the distributed system.

**Directed Homotopy (Dihomotopy):** Two dipaths,  $\gamma_0$  and  $\gamma_1$ , with the same start point  $x$  and end point  $y$ , are said to be dihomotopic if there exists a continuous map  $H : [0, 1] \times [0, 1] \rightarrow X$  such that: 1.  $H(s, 0) = \gamma_0(s)$  and  $H(s, 1) = \gamma_1(s)$  for all  $s \in [0, 1]$ . (The map deforms  $\gamma_0$  into  $\gamma_1$ ). 2.  $H(0, t) = x$  and  $H(1, t) = y$  for all  $t \in [0, 1]$ . (The endpoints are fixed throughout the deformation). 3. **Crucially**, for each fixed  $t \in [0, 1]$ , the path  $s \mapsto H(s, t)$  is itself a dipath.

This third condition is the essence of dihomotopy. The continuous deformation from one execution trace to another must consist entirely of valid execution traces. It prevents deformations that would violate causality.

**Composition and Identity:** \* **Composition:** The composition of two morphisms  $[\gamma_1] : x \rightarrow y$  and  $[\gamma_2] : y \rightarrow z$  is defined by the concatenation of their representative paths. If  $\gamma_1$  is a path from  $x$  to  $y$  and  $\gamma_2$  is a path from  $y$  to  $z$ , their concatenation  $\gamma_2 \circ \gamma_1$  is a dipath from  $x$  to  $z$ . This composition is well-defined on dihomotopy classes. \* **Identity:** For any object  $x \in X$ , the identity morphism  $id_x : x \rightarrow x$  is the dihomotopy class of the constant path  $\gamma(s) = x$  for all  $s \in [0, 1]$ .

**The Causal and Noncommutative Structure of  $\text{Cat}_1(X)$**  The fundamental category is not merely an abstract construction; it is a direct encoding of the causal and logical structure of a distributed computation.

- **Causality Encoded:** The existence of a morphism in  $\text{Cat}_1(X)$  is a statement about causality. The set  $\text{Hom}_{\text{Cat}_1(X)}(x, y)$  is non-empty if and only if the state  $y$  is causally reachable from the state  $x$ . The partial order of computation is thus embedded in the connectivity of the category. If this set is empty, no sequence of valid operations can transform the system from state  $x$  to state  $y$ .
- **Irreversibility:** The non-invertibility of morphisms directly models the arrow of time. The existence of a morphism from  $x$  to  $y$  makes no claim

about the existence of one from  $y$  to  $x$ . This distinguishes  $\text{Cat}_1(X)$  from the standard fundamental groupoid, which is formed by considering all paths to be reversible and is a less faithful model for computation.

- **Path Dependence and Non-Uniqueness:** A key feature of asynchronous systems with non-commuting operations is that the final state may depend on the order of execution.  $\text{Cat}_1(X)$  captures this phenomenon with precision. The existence of two *distinct* morphisms in  $\text{Hom}_{\text{Cat}_1(X)}(x, y)$ —that is, two dipaths from  $x$  to  $y$  that are not dihomotopic—signifies that there are at least two fundamentally different, causally sound ways for the system to evolve from state  $x$  to state  $y$ . These different histories, corresponding to different interleavings of concurrent operations, are the root cause of path-dependent disagreement. Consensus protocols are, in essence, attempts to ensure that despite the potential for multiple execution paths, all reachable decision states are equivalent in some meaningful way.

**Causal Cycles and Holonomy** In a directed space, it is still possible to have directed loops—dipaths that start and end at the same point,  $\gamma : x \rightarrow x$ . These are not time-travel paradoxes. Rather, they represent sequences of events (e.g., concurrent operations by different processes) that return the system to the same global configuration. For each state  $x$ , the set of morphisms from  $x$  to itself,  $\text{End}_{\text{Cat}_1(X)}(x) = \text{Hom}_{\text{Cat}_1(X)}(x, x)$ , forms a monoid under composition (a semigroup with an identity element, but not necessarily with inverses).

These endomorphism monoids are of profound importance. They are the directed analogue of the fundamental group at a point. A non-trivial element in  $\text{End}_{\text{Cat}_1(X)}(x)$  corresponds to a causal cycle in the state space. While such a cycle may return the system to global state  $x$ , the local views of the constituent processes may have been altered in a path-dependent manner.

This leads us to the crucial concept of **holonomy**. In differential geometry, holonomy measures the transformation of a vector as it is parallel-transported around a closed loop. In our context, we are not transporting vectors, but **local process states**, which are elements of some, possibly nonabelian, group  $A$ .

Let’s imagine that each directed edge in our protocol complex is labeled with an element from a group  $A$ , representing the state update performed by a process. A dipath  $\gamma$  in the space  $X$  corresponds to a sequence of such updates. The **holonomy of the path**  $\gamma$  is the total composed transformation obtained by traversing it.

Now consider a causal cycle  $\gamma : x \rightarrow x$ . If the group  $A$  of state updates is nonabelian, the holonomy around this cycle, denoted  $\text{hol}(\gamma) \in A$ , may be non-trivial. That is,  $\text{hol}(\gamma) \neq e$ , where  $e$  is the identity in  $A$ .

**Interpretation for Consensus Failures:** A non-trivial holonomy around a causal cycle is a definitive signature of disagreement. It means that there exists

a valid sequence of operations after which the global state appears unchanged ( $x \rightarrow x$ ), but a process’s local view of the state has been irrecoverably altered.

For instance, consider two processes, P1 and P2, starting in a state where their local values are  $v_1$  and  $v_2$ . Suppose they are part of a system that undergoes a sequence of operations corresponding to a causal cycle  $\gamma$  based at the initial state  $x$ . If P1’s local state is transformed by  $\text{hol}(\gamma)$  while P2’s is not (perhaps P2 was not involved in all operations of the cycle), their final states will be  $v'_1 = v_1 \cdot \text{hol}(\gamma)$  and  $v'_2 = v_2$ . Since  $\text{hol}(\gamma) \neq e$ , we have  $v'_1 \neq v_1$ . If consensus required P1 and P2 to agree, this path-dependent discrepancy, invisible at the level of the global state configuration, makes agreement impossible. The system contains a hidden “topological defect” in its execution space.

**Bridging to Nonabelian Cohomology** The structure of the fundamental category provides the precise geometric framework for understanding these defects. The holonomy map, which assigns a group element in  $A$  to each causal cycle, is the key.

If we temporarily ignore directionality and consider the fundamental group  $\pi_1(X, x)$ , a holonomy map is precisely a group homomorphism  $\rho : \pi_1(X, x) \rightarrow A$ . The set of such homomorphisms, up to conjugacy in  $A$ , is in bijective correspondence with the first nonabelian cohomology set,  $H^1(X, A)$ .

Therefore, the study of  $\text{Cat}_1(X)$  and its endomorphism monoids naturally leads us to cohomology. The existence of causal cycles with non-trivial holonomy implies that  $H^1(X, A)$  is non-trivial. The elements of this cohomology set classify the fundamental, topologically-enshrined sources of disagreement within the system. Each distinct element of  $H^1(X, A)$  corresponds to a unique “twist” in the state space—a global, self-consistent pattern of potential disagreement that cannot be eliminated by any local protocol actions.

In summary, the fundamental category  $\text{Cat}_1(X)$  serves as the essential bridge between the intuitive, operational model of asynchronous computation and the abstract, powerful machinery of nonabelian cohomology. It replaces the notion of reversible paths with causal morphisms, capturing the structure of all possible executions. Its internal structure, particularly the existence of multiple, non-dihomotopic paths and the holonomy of causal cycles, provides a direct topological interpretation for path-dependence and consensus failures. The following chapters will formalize this connection by defining cohomology directly on the protocol space, using the insights gained from the causal structure revealed by  $\text{Cat}_1(X)$ .

### Chapter 2.3: Geometric Semantics: Concurrency, Deadlock, and Directed Holes

Geometric Semantics: Concurrency, Deadlock, and Directed Holes

Having established the directed space of executions,  $\mathbf{X}$ , and its fundamental cat-

egory,  $\text{Cat}_1(X)$ , as our formal model, we now move from abstract structure to computational meaning. The power of the topological approach lies in its ability to provide a *geometric semantics* for the behavior of asynchronous systems. In this framework, fundamental computational concepts such as concurrency, resource contention, deadlock, and liveness failures are not merely properties to be verified but are manifest as intrinsic, visualizable geometric features of the protocol complex itself.

This chapter establishes the core of this geometric dictionary. We will demonstrate that concurrency is embodied by the “volume” of higher-dimensional cells, deadlocks correspond to terminal boundaries in the execution space, and more subtle consensus failures—states of irreconcilable disagreement—are captured by a novel concept: the *directed hole*. These holes are not merely absent regions but are specific topological configurations that signify a failure of local information to cohere into a global reality, a geometric precursor to the cohomological obstructions we will develop in later chapters.

### The Geometry of Concurrency: Dihomotopy and Higher-Dimensional Cells

In an asynchronous system, the defining characteristic of concurrency is the lack of a predetermined total order on independent events. If process  $P_1$  executes action  $a$  and process  $P_2$  executes action  $b$ , and these actions do not causally depend on each other (e.g., they operate on different data or non-contentious resources), then the system may evolve through the sequence  $a$  then  $b$ , or  $b$  then  $a$ . Both execution paths are valid and, crucially, should lead to the same resulting state.

This scenario has a direct and elegant geometric representation. Let the initial state be a 0-simplex (a vertex)  $v_0$ . The execution of  $a$  corresponds to a 1-simplex (a directed edge)  $e_a: v_0 \rightarrow v_a$ , and the execution of  $b$  corresponds to another 1-simplex  $e_b: v_0 \rightarrow v_b$ . Since the actions are independent, executing  $b$  from state  $v_a$  is possible, leading to a state  $v_{\{ab\}}$  via an edge  $e'_b$ . Similarly, executing  $a$  from  $v_b$  leads to  $v_{\{ba\}}$  via  $e'_a$ . The fundamental property of concurrency is that  $v_{\{ab\}} = v_{\{ba\}}$ .

The resulting geometric figure is a commutative square:

$$\begin{array}{ccc}
 & e'_b & \\
 v_a & \dashrightarrow & v_{ab} \\
 e_a \downarrow & & \downarrow e'_a \\
 v_0 & \dashrightarrow & v_b \\
 & e_b & 
 \end{array}$$

This square can be interpreted as a 2-dimensional cell, specifically a 2-simplex (if triangulated) or a 2-cube. The two distinct directed paths from  $v_0$  to  $v_{\{ab\}}$ — $e_a$  followed by  $e'_b$ , and  $e_b$  followed by  $e'_a$ —are *dihomotopic*. They are two sides of the same “fabric” of execution. The existence of this filled-in 2-cell



is the geometric signature of the concurrency of **a** and **b**. The interior of the cell represents the continuous deformation of one execution schedule into another, embodying the semantic equivalence of all interleavings.

This generalizes to higher dimensions. The concurrent execution of **n** independent actions corresponds to the existence of an **n**-dimensional cube (or **n**-simplex) in the protocol complex **X**. The “volume” of the space of executions is a direct measure of its degree of concurrency. In the language of the fundamental category  $\mathbf{Cat}_1(\mathbf{X})$ , concurrency manifests as the existence of multiple, dihomotopic morphisms between two objects (states).

**Deadlock as a Terminal Boundary** Concurrency represents the ideal, conflict-free nature of asynchronous execution. The first and most classic deviation from this ideal is deadlock, which arises from resource contention. Consider two processes, **P<sub>1</sub>** and **P<sub>2</sub>**, competing for two mutex locks, **L<sub>1</sub>** and **L<sub>2</sub>**. **P<sub>1</sub>**’s protocol is `lock(L1); lock(L2)`, while **P<sub>2</sub>**’s is `lock(L2); lock(L1)`.

Let us trace the geometry of this system’s execution space: 1. **Initial State v<sub>0</sub>**: Neither process holds a lock. 2. **Path 1**: **P<sub>1</sub>** acquires **L<sub>1</sub>**. The system moves to state **v<sub>1</sub>**. This is a directed edge **e<sub>1</sub>**: **v<sub>0</sub>** → **v<sub>1</sub>**. 3. **Path 2**: **P<sub>2</sub>** acquires **L<sub>2</sub>**. The system moves to state **v<sub>2</sub>**. This is a directed edge **e<sub>2</sub>**: **v<sub>0</sub>** → **v<sub>2</sub>**.

These two initial actions are concurrent. We might therefore expect them to form the edges of a 2-cell, as in the previous example. However, let’s examine the subsequent states: \* From state **v<sub>1</sub>** (**P<sub>1</sub>** holds **L<sub>1</sub>**), **P<sub>2</sub>** attempts to execute `lock(L2)`. This succeeds, moving the system to state **v<sub>{12}</sub>** (**P<sub>1</sub>** holds **L<sub>1</sub>**, **P<sub>2</sub>** holds **L<sub>2</sub>**). \* From state **v<sub>2</sub>** (**P<sub>2</sub>** holds **L<sub>2</sub>**), **P<sub>1</sub>** attempts to execute `lock(L1)`. This succeeds, moving the system to state **v<sub>{21}</sub>** (**P<sub>1</sub>** holds **L<sub>1</sub>**, **P<sub>2</sub>** holds **L<sub>2</sub>**).

So far, the states **v<sub>{12}</sub>** and **v<sub>{21}</sub>** are identical, let’s call it **v<sub>{final}</sub>**. The issue arises when we consider the *other* possible executions from **v<sub>1</sub>** and **v<sub>2</sub>**.

- **From v<sub>1</sub>**: **P<sub>1</sub>** holds **L<sub>1</sub>**. **P<sub>1</sub>** attempts to acquire **L<sub>2</sub>**. Simultaneously, **P<sub>2</sub>** might attempt to acquire **L<sub>1</sub>**. If **P<sub>1</sub>**’s request is blocked waiting for **L<sub>2</sub>** (held by **P<sub>2</sub>**) and **P<sub>2</sub>**’s request is blocked waiting for **L<sub>1</sub>** (held by **P<sub>1</sub>**), the system has entered a deadlocked state.
- Let’s model the “bad” interleaving:
  1. **v<sub>0</sub>** → **v<sub>1</sub>** (**P<sub>1</sub>** acquires **L<sub>1</sub>**).
  2. **v<sub>0</sub>** → **v<sub>2</sub>** (**P<sub>2</sub>** acquires **L<sub>2</sub>**).
  3. From **v<sub>1</sub>**, **P<sub>2</sub>** cannot proceed (`lock(L1)` fails). **P<sub>1</sub>** attempts `lock(L2)`.
  4. From **v<sub>2</sub>**, **P<sub>1</sub>** cannot proceed (`lock(L2)` fails). **P<sub>2</sub>** attempts `lock(L1)`.

If the interleaving is **P<sub>1</sub>**:`lock(L1)` followed by **P<sub>2</sub>**:`lock(L2)`, the system

reaches a state where  $P_1$  is blocked on  $L_2$  and  $P_2$  is blocked on  $L_1$ . This is a terminal state  $v_{\text{deadlock}}$ . No further directed edges emanate from  $v_{\text{deadlock}}$ .

Geometrically, the state space which should have been a simple, filled square now has a “hole” or a “missing corner.” The paths  $v_0 \rightarrow v_1$  and  $v_0 \rightarrow v_2$  exist. But the composite path corresponding to the deadlock interleaving leads to a vertex  $v_{\text{deadlock}}$  from which progress is impossible. This vertex lies on the **boundary** of the space of reachable, non-terminal states. A deadlock is a point, or a region, in the protocol complex  $X$  that is a sink in the directed graph of states—a geometric cul-de-sac.

**Directed Holes: The Geometry of Irreconcilable Views** Deadlock is a stark failure, a complete cessation of progress. However, distributed systems are plagued by more subtle failures, particularly those related to achieving consensus. A system may remain “live” (processes continue to execute steps) but fail to achieve its collective goal. This is the situation modeled by a *directed hole*.

A directed hole is a more sophisticated structure than a simple boundary. It represents a situation where multiple execution paths diverge and can never be reconciled into a single, consistent future state, *even though locally it appears they might*. It is a global inconsistency arising from locally consistent information.

**Definition (Intuitive):** A directed  $n$ -hole in a protocol complex  $X$  is an arrangement of  $(n-1)$ -dimensional execution cells that form a closed, directed “surface” (analogous to a sphere  $S^{n-1}$ ), but which do not form the boundary of any  $n$ -dimensional cell in  $X$ . The absence of this “filling”  $n$ -cell signifies that the state of agreement it would represent is forbidden by the protocol.

### The Canonical Example: A 2-Hole and Consensus Failure

Consider a task of  **$k$ -set agreement** among  $n$  processes ( $k < n$ ). Each process  $P_i$  starts with an initial value  $x_i$ . After communication, each process must decide on a value, and the set of all decided values must have a cardinality of at most  $k$ .

Let’s analyze a potential failure for  $k=1$  (standard consensus) with  $n=3$  processes  $P_1, P_2, P_3$ . \* Imagine an execution  $E_1$  where  $P_1$  fails early.  $P_2$  and  $P_3$  communicate and decide on value  $v_A$ . \* Imagine an execution  $E_2$  where  $P_2$  fails early.  $P_1$  and  $P_3$  communicate and decide on value  $v_B$ . \* Imagine an execution  $E_3$  where  $P_3$  fails early.  $P_1$  and  $P_2$  communicate and decide on value  $v_C$ .

Assume  $v_A, v_B$ , and  $v_C$  are all distinct. Geometrically: \* There is a directed path (a 1-cell) from the initial state to a state where the decision is  $\{v_A\}$ . \* There is another path to a state where the decision is  $\{v_B\}$ . \* There is a third path to a state where the decision is  $\{v_C\}$ .

Now, consider an execution where no process fails, but messages are delayed such that: \* P<sub>1</sub> “thinks” P<sub>2</sub> might have failed, and is leaning towards deciding  $v_C$  (as in E<sub>3</sub>). \* P<sub>2</sub> “thinks” P<sub>3</sub> might have failed, and is leaning towards deciding  $v_A$  (as in E<sub>1</sub>). \* P<sub>3</sub> “thinks” P<sub>1</sub> might have failed, and is leaning towards deciding  $v_B$  (as in E<sub>2</sub>).

This situation forms the boundary of our directed hole. We have three divergent “possible futures.” The system state is on a “rim” defined by these three possibilities. The three executions (E<sub>1</sub>, E<sub>2</sub>, E<sub>3</sub>) form a directed, triangular boundary. Can this boundary be “filled”?

A “filling” would correspond to a 2-dimensional execution cell, representing a set of interleavings where the three processes communicate further and reconcile their views to a single decision value (since  $k=1$ ). However, the famous **FLP Impossibility Result** states that no such deterministic algorithm can exist in the presence of even one potential crash failure.

Therefore, for any such protocol, its execution complex  $X$  *must* contain this configuration: a directed, triangular boundary of states (where processes have conflicting “views” of the system) that is not filled in. This empty triangle is a **directed 2-hole**. It is the geometric manifestation of the FLP impossibility theorem. The hole represents the fundamental ambiguity of an asynchronous system: the inability to distinguish a crashed process from a very slow one. The processes are trapped on the boundary of this hole, able to communicate locally but unable to bridge the global divide to reach the (missing) consensus state.

### Higher-Dimensional Holes and Higher-Order Failures

The concept generalizes. \* A **directed 1-hole** is a cycle of directed paths  $v_0 \rightarrow v_1 \rightarrow \dots \rightarrow v_n \rightarrow v_0$  that does not bound a 2-cell. Computationally, this can represent a form of livelock, where the system cycles through a set of states without making forward progress towards its goal.

- A **directed 3-hole** would be a hollow tetrahedron, representing a failure of four-way agreement, and so on. These higher-order failures often correspond to obstructions in composing protocols or reasoning about systems with multiple, interacting consensus groups.

**From Geometric Holes to Algebraic Obstructions** The discovery of a directed hole in the protocol complex is a powerful diagnostic tool. It is a precise, geometric witness to a protocol’s flaw. This geometric intuition provides the foundation for the algebraic methods we will introduce next. The central thesis of this work is that:

**A directed hole in the geometry of  $X$  corresponds to a nontrivial class in the nonabelian cohomology of  $X$ .**

The boundary of a directed hole can be described by a **cocycle**: an assignment of data (state transformations, group elements) to the cells of the boundary that

is locally consistent but cannot be derived from a single, global piece of data. For instance, the “rim” of the 2-hole in the consensus example can be annotated with the different decision values  $\{v_A, v_B, v_C\}$ . This assignment is a 1-cocycle. The fact that it cannot be reconciled to a single value means this cocycle is not a coboundary. Its corresponding class in  $H^1(X, A)$  (where  $A$  is the group of state differences) is nontrivial.

In conclusion, the directed topology of asynchronous computations provides a rich semantic landscape. Concurrency gives the space volume, deadlocks form its terminal edges, and the fundamental impossibilities of distributed consensus carve out directed holes. These holes are the geometric loci of disagreement. By studying their structure, we can classify failure modes, and by translating them into the language of algebra, we can develop a powerful calculus for detecting and reasoning about them.

## Chapter 2.4: Directed Homology: Quantifying Causal Anomalies

In the preceding chapters, we established that the inherent causal structure of asynchronous computations—the directedness of time and information flow—demands a move from standard topology to directed topology. We introduced the execution space  $X$  as a directed space (often modeled as a simplicial or cubical complex with a preferred direction on its cells) and the fundamental category  $\text{Cat}_1(X)$  as the repository of all valid causal pathways. These tools provide a rich, qualitative language for describing computational phenomena: concurrent execution paths appear as “diamonds,” deadlocks as “unfillable” directed cycles, and irreversible state changes as oriented edges.

However, a complete description of all causal paths via  $\text{Cat}_1(X)$  is often intractable and overly detailed for comparative analysis. To move from qualitative description to quantitative measurement, we require algebraic invariants that can distill the essential causal features of an execution space into a computable and comparable form. Just as classical homology theory counts holes and connected components in a space, we need an analogous tool that respects the arrow of time. This tool is **directed homology**. Its purpose is not merely to detect holes, but to detect and quantify *causal anomalies*: non-confluent execution paths, resource deadlocks, and other path-dependent inconsistencies that are invisible to classical methods. This chapter develops the theory of directed homology and interprets its groups as direct measures of the structural origins of consensus failures.

### The Inadequacy of Standard Homology

Before constructing a new theory, it is crucial to understand why the standard, powerful machinery of singular or simplicial homology is insufficient for our purposes. Classical homology is fundamentally insensitive to direction. It is built upon the notion of a chain, a formal sum of simplices  $c = \sum a_i \sigma_i$ , and a boundary operator  $\partial$  that computes the oriented boundary of a simplex. The

key word is *oriented*, but this orientation is purely combinatorial—an arbitrary ordering of vertices ( $v_0, v_1, \dots, v_n$ )—used to ensure that the boundary of a boundary is zero ( $\partial^2 = 0$ ).

Consider two fundamental scenarios in an asynchronous system with processes P1 and P2: 1. **Safe Concurrency:** P1 executes action  $a$ , and P2 executes action  $b$ . In an asynchronous setting, the system can transition from an initial state  $S_0$  to a final state  $S_f$  via two paths:  $S_0 \rightarrow a \rightarrow S_a \rightarrow b \rightarrow S_f$  and  $S_0 \rightarrow b \rightarrow S_b \rightarrow a \rightarrow S_f$ . Topologically, this forms a commutative square or “diamond.” In standard homology, this square is a 2-chain whose boundary is zero. It is a filled, trivial structure. 2. **Deadlock:** P1 holds resource R1 and requests R2, while P2 holds R2 and requests R1. This creates a directed cycle of dependencies:  $P1\_waits\_for\_P2 \rightarrow P2\_waits\_for\_P1 \rightarrow \dots$ . This is a 1-cycle in the state space.

The problem is that classical homology cannot reliably distinguish these scenarios. If the concurrent diamond is modeled as the boundary of a 2-simplex, its first homology  $H_1$  is trivial. If the deadlock cycle is also modeled as the boundary of some higher-dimensional simplex representing a “protocol recovery” state, its  $H_1$  also becomes trivial. Standard homology is designed to detect holes that exist regardless of the direction of traversal. It equates a path  $A \rightarrow B$  with its inverse  $B \rightarrow A$  at the chain level, effectively erasing the very essence of computation: causality and irreversibility. We need a homology theory where chains are intrinsically directed and the boundary operator respects this flow.

## Constructing Directed Homology

The core idea behind directed homology is to restrict the notion of chains and redefine the boundary operator to be consistent with the causal structure of the execution space  $X$ . While several formalisms exist (e.g., path homology, natural homology), they share a common spirit.

Let  $X$  be a directed space, modeled for simplicity as a simplicial complex where each simplex  $\sigma = (v_0, \dots, v_n)$  has an intrinsic partial order on its vertices, representing a valid sequence of events.

**1. Directed Chains:** A *directed  $n$ -chain* is a formal linear combination of *globally consistent* directed  $n$ -simplices. A directed  $n$ -simplex  $\sigma = (v_0, \dots, v_n)$  is not just a set of vertices but represents a specific causal pathway, often written  $v_0 \rightarrow v_1 \rightarrow \dots \rightarrow v_n$ . The chain group,  $C_n^{\text{dir}}(X)$ , is the free abelian group generated by these directed  $n$ -simplices. This is in contrast to standard homology where any permutation of vertices defines a simplex, related by a sign. Here,  $(v_0, v_2, v_1)$  might not be a valid simplex at all if it violates the causal order.

**2. The Directed Boundary Operator  $\partial^{\text{dir}}$ :** This is the most critical modification. The boundary of a directed path should be its directed extremities: its final state minus its initial state. This principle is extended to higher

dimensions. For a directed 1-simplex (a path)  $= (v_0 \rightarrow v_1)$ , its boundary is:  $\partial_1^{\text{dir}}() = v_1 - v_0$

For a directed 2-simplex  $= (v_0 \rightarrow v_1 \rightarrow v_2)$ , which can be viewed as a composition of paths, its boundary is composed of its constituent 1-paths, respecting their role as “inputs” and “outputs”. A common definition is:  $\partial_2^{\text{dir}}() = (v_0 \rightarrow v_2) - (v_1 \rightarrow v_2) - (v_0 \rightarrow v_1)$  *Note the signs.* This is not the alternating sum  $(v_1, v_2) - (v_0, v_2) + (v_0, v_1)$  of standard homology. This definition expresses that the “shortcut” path  $(v_0 \rightarrow v_2)$  is equivalent to the composite path  $(v_0 \rightarrow v_1)$  followed by  $(v_1 \rightarrow v_2)$ . The operator is constructed such that  $\partial_1^{\text{dir}} \partial_2^{\text{dir}} = 0$  holds, allowing for the definition of homology groups.

**3. Directed Homology Groups:** With the directed chain complex  $(C_*^{\text{dir}}, \partial_*^{\text{dir}})$  in place, the directed homology groups are defined in the usual way:  $H_n^{\text{dir}}(X) = \text{Ker}(\partial_n^{\text{dir}}) / \text{Im}(\partial_{n+1}^{\text{dir}})$  An element of  $H_n^{\text{dir}}(X)$  is an equivalence class of directed  $n$ -cycles (chains with no boundary) where two cycles are equivalent if they form the boundary of a directed  $(n+1)$ -chain.

### Geometric Interpretation of Directed Homology Groups

The power of these groups lies in their direct computational interpretation. They quantify specific types of causal structures and anomalies.

- **$H_0^{\text{dir}}(X)$ : Causal Sinks and Sources.** In standard homology,  $H_0(X)$  counts the number of path-connected components. In directed homology,  $H_0^{\text{dir}}(X)$  is more subtle. Its rank is related to the number of terminal components—subsets of the state space from which no further progress is possible—and initial components. It distinguishes between states that can only be started from and states that can only be reached, providing a global measure of the system’s entry and exit points.
- **$H_1^{\text{dir}}(X)$ : Quantifying Non-Confluence and Deadlock.** This is the most informative group for analyzing consensus protocols. A non-trivial class  $[z] \in H_1^{\text{dir}}(X)$  corresponds to a directed 1-cycle  $z$  that is not the boundary of any directed 2-chain. This represents a “causal hole” in the execution space and manifests in two primary ways:

1. **Non-confluence:** Consider two distinct causal paths,  $\alpha_1$  and  $\alpha_2$ , from a state  $S_i$  to a state  $S_f$ . The chain  $z = \alpha_1 - \alpha_2$  is a cycle, as  $\partial^{\text{dir}}(z) = (S_f - S_i) - (S_f - S_i) = 0$ . This cycle  $z$  represents a non-trivial homology class if there is no directed 2-chain (i.e., no “filling” execution) such that  $\partial^{\text{dir}}() = z$ . This is precisely the geometric picture of path-dependence: two different histories lead from the same past to the same future, but the difference between them is topologically significant and cannot be reconciled. The system state depends on the path taken. The rank of

$H_1^{\text{dir}}(X)$  provides a quantitative measure of the degree of such path-dependent ambiguity in the system.

2. **Deadlock:** A cycle of resource dependencies,  $P1\_waits \rightarrow P2\_waits \rightarrow \dots \rightarrow P1\_waits$ , forms a directed 1-cycle  $z$ . If the protocol provides no mechanism to break this deadlock, there is no higher-dimensional execution whose boundary is  $z$ . Therefore, the deadlock cycle persists as a non-trivial element in  $H_1^{\text{dir}}(X)$ . The presence of such classes is a definitive signature of deadlock vulnerabilities.

- **$H_n^{\text{dir}}(X)$  for  $n > 1$ : Higher-Order Causal Anomalies.** Higher directed homology groups capture more complex causal structures. A non-trivial class in  $H_2^{\text{dir}}(X)$  can be interpreted as an “ambiguity in how ambiguities are resolved.” For instance, imagine two different non-confluent diamonds (1-cycles) starting and ending at the same states. If the “surface” between these two diamonds cannot be filled by a consistent 3-dimensional execution, it signifies a higher-order obstruction. These are subtler anomalies in the protocol’s logic, relating to how different failure scenarios are related to each other.

### Duality with Nonabelian Cohomology

Directed homology, as constructed with integer coefficients, excels at identifying the *existence and number* of causal anomalies. It provides the geometric “scaffolding” of the problem. Nonabelian cohomology, the central theme of this work, provides the *semantics* of these anomalies. The two are dual perspectives on the same problem.

- **Homology finds the path:** A non-trivial cycle  $[z] \in H_1^{\text{dir}}(X)$  identifies a loop  $z$  in the execution space, such as a non-confluent pair of paths. This tells us *where* a problem might occur.
- **Cohomology measures the effect:** A nonabelian cocycle, as we have seen, assigns an element from a group  $A$  (representing state transformations) to each path segment. The holonomy of the cocycle around the cycle  $z$  is the total transformation accumulated by traversing it.  $\text{Hol}(z) \in A$ . This tells us *what* the consequence of the anomaly is.

If  $\text{Hol}(z)$  is the identity element of  $A$ , then the path-dependence identified by homology is benign; the final state is unaffected. If  $\text{Hol}(z)$  is non-identity, the system is in a state of disagreement. The non-triviality of  $[z] \in H_1^{\text{dir}}(X)$  is a necessary condition for a non-trivial holonomy to manifest.

Therefore, the complete picture emerges from the interplay: **> Directed homology detects the geometric loci of potential disagreement (cycles), while directed cohomology with nonabelian coefficients quantifies the magnitude and nature of that disagreement (holonomy).**

A non-trivial class in  $H_1^{\text{dir}}(X)$  is a “disagreement waiting to happen.” A non-trivial class in  $H^1(X, A)$  is the realization of that disagreement, where  $A$  is the group of state-altering, non-commutative operations. Similarly, the 2-dimensional voids detected by  $H_2^{\text{dir}}(X)$  are the geometric structures upon which the non-trivial obstructions classified by  $H^2(X, A)$ —the gerbes representing protocol inconsistencies—are supported.

### Application: Probing the Ambiguity of Asynchronous Protocols

We can apply this machinery to analyze real-world protocols. Consider the  $k$ -set consensus problem, where  $n$  processes must decide on at most  $k$  different values. The FLP impossibility result and its generalizations show that for  $k$  less than the number of potential crash faults  $f+1$ , consensus is impossible in an asynchronous system.

We can model the execution space  $X$  of a  $k$ -set consensus protocol. The vertices are global states, and directed edges represent process steps (sending/receiving messages, changing local state). - Schedules where messages are delivered in different orders create different paths in  $X$ . - If two different schedules  $_1$  and  $_2$  lead from a common initial state to states where different sets of  $k$  values have been decided upon, the paths are non-confluent. - The cycle  $z = _1 - _2$  might be non-trivial in  $H_1^{\text{dir}}(X)$ .

The rank of  $H_1^{\text{dir}}(X)$  can thus be interpreted as a quantitative measure of the protocol’s inherent **schedule ambiguity**. A protocol with a higher-rank first directed homology group is one with more fundamental ways for executions to diverge due to asynchrony. A “robust” protocol would be one designed to introduce higher-dimensional “filler” executions that make these cycles trivial, ensuring that all paths eventually converge in a consistent manner. By computing these invariants, we can move beyond testing specific scenarios and make formal, quantitative statements about the intrinsic topological fragility of a given distributed algorithm.

## Chapter 2.5: The Topology of Impossibility: Analyzing $k$ -Set Consensus Protocols

### The Topology of Impossibility: Analyzing $k$ -Set Consensus Protocols

The preceding chapters have established the framework of directed topology as the natural language for describing the causal structure of asynchronous computations. By moving from undirected simplicial complexes to directed spaces, and from the fundamental group  $\pi_1(X)$  to the fundamental category  $\text{Cat}_1(X)$ , we have gained a finer lens through which to view computational phenomena like concurrency and deadlock. We now apply this powerful apparatus to one of the central problems in fault-tolerant distributed computing: *k-set consensus*. This problem serves as a crucible for theories of distributed computability, forming a precise hierarchy of impossibility that perfectly illustrates the interplay between fault levels, asynchrony, and topological obstruction.



The classic impossibility results for  $k$ -set consensus are among the foundational theorems of the field, elegantly proven using standard algebraic topology. In this chapter, we will revisit these results through the lens of directed topology and non-abelian cohomology. We will demonstrate that this new perspective does not merely re-prove a known theorem but enriches its meaning, recasting the static, topological “hole” of the classical proof as a dynamic, causal “vortex”—an inescapable ambiguity in the flow of information, quantified by cohomological invariants.

**The  $k$ -Set Consensus Problem** The standard consensus problem requires all correct processes in a distributed system to agree on a single value, chosen from a set of initial values proposed by the processes. The FLP impossibility result famously proved this problem unsolvable in an asynchronous system with even a single crash failure. The  $k$ -set consensus problem is a strategic relaxation of this requirement.

**Definition ( $k$ -Set Consensus):** In a system of  $n$  processes, each proposing an initial value, a protocol solves  $k$ -set consensus if: 1. **Termination:** Every correct process eventually decides on a value and terminates. 2. **Validity:** Any decided value must have been one of the initial values proposed by some process. 3. **Agreement:** The set of all values decided by correct processes has a size of at most  $k$ .

Here,  $k$  is an integer,  $1 \leq k \leq n$ . The case  $k=1$  corresponds to the standard consensus problem. The case  $k=n$  is trivial, as processes can simply decide their own initial value. The problem’s interest lies in the intermediate values of  $k$ . It establishes a hierarchy of agreement problems, where lower values of  $k$  represent “harder” problems. The seminal result by Herlihy and Shavit, building on work by Chaudhuri, established a tight bound: in an asynchronous system of  $n$  processes that is resilient to  $t$  crash failures,  $k$ -set consensus is solvable if and only if  $k > t$ . Consequently,  $t$ -set consensus is impossible.

**The Classical Topological Proof: Connectivity vs. Obstruction** The classical proof of the  $k \leq t$  impossibility is a masterpiece of applied topology. It proceeds by modeling the computational task itself as a relationship between topological spaces.

1. **The Input Complex ( $I$ ):** This space represents all possible initial configurations. For a system of  $n=t+1$  processes where each proposes a unique value from a set  $V$ , the input complex can be seen as the set of all  $t+1$ -subsets of  $V$ . The nerve of this configuration is a  $t$ -dimensional simplex, and its boundary, representing states where at least one initial value is not present (e.g., because a process has not yet run), is topologically equivalent to a  $t$ -sphere,  $S^t$ . More generally, for  $n$  processes and  $t$  failures, the protocol must be able to handle any input configuration involving  $n-t$  running processes, giving the protocol complex a topological complexity related to  $t$ -connectivity.

2. **The Output Complex ( $O_k$ ):** This space represents all valid final configurations. For  $k$ -set consensus among  $n$  processes, a vertex in this complex is a pair  $(p_i, v_j)$  where process  $p_i$  decides value  $v_j$ . A simplex  $\sigma = \{(p_{i_0}, v_{j_0}), \dots, (p_{i_m}, v_{j_m})\}$  is in  $O_k$  if the set of decided values  $\{v_{j_0}, \dots, v_{j_m}\}$  has size at most  $k$ . A key topological property of this output complex is that it is  $(k-1)$ -connected. This means that any continuous map from an  $m$ -sphere  $S^m$  into  $O_k$  can be extended to a continuous map of the  $(m+1)$ -disk  $D^{m+1}$  for all  $m < k$ . Intuitively, it has no “holes” in dimensions up to  $k-1$ .
3. **The Protocol Complex ( $P$ ):** This space represents all possible executions of a given protocol. Its vertices are local states of processes, and its simplices represent collections of processes that are mutually consistent at some point in an execution. A wait-free protocol tolerant of  $t$  failures must produce a protocol complex  $P$  that is  $t$ -connected.

**The Topological Argument:** A protocol solves the task if there exists a continuous, color-preserving (respecting process IDs) decision map  $\delta : P \rightarrow O_k$ . This map must be compatible with the initial values specified on the boundary of  $P$ . The impossibility for  $k \leq t$  arises from a topological contradiction.

Consider a system with  $n=t+1$  processes. The protocol complex  $P$  must be  $t$ -connected. The input configuration space contains a topological  $t$ -sphere,  $S^t$ . The decision map  $\delta$  must map this  $S^t$  (embedded in  $P$ ) into the output complex  $O_k$ . However, since  $k \leq t$ , we have  $k-1 < t$ . The output complex  $O_k$  is  $(k-1)$ -connected, meaning it has no holes of dimension  $t$ . But  $\delta$  restricted to the input sphere would be a map from  $S^t$  to a space with no  $t$ -dimensional holes. Such a map cannot exist without being contractible to a point, which would violate the Validity condition (forcing all processes to decide the same value, regardless of input). The  $t$ -dimensional hole in the input/protocol space has nowhere to go in the  $(k-1)$ -connected output space.

**A Directed Reinterpretation: Causal Vortices** The classical model, for all its power, is fundamentally static. It identifies the existence of a hole but does not describe the computational dynamics that create and sustain it. Directed topology provides this dynamic narrative.

We replace the protocol complex  $P$  with its directed version, the protocol **d-space**  $dP$ , where paths are irreversible sequences of events. The fundamental invariant is no longer the homotopy group  $\pi_n(X)$  but the fundamental category  $\text{Cat}_1(dP)$ . The “directed holes” of the previous chapter—diverging causal paths that cannot reconverge—are the geometric manifestation of computational ambiguity.

Let’s re-examine the  $k \leq t$  impossibility. Consider  $n = t+1$  processes,  $\{p_0, \dots, p_t\}$ , each starting with a unique value  $\{v_0, \dots, v_t\}$ . We are trying to solve  $t$ -set consensus. \* **Divergence:** In an asynchronous execution, process  $p_i$  might run to completion before any other process even starts. In this “solo”

execution path,  $p_i$  sees only its own value  $v_i$  and, by the validity property, must decide  $v_i$ . \* **Concurrency:** Now consider an execution where  $p_i$  and  $p_j$  run concurrently. Other processes may observe their actions in different orders, leading to different future states. For example, process  $p_k$  might see  $p_i$ 's proposal before  $p_j$ 's, while  $p_l$  sees  $p_j$ 's before  $p_i$ 's. These scenarios trace out distinct, non-interchangeable directed paths in  $dP$ . \* **The Causal Vortex:** The collection of all possible interleavings of  $t+1$  processes generates a complex web of directed paths. The solo executions  $_0, \dots, _t$  represent  $t+1$  divergent futures, each leading to a different decision value. The problem of  $t$ -set consensus is to ensure that no matter what execution path is taken, the final set of decisions across all processes has size at most  $t$ . \* **Directed Obstruction:** The impossibility can now be stated in terms of directed maps. The protocol  $d\text{-space } dP$  contains a substructure corresponding to the initial ambiguity—a “di-sphere” of sorts, where  $t+1$  causally distinct execution families diverge. Any directed continuous decision map  $dP \rightarrow dO_t$  must map these divergent paths into the directed output space for  $t$ -set consensus. However, the structure of  $dO_t$  is such that it accommodates convergence to *at most*  $t$  distinct outcomes. It is “causally  $(t-1)$ -connected”. The attempt to map the  $(t+1)$ -fold ambiguity of  $dP$  into the  $t$ -outcome space of  $dO_t$  fails. There is no way to “tame” the  $t+1$  divergent causal paths into  $t$  final states without creating a discontinuity in the directed sense—a computational jump that is forbidden in an asynchronous model.

The static hole of the classical proof is thus re-imagined as a **causal vortex**: a region in the state space where multiple, mutually incompatible futures ( $t+1$  of them) are simultaneously possible due to asynchrony, and from which the system cannot escape into a configuration with fewer than  $t+1$  outcomes without violating causality.

**Cohomological Framing: The Torsor of Indecision** We can elevate this geometric intuition to an algebraic one using the framework of non-abelian cohomology. Let the state of the system be described by a sheaf  $\mathcal{F}$  over the protocol  $d\text{-space } dP$ . The stalks of this sheaf,  $\mathcal{F}_x$ , represent the information available at an execution state  $x$ . The transitions along directed paths are given by a non-abelian group  $A$  of transformations.

For 1-set consensus ( $k=1$ ), as discussed in previous chapters, the impossibility is captured by the non-triviality of the first cohomology set,  $H^1(dP, \mathcal{A})$ , where  $\mathcal{A}$  is the sheaf of groups of permutations on the decision values. A non-trivial class in  $H^1(dP, \mathcal{A})$  corresponds to an  $\mathcal{A}$ -torsor—a “twisted” version of the state space where there is no globally consistent choice of a “base” decision. The holonomy around a directed loop, representing an information cycle like  $p_i$  sees  $x$ ,  $p_j$  sees  $y$ ,  $p_i$  learns  $p_j$ 's state,  $p_j$  learns  $p_i$ 's state—results in a non-identity element of  $A$ , a tangible measure of accumulated disagreement.

For  $k$ -set consensus, the obstruction is more subtle and lies in higher cohomology. It is not that agreement is completely impossible, but that the possible scope

of agreement is limited. The obstruction is not a single torsor, but a higher-dimensional algebraic structure.

Let us posit the existence of a **characteristic class of  $k$ -agreement**,  $c_k \in H^k(dO_k, \mathcal{G})$ , for some appropriate coefficient sheaf of (possibly non-abelian) groups or higher-groupoids  $\mathcal{G}$ . This class  $c_k$  is a cohomological object that fundamentally defines the  $k$ -agreement property. It is trivial for any subspace corresponding to agreement on fewer than  $k$  values, but non-trivial over regions requiring exactly  $k$  distinct values.

The impossibility of  $t$ -set consensus for  $n=t+1$  processes can then be framed as follows: 1. The protocol  $d$ -space  $dP$  has a rich directed topology. Its  $t$ -dimensional causal structure, stemming from the  $t+1$  independent processes, is non-trivial. This can be expressed as the non-vanishing of some directed homology group,  $H_t^{\text{dir}}(dP)$ . 2. The decision map  $dP \rightarrow dO_t$  induces a map in cohomology,  $\hat{*}: H^t(dO_t, \mathcal{G}) \rightarrow H^t(dP, \mathcal{G})$ . 3. The characteristic class of  $t$ -agreement,  $c_t \in H^t(dO_t, \mathcal{G})$ , is by definition non-trivial. 4. For a valid protocol to exist, the pullback  $\hat{*}(c_t)$  must be compatible with the system's initial state. However, the initial configuration, which contains the  $t$ -sphere of ambiguity, forces  $\hat{*}(c_t)$  to pair non-trivially with the  $t$ -dimensional directed homology class of the protocol space. 5. This leads to an obstruction:  $\hat{*}(c_t), [S_t^{\text{dir}}] \neq 0$ . The  $t$ -dimensional causal vortex in the protocol space  $dP$  “activates” the  $t$ -dimensional cohomological obstruction  $c_t$  in the output space  $dO_t$ . The protocol is impossible because it is being asked to carry a non-trivial  $t$ -cocycle, but its  $t$ -connectivity properties, required for fault tolerance, mean it cannot support such a cocycle in a way that is consistent across all executions.

The non-abelian nature of the group  $\mathcal{G}$  is crucial here. It reflects the path-dependence of information. The order in which a process learns of the initial values  $\{v_0, \dots, v_t\}$  matters. One ordering leads down one causal path, another ordering leads down another. The discrepancy between these paths is an element of  $\mathcal{A}$ . The  $t$ -set impossibility arises because there are  $t+1$  fundamental “discrepancy generators” (the  $t+1$  initial values) and the algebraic structure of the output space  $dO_t$  only provides the means to resolve  $t$  of them. The  $(t+1)$ -th source of disagreement creates a global, non-trivial cocycle that cannot be made trivial by any protocol.

In conclusion, the analysis of  $k$ -set consensus through the prism of directed topology and non-abelian cohomology transforms a static impossibility proof into a dynamic story of causality, information flow, and irreducible ambiguity. The topological hole becomes a causal vortex, a structure in the state space sustained by asynchrony. The obstruction to consensus becomes an algebraic object—a non-trivial cocycle—that measures the degree to which local, path-dependent views of the system fail to cohere into a global state that satisfies the constraints of the problem. This powerful synthesis of ideas not only deepens our understanding of this specific problem but also provides a general and extensible language for analyzing the fundamental limits of distributed computation.

## Part 3: Nonabelian State Symmetries and Path-Dependent Updates

### Chapter 3.1: The Algebra of State Updates: Modeling Operations with Non-Abelian Groups

#### The Algebra of State Updates: Modeling Operations with Non-Abelian Groups

In our prior exploration of the topological nature of distributed executions, we have largely treated the system's state abstractly. We identified how the topology of a protocol complex  $X$ , particularly its fundamental group  $\pi_1(X)$ , can obstruct the existence of a continuous global choice, thereby dooming consensus. However, this analysis remains incomplete without a precise characterization of the *values* over which processes must agree and the *operations* that modify these values. The nature of the state space and its transformations is not merely incidental; it is a fundamental determinant of the system's susceptibility to failure.

Many simple models of distributed state, such as replicated counters or registers, implicitly assume that the operations performed by processes are commutative. Incrementing a counter by 5 and then by 3 yields the same result as incrementing by 3 and then by 5. This property, captured by an Abelian group structure  $(G, +)$ , simplifies reconciliation enormously. If the set of applied operations is known, their execution order is irrelevant to the final state. However, a vast class of critical computational problems involves operations that do not commute. This chapter introduces the algebraic framework necessary to model such systems, demonstrating that when the algebra of state updates is non-Abelian, path-dependence emerges as a primary and irreducible source of disagreement.

**From State Machines to Group Actions** Let us formalize the components of our system. We have a set of possible global states,  $S$ , and a set of elementary operations,  $O = \{o_1, o_2, \dots\}$ , that can be applied to these states. A distributed system can be conceptualized as a collection of replicas, each maintaining a local copy of the state  $s \in S$ , which they update in response to a stream of operations.

To capture the algebraic nature of these updates, we model the set of reversible operations as a group  $A$ , which we will call the **group of state transformations**. Each element  $a \in A$  corresponds to a bijective transformation of the state space,  $a: S \rightarrow S$ . The group law is function composition: applying operation  $b$  followed by operation  $a$  is equivalent to the single composite operation  $a \circ b$ . The group identity element,  $e$ , represents the null operation, which leaves the state unchanged. The inverse  $a^{-1}$  represents the undoing of operation  $a$ .

This group structure  $A$  is the algebraic fingerprint of the system's dynamics. Its properties dictate the fundamental rules of state evolution. The crucial distinction is whether  $A$  is Abelian or non-Abelian.

- **Abelian (Commutative) Updates:** If for all  $a, b \in A$ ,  $ab = ba$ , the group is Abelian. The order of operations is irrelevant.
  - **Example:** A replicated bank account ledger. The state  $s$  is the balance. The operations are deposits and withdrawals, which can be modeled by the group  $A = (\mathbb{R}, +)$ . A deposit of  $d$  and a deposit of  $d$  result in the same final balance regardless of order:  $s + d + d = s + d + d$ .
- **Non-Abelian (Non-Commutative) Updates:** If there exist  $a, b \in A$  such that  $ab \neq ba$ , the group is non-Abelian. The order of operations is critical.
  - **Example 1: Configuration of a Rigid Body.** Let the state  $s \in S$  be the orientation of a satellite in space. The operations  $a, b \in A = SO(3)$  are rotations. A 90-degree rotation about the x-axis followed by a 90-degree rotation about the y-axis results in a different final orientation than performing the rotations in the reverse order.
  - **Example 2: Document Editing.** Let  $S$  be the set of all possible text documents. Consider two operations:  $a =$  “replace all instances of ‘color’ with ‘colour’” and  $b =$  “append the sentence: ‘The color is vibrant.’”. Applying  $a$  then  $b$  yields a document containing “colour” and the sentence “The colour is vibrant.”. Applying  $b$  then  $a$  yields a document containing “colour” and the sentence “The color is vibrant.”. The results are different. While not a perfect group structure (not all edits are easily reversible), this illustrates the core non-commutative principle.
  - **Example 3: Replicated Database with Constraints.** Consider a database state  $s$  and two transactions  $T$  and  $T'$ .  $T$  transfers funds from account  $X$  to  $Y$ .  $T'$  checks if  $Y$ ’s balance is above a threshold and, if so, grants a privilege. If  $T$  and  $T'$  are executed concurrently, their interleaving matters. If  $T$  runs first,  $T'$  might succeed; if  $T'$  runs first, it might fail. The state includes not just account balances but also privileges, and the transformation group is non-Abelian.

In a distributed setting, asynchrony guarantees that different replicas will observe and apply concurrent operations in different orders. When the group  $A$  is non-Abelian, this temporal ambiguity translates directly into state divergence.

**Path-Dependence and the Geometry of Disagreement** Let us connect this algebraic structure to the topological model of an execution. An execution path for a single replica  $i$  is a sequence of states and applied operations:  $s \rightarrow a \rightarrow s \rightarrow a \rightarrow s \rightarrow \dots \rightarrow a \rightarrow s$ , where  $s = a \cdot s$ . The final state is the result of applying the composite transformation  $h = a \cdot \dots \cdot a$  to the initial state  $s$ . We write this as  $s = h \cdot s$ .

Now, consider two replicas,  $P$  and  $P'$ , starting in a common initial state  $s$ . Due to network latency and scheduling uncertainties, they receive requests for two concurrent, non-commuting operations,  $a$  and  $b$  ( $ab \neq ba$ ). \*  $P$  observes the

sequence  $(a, b)$ . Its local history corresponds to a path  $p$  in the protocol complex  $X$ . Its final state is  $s = (b \ a) \ s$ .  $P$  observes the sequence  $(b, a)$ . Its local history corresponds to a different path  $p$  in  $X$ . Its final state is  $s = (a \ b) \ s$ .

Since  $ab \neq ba$ , it follows that  $s \neq s$ . The replicas have irrevocably diverged. This phenomenon is **path-dependence**: the final state of the system is a function not merely of the *set* of operations executed, but of the *path* taken through the execution space.

This simple “diamond” scenario, formed by two concurrent operations, is the archetypal geometric feature of non-commutative disagreement.

$$\begin{array}{c} \text{/-- } a \text{ --} \backslash \\ s \qquad \qquad s = (a \ b) \ s \\ \backslash \text{-- } b \text{ --} / \\ s = (b \ a) \ s \end{array}$$

The discrepancy between the two final states is not random. It is an algebraic quantity. To transform state  $s$  into  $s$ , one must apply the transformation  $g = (a \ b) \ (b \ a)^{-1}$ :  $g \ s = ((a \ b) \ (b \ a)^{-1}) \ ((b \ a) \ s) = (a \ b) \ s = s$ .

This transformation,  $g = a \ b \ a^{-1} \ b^{-1}$ , is the **commutator** of  $a$  and  $b$ , often denoted  $[a, b]$ . The non-triviality of the commutator ( $[a, b] \neq e$ ) is the algebraic measure of the disagreement generated by the concurrency of  $a$  and  $b$ .

**Holonomy: The Algebraic Residue of Execution Cycles** The concept of the commutator generalizes to arbitrary loops in the execution space  $X$ . A path in  $X$  represents a causal sequence of events (e.g., sending/receiving messages, applying operations). A loop in  $X$  represents a scenario where a process returns to a state of knowledge equivalent to a previous one, but potentially after a complex series of intermediate events have unfolded elsewhere in the system.

Let  $\gamma$  be a loop in the protocol complex  $X$  based at some initial vertex (system state). Traversing this loop corresponds to applying a sequence of operations  $a_1, a_2, \dots, a_n$ . The net effect of traversing this loop is the composite operation  $h_\gamma = a_1 \dots a_n \ a_1^{-1}$ . This element  $h_\gamma \in A$  is the **holonomy** of the path  $\gamma$ .

- If  $A$  is Abelian, the holonomy depends only on the homology class of the loop. For any contractible loop (a loop that can be continuously shrunk to a point), the holonomy is the identity  $e$ .
- If  $A$  is non-Abelian, the situation is far richer. Even for a contractible loop, the holonomy may be non-trivial. The holonomy of the “diamond” path  $p^{-1} \circ p$  (going up via  $b$  and  $a$ , and back down via  $a^{-1}$  and  $b^{-1}$ ) is precisely the commutator  $[a, b]$ .

The holonomy group, the set of all  $h_\gamma$  for loops  $\gamma$  based at a point, captures the

“twist” or “charge” of disagreement that can be accumulated by traversing cycles in the space of possible executions. If a protocol permits execution paths that form a loop with non-trivial holonomy, it means that a subset of processes can follow a sequence of operations that brings their local view of the protocol state back to where it started, yet the global data state has been irrecoverably altered by a transformation  $h \neq e$ . This is a fundamental obstruction to reaching a consistent state. Any attempt to “reconcile” states by simply comparing final values is doomed to fail, as the discrepancy is woven into the very fabric of the execution’s geometry and the state’s algebraic symmetries.

**A Formal Model of Disagreement** We can now state the problem of non-Abelian consensus with greater precision. Let the system be defined by: 1. A protocol complex  $X$ , representing the space of all possible executions. 2. A set of states  $S$ . 3. A non-Abelian group  $A$  of state transformations, acting on  $S$ . This is formally a group action  $\alpha : A \times S \rightarrow S$ .

Each process  $P$  traverses a path  $p$  in  $X$  starting from a common initial vertex corresponding to the initial state  $s \in S$ . The path  $p$  is determined by the sequence of messages  $P$  sends and receives. To each path  $p$ , we can associate a total transformation  $h \in A$  representing the ordered product of all operations  $P$  has applied. The local state at process  $P$  is  $s' = h \cdot s$ .

Consensus is the condition that  $s_i = s_j$  for all pairs of correct processes  $(i, j)$ . This implies  $h_j \cdot s = h_i \cdot s$ . This does not necessarily mean  $h_i = h_j$ . It means that the transformation  $g = h_i \cdot h_j^{-1}$  must be in the **stabilizer** of  $s$ , i.e.,  $g \cdot s = s$ . For many systems, the state space  $S$  is “homogeneous,” meaning the action is free (the stabilizer of any point is trivial,  $\{e\}$ ). In this common and important case, consensus  $s_i = s_j$  is equivalent to the condition  $h_i = h_j$ .

Disagreement, therefore, can be precisely quantified. For any two processes  $P_i$  and  $P_j$ , their disagreement is captured by the group element  $g = h_i \cdot h_j^{-1}$ . This element represents the “gauge transformation” needed to map  $P_i$ ’s state to  $P_j$ ’s state:  $s_j = g \cdot s_i$ . The system is in a state of disagreement if there exists any pair  $(i, j)$  for which  $g \neq e$ .

**Conclusion: Setting the Stage for Cohomology** This chapter has established a critical link: the non-commutativity of state-update operations is a fundamental algebraic source of path-dependent disagreement in distributed systems. We have moved from an abstract topological picture to a concrete algebraic model where states are elements of a set  $S$  and updates are elements of a non-Abelian group  $A$  acting on  $S$ .

The key insights are: 1. **Algebra of Updates:** Modeling state transformations as a group  $A$  allows us to precisely characterize the system’s potential for disagreement. If  $A$  is non-Abelian, order matters. 2. **Path-Dependence:** Asynchronous execution creates a multitude of paths through the protocol complex.



Non-commutativity turns this path ambiguity into concrete state divergence. 3. **Holonomy as Obstruction:** The holonomy  $h_\gamma \in A$  of a loop  $\gamma$  in the execution space measures the net state transformation accumulated. A non-trivial holonomy represents a topologically protected “scar” of disagreement. 4. **Local Discrepancies:** The disagreement between any two processes  $P$  and  $P'$  can be encapsulated in a single group element  $g = h_\gamma h_\gamma^{-1}$ .

These elements  $g$  are the crucial building blocks for the next stage of our analysis. They are not independent; they are related by the geometry of the protocol complex. For any three processes  $i, j, k$ , the transformations must satisfy the relation  $g_{ij} = g_{ik} g_{kj}$ . This is precisely the **1-cocycle condition**.

We have successfully translated the problem of consensus failure into the language of algebra and topology. The local, pairwise disagreements  $g$  are the raw data. The global question of whether these local disagreements can be “gauged away”—that is, whether a global state  $s$  can be defined such that all local states can be seen as transformations of  $s$ —is a question of non-Abelian cohomology. The next chapter will formalize this, showing how the set of all such discrepancies  $\{g\}$  defines a 1-cocycle, and how the possibility of consensus is equivalent to this cocycle being a trivial 1-coboundary. A non-trivial class in the first cohomology group  $H^1(X, A)$  will represent a fundamental, topologically enforced obstruction to agreement.

## Chapter 3.2: State Symmetries and the Action of the Update Groupoid

### State Symmetries and the Action of the Update Groupoid

In the previous chapter, we established the algebraic foundation for path-dependence by modeling system updates as a non-abelian group,  $A$ . This algebraic structure captures the crucial fact that the order of operations matters. However, this model remains static; it describes the *types* of operations available but does not yet integrate them with the *structure* of possible execution paths. The critical step is to combine the algebra of updates ( $A$ ) with the topology of computation ( $X$ ). This chapter formalizes this synthesis by introducing the **Update Groupoid**, a structure that encodes how transformations from  $A$  are applied along the causal paths of a distributed execution. Furthermore, we will explore the concept of **state symmetries**—transformations that leave the essential aspects of a state invariant—and demonstrate how their interplay with the groupoid action becomes a primary source of ambiguity and consensus failure.

**The Nature of State Symmetries** In any complex system, not all state configurations are meaningfully distinct. Certain transformations may alter the raw data of a state  $s \in S$  to a new state  $s'$  while preserving the properties relevant to the protocol’s logic. This notion of “indistinguishability” is formalized by the concept of symmetry.

**Definition 3.2.1 (State Space and Group Action):** Let  $S$  be the set of all possible global states of the system. The update group  $A$  acts on  $S$  from the left. For any operation  $a \in A$  and state  $s \in S$ , the application of the operation yields a new state  $a \cdot s \in S$ . This action must be compatible with the group structure:  $(a \cdot a') \cdot s = a \cdot (a' \cdot s)$  for all  $a, a' \in A$ , and  $e \cdot s = s$  where  $e$  is the identity of  $A$ .

The most direct form of symmetry is an operation that leaves a state completely unchanged.

**Definition 3.2.2 (Stabilizer Group):** For a given state  $s \in S$ , the **stabilizer** of  $s$  in  $A$ , denoted  $\text{Stab}_A(s)$ , is the subgroup of  $A$  consisting of all elements that fix  $s$ :  $\text{Stab}_A(s) = \{a \in A \mid a \cdot s = s\}$ . The elements of  $\text{Stab}_A(s)$  are symmetries of the specific state  $s$ . For example, in a system managing resource locks, re-acquiring a lock that a process already holds might be an identity operation on the state, making it an element of the stabilizer.

However, a more powerful and often more relevant concept of symmetry arises when we consider states to be equivalent without being identical. An observer or a protocol participant may not be able to distinguish between  $s$  and  $s'$  based on available information.

**Definition 3.2.3 (Symmetry Group of Indistinguishability):** Let  $\sim$  be an equivalence relation on the state space  $S$ , where  $s \sim s'$  signifies that states  $s$  and  $s'$  are indistinguishable from the perspective of the protocol's decision logic. The **symmetry group**  $A_{\sim}$  with respect to  $\sim$  is the largest subgroup of  $A$  whose elements map any state to an equivalent state:  $A_{\sim} = \{a \in A \mid s \in S, a \cdot s \sim s\}$ .

The existence of a non-trivial symmetry group  $A_{\sim}$  is a fundamental source of ambiguity. If an operation  $a \in A_{\sim}$  can be executed, the system transitions to a new state  $a \cdot s$  that is observationally identical to  $s$ . Processes that are unaware that  $a$  has occurred may proceed with incorrect assumptions about the system's true configuration. This is particularly pernicious in asynchronous systems where a sequence of operations may have a net effect  $a \in A_{\sim}$ , effectively hiding a complex history of state changes within a “phantom” operation.

**From Update Group to Update Groupoid** To model the application of updates along execution paths, we must elevate the update group  $A$  to a structure that lives on the protocol complex  $X$ . This structure is the **update groupoid**, which we will denote  $G(X, A)$ . While the fundamental groupoid  $\pi_1(X)$  captures the topology of paths, the update groupoid decorates these paths with the algebraic transformations they induce.

**Construction of the Update Groupoid  $G(X, A)$ :**

Let  $X$  be the protocol complex, a simplicial or cubical complex representing the space of possible executions. 1. **Objects:** The objects of  $G(X, A)$  are the vertices of  $X$ . Each vertex  $v \in V(X)$  represents a local or global state of

the system at a particular point in a partial execution. 2. **Morphisms:** A morphism in  $G(X, A)$  from vertex  $u$  to vertex  $v$  is a pair  $(\gamma, a)$ , where  $\gamma$  is a path in  $X$  from  $u$  to  $v$ , and  $a \in A$  is the cumulative transformation associated with that path. 3. **Composition:** Given two morphisms,  $(\gamma : u \rightarrow v, a)$  and  $(\delta : v \rightarrow w, b)$ , their composition is defined as  $(\gamma \delta, a \cdot b)$ , where  $\gamma \delta$  is the concatenated path from  $u$  to  $w$ . The group operation in  $A$  directly corresponds to the composition of transformations along the path. 4. **Identity:** For each object  $u$ , the identity morphism is  $(\text{id}_u, e)$ , where  $\text{id}_u$  is the constant path at  $u$  and  $e$  is the identity of  $A$ .

This construction is often simplified by considering a presentation of the groupoid. Let  $X$  be represented as a directed graph where edges correspond to atomic operations. For each directed edge  $e = (u, v)$ , we associate an element  $g_e \in A$  representing the update. A path  $\gamma = e_1 e_2 \dots e_k$  then corresponds to a total transformation  $a_\gamma = g_{e_k} \dots g_{e_1} g_{e_1}$ . The groupoid  $G(X, A)$  is then the path groupoid of this  $A$ -labeled graph.

The crucial insight is that  $G(X, A)$  is not simply  $(X) \times A$ . The algebraic and topological structures are intertwined. Two paths  $\gamma$  and  $\delta$  might be homotopic in  $X$  (i.e., they represent the same morphism in  $G(X)$ ), but their associated transformations  $a_\gamma$  and  $a_\delta$  may be different in  $A$ . This discrepancy, where topologically equivalent paths lead to algebraically distinct outcomes, is the essence of non-abelian holonomy and a primary focus of our analysis.

**The Action of the Groupoid on State Fibers** With the update groupoid defined, we can now formalize how the system's state evolves across the execution space. We conceptualize the state space not as a single monolithic set  $S$ , but as a collection of "state fibers" distributed over the protocol complex.

**Definition 3.2.4 (Bundle of States):** A **bundle of states** over  $X$  is a collection of sets  $\{S_v\}_{v \in V(X)}$ , where  $S_v$  is the set of possible system states as viewed from the local perspective of vertex  $v$ . The total state space is the disjoint union  $\text{mathcal{S}} = \bigsqcup_{v \in V(X)} S_v$ .

The update groupoid  $G(X, A)$  acts on this bundle of states.

**Definition 3.2.5 (Groupoid Action):** An action of  $G(X, A)$  on the bundle  $\text{mathcal{S}}$  is a map that takes a morphism  $(\gamma : u \rightarrow v, a)$  in  $G(X, A)$  and a state  $s_u \in S_u$  to a new state  $s_v \in S_v$ , denoted  $(\gamma, a) \cdot s_u$ . This action must satisfy: 1. **Compatibility:**  $(\gamma, a) \cdot s_u = a \cdot s_u$ . The action on the state is determined solely by the algebraic element  $a$  of the morphism. 2. **Functoriality:** If we compose morphisms  $(\gamma : u \rightarrow v, a)$  and  $(\delta : v \rightarrow w, b)$ , the action of the composite morphism is the composition of the actions:  $(\gamma \delta, b \cdot a) \cdot s_u = b \cdot (a \cdot s_u)$ .

This action provides a precise, path-dependent evolution rule for states. Given an initial state  $s_u$  at vertex  $u$ , the state  $s_v$  at vertex  $v$  depends entirely on the chosen path  $\gamma$  from  $u$  to  $v$  and its associated transformation  $a_\gamma$ .

**Orbits, Holonomy, and the Emergence of Disagreement** The action of  $G(X, A)$  partitions the total state space  $\mathcal{S}$  into orbits. Two states  $s_u \in S_u$  and  $s_v \in S_v$  are in the same orbit if there exists a morphism in  $G(X, A)$  connecting them. A consensus protocol can be viewed as an attempt to drive the system into an orbit containing only a single decision value (or a set of indistinguishable decision values).

The structure of these orbits is profoundly affected by the topology of  $X$  and the non-commutativity of  $A$ . Consider a loop  $\gamma$  in  $X$  based at a vertex  $v$ . This loop corresponds to a morphism in  $G(X, A)$  from  $v$  to  $v$ . The set of all such algebraic transformations associated with loops at  $v$  forms a group, the **holonomy group**  $\text{Hol}_v(X, A)$ .

$$\text{Hol}_v(X, A) = \{a_\gamma \in A \mid \gamma \text{ is a loop in } X \text{ based at } v\}.$$

This group is a subgroup of  $A$  and is isomorphic to the image of the fundamental group  $\pi_1(X, v)$  under the homomorphism induced by the  $A$ -labeling of edges.

The holonomy group captures the path-dependent ambiguity at the heart of consensus failures: - If  $\text{Hol}_v(X, A)$  is trivial (i.e., contains only the identity  $e$ ), then the transformation around any loop is the identity. The state  $s_v$  is restored to itself regardless of the cycle of operations. In this case, the state at any point  $u$  is uniquely determined by the initial state  $s_v$  and the homotopy class of the path from  $v$  to  $u$ . Path-dependence is erased by topological equivalence. - If  $\text{Hol}_v(X, A)$  is non-trivial, a process can execute a sequence of operations corresponding to a loop  $\gamma$  and return to its starting point  $v$  in the execution complex, but find its local state transformed from  $s_v$  to  $a_\gamma \cdot s_v$ , where  $a_\gamma \neq e$ .

Now, consider the role of symmetries. If the holonomy element  $a_\gamma$  belongs to the symmetry group  $A_\sim$ , the resulting state  $a_\gamma \cdot s_v$  is indistinguishable from the original state  $s_v$ . The process at  $v$  has no direct way of knowing that its state has been altered. It may proceed to broadcast information or make decisions based on the assumption that its state is  $s_v$ , while the “true” state is  $a_\gamma \cdot s_v$ . This discrepancy, invisible locally but globally significant, can propagate through the system, creating irresolvable disagreements.

**Example Revisited:** Consider the system with registers  $(x, y)$  and  $A = \langle \text{inc}_x, \text{swap} \rangle$ . Let the protocol complex  $X$  contain a square path representing the concurrent and asynchronous execution of  $\text{inc}_x$  by Process 1 and  $\text{swap}$  by Process 2. - Path 1 ( $\text{inc}_x$  then  $\text{swap}$ ):  $(x, y) \rightarrow (x+1, y) \rightarrow (y, x+1)$ . Total transformation:  $a = \text{swap} \cdot \text{inc}_x$ . - Path 2 ( $\text{swap}$  then  $\text{inc}_x$ ):  $(x, y) \rightarrow (y, x) \rightarrow (y+1, x)$ . Total transformation:  $a = \text{inc}_x \cdot \text{swap}$ . These two paths form a loop  $\gamma$  in  $X$ . The holonomy element associated with this loop is  $a_\gamma = a a^{-1} = (\text{inc}_x \cdot \text{swap}) \cdot (\text{swap} \cdot \text{inc}_x)^{-1} = \text{inc}_x \cdot \text{swap} \cdot \text{inc}_x^{-1} \cdot \text{swap}$ . This is a non-trivial element of  $A$ . If a process traverses this execution loop, its state  $(x, y)$  transforms into  $(y-1, x+1)$ . If the protocol logic only depends on the sum  $x+y$ , then  $\text{swap}$  is a symmetry. The state  $(y, x)$  is indistinguishable from  $(x, y)$ . The holonomy  $a_\gamma$  is *not* a symmetry, yet

its generation was mediated by the symmetry **swap**, illustrating how local indistinguishability can combine with non-commutativity to produce global, tangible discrepancies.

**Bridge to Cohomology** The framework of the update groupoid and its action on state fibers provides the precise geometric and algebraic language needed to formulate consensus failures in terms of non-abelian cohomology.

A global, consistent state assignment would be a function  $f: V(X) \rightarrow S$  from the vertices of the complex to a single, unified state space  $S$ , such that for any directed edge  $e = (u, v)$  with associated transformation  $g_e \in A$ , we have  $f(v) = g_e \cdot f(u)$ . This is known as an **equivariant section**.

The existence of such a section is a strong condition. A consensus failure occurs when no such equivariant section can be found. The obstruction to finding this section is measured by cohomology. Specifically, the set of all possible “twisted” but internally consistent state assignments that fail to be globally synchronized corresponds bijectively to the elements of the first non-abelian cohomology set,  $H^1(G(X, A), S)$ . Each non-trivial element of this set represents a distinct mode of systemic disagreement—a **torsor** of states, which has local solutions but no global one. The structure of the update groupoid, its holonomy, and its symmetries are precisely what give  $H^1$  its non-trivial structure, thus providing a direct, mathematical classification of why consensus can fail. The following chapters will develop this cohomological interpretation in detail.

### Chapter 3.3: Path-Dependence as Non-Trivial Holonomy in the State Space

#### Path-Dependence as Non-Trivial Holonomy in the State Space

In the preceding chapters, we established that modeling state updates with a non-abelian group,  $A$ , provides an algebraic foundation for path-dependence. The order in which operations are applied fundamentally alters the final state of the system, a direct consequence of the non-commutativity of  $A$ . We also introduced the concept of an update groupoid, whose paths correspond to computational executions, acting on the space of possible system states. We now elevate this connection to a more geometric and topological level, demonstrating that path-dependence in distributed systems can be rigorously understood as the phenomenon of **non-trivial holonomy**.

This concept, borrowed from differential geometry, describes how a vector, when “parallel transported” around a closed loop on a curved manifold, may not return to its original orientation. The “twist” it accumulates is its holonomy, a measure of the manifold’s intrinsic curvature. In our context, the “manifold” is the topological space of all possible execution paths, the protocol complex  $X$ . The “vector” is the state of a process,  $s \in S$ . The “parallel transport” is the evolution of this state under a sequence of computational operations. The “curvature” is an emergent property arising from the non-commutativity of the

update group  $A$ . This chapter will formalize this analogy, showing that the holonomy of computational loops is precisely what is measured by the first non-abelian cohomology group,  $H^1(X, A)$ .

### A Connection on the State Space

To formalize the idea of “transporting” a state along a computational path, we must first define a structure analogous to a connection in geometry. Let  $X$  be the protocol complex, whose vertices represent partial states of the system (e.g., the state of a single process after a sequence of events) and whose directed edges represent elementary computational steps (e.g., sending a message, performing a local operation).

The total state of the system resides in a state space  $S$ . The group of state updates,  $A$ , acts on this space. We can visualize this as a fiber bundle over  $X$ , where the fiber over each vertex  $v \in X_0$  is the set of possible states a process can be in at that point in the execution. For our purposes, we focus on the transformations themselves and model the situation as a principal  $A$ -bundle, where the fiber is the group  $A$  itself.

A **connection** in this discrete setting is a map that allows us to relate the states (or the frames of reference for states) at adjacent vertices.

**Definition (Discrete Connection):** A connection on the state space over the protocol complex  $X$  is a function  $\phi : E(X) \rightarrow A$  that assigns to each directed edge  $e = (v, w)$  in  $X$  an element  $\phi(e) \in A$ . This element,  $\phi(v, w)$ , represents the state transformation that occurs when the system evolves from configuration  $v$  to configuration  $w$ .

If the state at vertex  $v$  is  $s_v \in S$ , the state at vertex  $w$  is given by the action of  $\phi(v, w)$  on  $s_v$ :  $s_w = \phi(v, w) \cdot s_v$  (We assume a left action for consistency, though the choice is a matter of convention).

This connection,  $\phi$ , is nothing more than a formalization of the system’s dynamics. It is a local rule specifying how the state changes over every possible atomic step of the computation. In the language of sheaf theory, this corresponds to the transition functions of a sheaf of groups over  $X$ .

### Holonomy of a Computational Path

With the concept of a connection, we can now define the transport of a state along an arbitrary path in the execution space.

**Definition (Parallel Transport):** Let  $\gamma = (v_0, e_1, v_1, e_2, \dots, e_k, v_k)$  be a directed path in the protocol complex  $X$ . The **parallel transport operator** along  $\gamma$ , denoted  $T_\gamma$ , is an element of  $A$  given by the ordered product of the connection elements along the path:  $T_\gamma = \phi(e_k) \circ \phi(e_{k-1}) \circ \dots \circ \phi(e_1) = \phi(v_{k-1}, v_k) \circ \dots \circ \phi(v_0, v_1) \in A$ .

If the system starts in a state  $s_0$  at vertex  $v_0$ , after executing the path  $\gamma$ , its state will be  $s_k = T_\gamma \cdot s_0$ . The path-dependence of the system is now a question about the properties of  $T_\gamma$ . Specifically, what happens when  $\gamma$  is a closed loop?

**Definition (Holonomy):** Let  $\gamma$  be a closed loop in  $X$  based at a vertex  $v_0$ , i.e.,  $\gamma = (v_0, e_1, v_1, \dots, e_k, v_0)$ . The **holonomy** of the loop  $\gamma$ , denoted  $\text{Hol}(\gamma, v_0)$ , is the parallel transport operator around this loop:  $\text{Hol}(\gamma, v_0) = T_\gamma = \phi(v_{k-1}, v_0) \circ \dots \circ \phi(v_0, v_1) \in A$ .

The holonomy is the net transformation a state undergoes after traversing a cycle of operations that returns the system to its initial configuration point.

- If  $\text{Hol}(\gamma, v_0) = \text{id}_A$  (the identity element of  $A$ ) for all loops  $\gamma$ , then the system is **path-independent**. The final state depends only on the end-point of the execution path, not the path taken. Any ambiguity arising from concurrent operations is always resolved consistently.
- If there exists even one loop  $\gamma$  for which  $\text{Hol}(\gamma, v_0) \neq \text{id}_A$ , the system possesses **non-trivial holonomy**. This is the formal signature of path-dependence. It signifies that there are cycles of operations that induce a permanent, non-identity change in the system state.

The set of all holonomy elements for loops based at a vertex  $v_0$ ,  $\{\text{Hol}(\gamma, v_0) \mid \gamma \text{ is a loop at } v_0\}$ , forms a subgroup of  $A$  called the **holonomy group** at  $v_0$ .

## Holonomy and First Non-Abelian Cohomology

The connection between holonomy and non-abelian cohomology is direct and profound. The framework of first cohomology,  $H^1(X, A)$ , provides the precise mathematical machinery for classifying systems with non-trivial holonomy.

Recall that for a group  $G$  (here, the fundamental group  $\pi_1(X, v_0)$ ) and a  $G$ -group  $A$ , the first cohomology set  $H^1(G, A)$  classifies crossed homomorphisms  $f : G \rightarrow A$  modulo an equivalence relation. In our simpler setting, where the action of  $\pi_1$  on  $A$  is trivial (a common case for modeling basic CRDT-like systems where operations are not context-dependent),  $H^1(\pi_1(X, v_0), A)$  simply classifies group homomorphisms  $\text{Hom}(\pi_1(X, v_0), A)$  modulo conjugation by elements of  $A$ .

Our connection map  $\phi : E(X) \rightarrow A$  naturally induces a holonomy map from loops to  $A$ . For this map to be well-behaved, we require it to depend only on the *homotopy class* of the loop. This is ensured by a local consistency condition. On any 2-simplex (a triangle)  $(v_0, v_1, v_2)$  in  $X$ , the path from  $v_0$  to  $v_2$  directly must be equivalent to the path through  $v_1$ . This translates to the **1-cocycle condition**:  $\phi(v_0, v_1) \circ \phi(v_1, v_2) = \phi(v_0, v_2)$ .

A connection  $\phi$  satisfying this condition is a **flat connection**. For a flat connection, the holonomy of a loop depends only on its homotopy class in  $\pi_1(X, v_0)$ . This allows us to define a map:  $h_\phi : \pi_1(X, v_0) \rightarrow A$  given by  $h_\phi([\gamma]) = \text{Hol}(\gamma, v_0)$ .

This map is a group homomorphism. The holonomy of the concatenation of two loops is the product of their individual holonomies:  $\text{Hol}(\gamma_2 \cdot \gamma_1) = \text{Hol}(\gamma_2)\text{Hol}(\gamma_1)$ . Thus, our connection  $\phi$  defines an element of  $\text{Hom}(\pi_1(X, v_0), A)$ .

This homomorphism,  $h_\phi$ , is precisely the **1-cocycle** that represents a class in  $H^1(X, A)$ .


- **Non-trivial Holonomy  $\iff$  Non-trivial Cocycle:** The system exhibits non-trivial holonomy if and only if the homomorphism  $h_\phi$  is non-trivial (i.e., its image is not just the identity element). This means the cocycle representing the system's dynamics is a non-trivial element in  $H^1(X, A)$ .
- **Trivial Holonomy  $\iff$  Trivial Cocycle (Coboundary):** The holonomy is trivial for all loops if and only if  $h_\phi$  maps every element of  $\pi_1(X, v_0)$  to the identity. This occurs when the cocycle  $\phi$  is a **1-coboundary**. A cocycle  $\phi$  is a coboundary if there exists a 0-cochain—a function  $\alpha : X_0 \rightarrow A$  assigning a group element to each vertex—such that for every edge  $(v, w)$ :  $\phi(v, w) = \alpha(v) \circ \alpha(w)^{-1}$ .

If such an  $\alpha$  exists, the holonomy of any loop  $\gamma = (v_0, v_1, \dots, v_k, v_0)$  is:  $\text{Hol}(\gamma) = \phi(v_{k-1}, v_0) \circ \dots \circ \phi(v_0, v_1) = (\alpha(v_{k-1})\alpha(v_0)^{-1}) \circ \dots \circ (\alpha(v_0)\alpha(v_1)^{-1}) = \alpha(v_{k-1}) \circ (\alpha(v_0)^{-1} \circ \alpha(v_0)) \circ \dots \circ (\alpha(v_1)^{-1} \circ \alpha(v_1)) \circ \alpha(v_1)^{-1} \dots = \text{id}_A$ . (Note: the exact formula depends on conventions, e.g.,  $\phi(v, w) = \alpha(v)^{-1}\alpha(w)$ , but the result is the same).

The existence of the 0-cochain  $\alpha$  means that the path-dependence is merely an artifact of our “coordinate system”. We can perform a “gauge transformation” at each vertex, defined by  $\alpha$ , to a new state representation in which all path-dependence vanishes. A non-trivial class in  $H^1(X, A)$  signifies that no such global re-coordination exists. The path-dependence is an intrinsic, topological property of the system.

### Interpretation: Holonomy as Irreconcilable Histories

A loop in the protocol complex  $X$  represents a cycle of causal ambiguity. The canonical example is the “diamond” shape common in models of concurrency:



Here, two concurrent operations,  $u_A$  and  $u_B$ , can be observed in different orders by different parts of the system. One computational path is  $\gamma_1 = v_0 \rightarrow v_A \rightarrow$



$v_{final}$ . Another is  $\gamma_2 = v_0 \rightarrow v_B \rightarrow v_{final}$ . The loop is  $\gamma = v_0 \rightarrow v_A \rightarrow v_{final} \rightarrow v_B \rightarrow v_0$ .

Let the state transformations be given by group elements  $g_A, g_B \in A$ . \* The state at  $v_{final}$  via path 1 is  $s_1 = (g_{A'} \circ g_A) \cdot s_0$ . \* The state at  $v_{final}$  via path 2 is  $s_2 = (g_{B'} \circ g_B) \cdot s_0$ .

If the operations are context-independent (the “constant sheaf” case), then  $g_A$  and  $g_B$  are the same along both parallel edges. The transformations to reach  $v_{final}$  are  $g_B \circ g_A$  and  $g_A \circ g_B$ . The disagreement between these two final states is a direct result of the non-commutativity of  $g_A$  and  $g_B$ .

The holonomy of the loop  $\gamma$  is  $\text{Hol}(\gamma) = g_B \circ g_A \circ (g_B)^{-1} \circ (g_A)^{-1} = [g_B, g_A]$ , the commutator of the two operations (up to inverses depending on path direction).

A non-trivial holonomy,  $[g_B, g_A] \neq \text{id}_A$ , means that the ambiguity inherent in the concurrent execution of  $u_A$  and  $u_B$  leads to an *irreconcilable* difference in the global state. Even after the system reaches a point  $v_{final}$  where the effects of both operations are known, there is no single, consistent state. Instead, there is a “torsor” of possible states, related to each other by the elements of the holonomy group. This is not a transient error; it is a permanent schism in the state space, an algebraic scar left by a topological feature of the execution.

This is precisely why systems like CRDTs (Commutative Replicated Data Types) succeed in providing eventual consistency: they are explicitly designed by restricting the update group  $A$  to be abelian. If  $A$  is abelian, all commutators are the identity, all holonomy is trivial, and  $H^1(X, A)$  must be trivial (for a connected complex  $X$ ). Path-dependence is eliminated by algebraic design.

## Conclusion: The Algebraic Signature of Path-Dependence

The geometric language of connections and holonomy provides a powerful and intuitive framework for understanding path-dependence in distributed systems. It translates the abstract algebraic property of non-commutativity into the concrete phenomenon of a state failing to return to its origin after a cycle of computations. This holonomy is the physical manifestation of a non-trivial cocycle in non-abelian cohomology.

We have established a key equivalence: **Path-Dependence**  $\iff$  **Non-Trivial Holonomy**  $\iff$  **Non-Trivial**  $H^1(X, A)$

This insight is central to our thesis. It demonstrates that certain consensus failures are not due to implementation errors or transient faults, but are fundamental obstructions encoded in the interplay between the system’s topology of execution ( $X$ ) and its algebra of updates ( $A$ ). The holonomy group, as the image of the cocycle  $h_\phi : \pi_1(X, v_0) \rightarrow A$ , quantifies the exact nature and extent of the disagreement. In the subsequent chapters, we will use this foundation to classify these obstructions as “torsors of disagreement,” linking the mathematical structure of  $H^1(X, A)$  directly to the analysis of specific consensus protocols

and their impossibility results.

### Chapter 3.4: Crossed Homomorphisms: The Algebraic Signature of Inconsistent Views

#### Crossed Homomorphisms: The Algebraic Signature of Inconsistent Views

In the preceding chapters, we established a framework where path-dependence in a distributed system is captured by a holonomy representation, a group homomorphism  $\rho : \pi_1(X) \rightarrow \text{Aut}(S)$ , where  $\pi_1(X)$  is the fundamental group of the execution space and  $S$  is the state space. This model powerfully describes how a globally consistent state can evolve in a non-trivial way along different computational paths. However, it presupposes the existence of such a globally consistent state. The most pernicious failures in distributed systems, particularly those characterized by asynchrony and faults, are precisely those where the very notion of a single global state breaks down.

In such scenarios, each process maintains a local, partial view of the system. These views are not merely out of sync; they can be fundamentally inconsistent. A process might receive contradictory information about the state of a distant peer, leading to a schizophrenic global picture that cannot be resolved into a single coherent snapshot. The holonomy representation, which describes the evolution of a *single* state, is insufficient. We require a more nuanced algebraic structure to capture the *relationship* between multiple, inconsistent local views—a structure that measures the discrepancy itself as a function of the computational path. This is the role of the crossed homomorphism, the algebraic signature of an inconsistent system view.

**The Twisted Homomorphism Condition** Let us formalize the scenario. We have a group  $G$  (which we will identify with the fundamental group of executions,  $\pi_1(X)$ ) acting on a non-abelian group  $A$  (the group of state differences or transformations). This action, denoted  $g \cdot a$  for  $g \in G, a \in A$ , is precisely the holonomy we have already discussed; it describes how the “reference frame” for state transformations changes as we traverse an execution loop  $g$ .

Now, imagine a function  $f : G \rightarrow A$  that assigns to each execution loop  $g$  an element  $f(g) \in A$ . We interpret  $f(g)$  as the *net discrepancy* or *error offset* observed by a process after traversing the loop  $g$ . It is the residual change in the local view that cannot be accounted for by the expected, globally consistent holonomy. If the system were perfectly consistent, composing two paths  $g$  and  $h$  would simply compose their discrepancies:  $f(gh) = f(g)f(h)$ . This would be a standard group homomorphism.

But in an inconsistent system, the discrepancy measured along path  $h$  is perceived differently after traversing path  $g$ . The path  $g$  alters the observer’s context by the holonomy action. Therefore, the discrepancy of  $h$ , when viewed from the “end” of path  $g$ , is not  $f(h)$  but rather  $g \cdot f(h)$ . The total discrepancy for the composite path  $gh$  is thus the discrepancy from  $g$ , followed by the  $g$ -transformed

discrepancy from  $h$ . This intuition is captured precisely in the definition of a **crossed homomorphism**.

**Definition:** A map  $f : G \rightarrow A$  is a **crossed homomorphism** (or a **1-cocycle**) if for all  $g, h \in G$ , it satisfies the condition:

$$f(gh) = f(g) \cdot (g \cdot f(h))$$

This “twisted” homomorphism law is the algebraic fingerprint of state inconsistency. The term  $g \cdot f(h)$  is the crucial addition that separates this concept from a simple holonomy map. It encodes the interaction between the path taken ( $g$ ) and the observed discrepancy ( $f(h)$ ) along a subsequent path. It tells us that the inconsistencies are not static but are themselves transformed by the execution flow.

**From Algebraic Condition to System Disagreement** To understand the operational meaning of this definition, consider a distributed system where consensus has failed. There is no global state  $s$ . Instead, each process  $i$  has a local state  $s_i$ . We can model these local states as being related to a hypothetical, unreachable consensus state  $s$  by some transformation  $a_i \in A$ , such that  $s_i = a_i \cdot s$ . The set of transformations  $\{a_i\}$  represents the total disagreement in the system.

Now, consider a process at a basepoint  $x_0$  in our execution space  $X$ . It has a local state  $s_0$ . Let’s say it embarks on a computational loop corresponding to an element  $g \in \pi_1(X, x_0)$ . This loop involves communicating with other processes and updating its state based on the messages received. Upon returning to its initial point in the execution topology, its state has changed. Part of this change is expected due to the non-abelian nature of the updates—this is the holonomy, which would transform its state to  $g \cdot s_0$ .

However, if the process receives conflicting information during the loop (e.g., from a Byzantine node, or due to race conditions leading to inconsistent reads), an additional, anomalous transformation occurs. The final state is not  $g \cdot s_0$ , but rather  $s' = a \cdot (g \cdot s_0)$  for some  $a \in A$ . The crossed homomorphism  $f$  maps the loop  $g$  to this anomalous transformation:  $f(g) = a$ .

The crossed homomorphism condition  $f(gh) = f(g) \cdot (g \cdot f(h))$  now has a clear physical meaning: 1. **Traverse loop  $h$ :** An anomalous transformation  $f(h)$  is accumulated. 2. **Traverse loop  $g$ :** An anomalous transformation  $f(g)$  is accumulated. Additionally, the previous anomaly  $f(h)$  is “viewed” through the lens of path  $g$ , transforming it to  $g \cdot f(h)$ . 3. **Total Anomaly:** The total anomaly for the path  $gh$  is the composition of the anomaly from  $g$  and the transformed anomaly from  $h$ , yielding  $f(g) \cdot (g \cdot f(h))$ .

This structure precisely captures a disagreement that is path-dependent and self-referential. The error accumulated along one path affects the measurement of error along subsequent paths.

**Cohomology: Classifying True Inconsistencies** Not all discrepancies are created equal. Some apparent inconsistencies might simply be artifacts of a poor choice of initial reference frame. For instance, if every process starts with an identical, but “wrong,” initial state, they might perceive discrepancies along execution paths that could be eliminated entirely if they all agreed to shift their initial state by some common transformation. These are trivial, or resolvable, inconsistencies.

Algebraically, this corresponds to the notion of a **1-coboundary**. A crossed homomorphism  $f : G \rightarrow A$  is called a principal crossed homomorphism or a 1-coboundary if there exists an element  $a \in A$  such that for all  $g \in G$ :

$$f(g) = a^{-1} \cdot (g \cdot a)$$

**Interpretation:** The discrepancy  $f(g)$  observed along any loop  $g$  is not a fundamental property of the system’s dynamics. Instead, it is entirely explained by having started with a “bad” initial reference point, represented by  $a$ . The term  $g \cdot a$  is the transformed version of this initial offset after traversing loop  $g$ . The discrepancy  $f(g)$  is simply the difference between the transformed offset and the original one ( $a^{-1}$ ). If we were to “correct” our initial state by applying  $a$ , all these apparent discrepancies would vanish. In such a case, a consistent global state *can* be defined, although we may have chosen a “gauge” that obscured it.

The set of all crossed homomorphisms (1-cocycles) is denoted  $Z^1(G, A)$ . The set of all 1-coboundaries is a subgroup denoted  $B^1(G, A)$ . The fundamental, unresolvable inconsistencies are captured by the quotient:

$$H^1(G, A) = Z^1(G, A) / B^1(G, A)$$

This is the **first cohomology group** of  $G$  with coefficients in the non-abelian  $G$ -module  $A$ . \* If  $H^1(G, A)$  is the trivial group (often denoted as 0 or  $\{*\}$ ), it means that every crossed homomorphism is a coboundary. Any observed systemic disagreement can be resolved by a global “change of coordinates” (choosing a different initial state reference). Consensus is achievable. \* If  $H^1(G, A)$  is non-trivial, there exists at least one crossed homomorphism  $f$  that is *not* a coboundary. This represents a “true” inconsistency—a topological obstruction woven into the fabric of the protocol’s execution space and the state update algebra. No simple re-initialization can eliminate this disagreement. This is the mathematical signature of an irrecoverable consensus failure.

**The Torsor of Disagreement: The Geometric View** What is the concrete manifestation of a non-trivial element in  $H^1(G, A)$ ? It is the geometric structure of the set of possible system states. When  $H^1(G, A)$  is trivial, the system possesses a valid global section, meaning a consistent global state can be chosen. When  $H^1(G, A)$  is non-trivial, no such global state exists. The collection of all possible local states form a geometric object known as an **A-torsor**.

An **A-torsor** (or a principal homogeneous space for  $A$ ) is a set  $S$  on which the group  $A$  acts freely and transitively. \* **Transitive:** For any two states  $s_1, s_2 \in S$ , there exists a unique group element  $a \in A$  such that  $s_2 = a \cdot s_1$ . The group  $A$  can connect any state to any other state. \* **Free:** If  $a \cdot s = s$  for some  $s \in S$ , then  $a$  must be the identity element  $e \in A$ . No non-trivial update can act as a stabilizer of a state.

An  $A$ -torsor is like the group  $A$  itself, but with its origin—the identity element—forgotten. There is no special, canonical “base state.” Any state is as valid as any other, and the only defined operation is finding the unique transformation  $a$  that relates two states.

The correspondence is profound: **the first cohomology set  $H^1(X, \mathcal{A})$  (where  $\mathcal{A}$  is the sheaf of groups of local states) is in one-to-one correspondence with the set of isomorphism classes of  $\mathcal{A}$ -torsors on the space  $X$ .**

A non-trivial class in  $H^1(X, A)$  corresponds to a non-trivial torsor of states—a “twisted” configuration space that admits no global section. The consensus problem, from this perspective, is the problem of determining whether the torsor of global states is trivial. A trivial torsor is one that is isomorphic to the group  $A$  itself, which *does* have a canonical origin (the identity element). This identity element corresponds to the agreed-upon consensus value. Therefore:

**Consensus is possible if and only if the torsor of global states is trivial, which is equivalent to the vanishing of the first cohomology group,  $H^1(X, A) = 0$ .**

The crossed homomorphism, therefore, is more than an algebraic curiosity. It is the cocycle representative of a cohomology class that directly measures the “twistedness” of the state space. It is the analytical tool that detects the presence of a non-trivial torsor of disagreement, providing the precise algebraic reason for a consensus failure. By calculating  $H^1(X, A)$ , we can move from a qualitative statement of impossibility (like the FLP result) to a quantitative classification of the structure of disagreement itself.

### Chapter 3.5: Case Study: Consensus Failure in Systems with Non-Commutative Operations

#### Case Study: Consensus Failure in Systems with Non-Commutative Operations

The theoretical framework developed in the preceding chapters posits that path-dependent inconsistencies in distributed systems can be rigorously described as non-trivial elements in the first non-abelian cohomology group,  $H^1(X, \mathcal{A})$ . This framework identifies the failure to reach a single, globally agreed-upon state with the existence of a non-trivial  $\mathcal{A}$ -torsor over the protocol’s execution space  $X$ . To ground these abstract concepts, this chapter presents a detailed case study of a distributed system where the state updates are inherently non-commutative. By analyzing a specific execution, we will trace the emergence of path-dependence,

calculate the resulting holonomy, and demonstrate how the global system state manifests as a cohomological obstruction to consensus.

**1. The System: Distributed Orientation Control** To construct a clear and intuitive example of non-commutative state updates, we will consider a system designed to manage the orientation of a rigid 3D object in space. This scenario is a practical analogue for various real-world problems, such as controlling a satellite’s attitude, coordinating robotic arms, or managing viewports in a collaborative 3D modeling application.

**The Components:**

- **Processes ( $P$ ):** A set of  $n$  server processes,  $\{P_1, P_2, \dots, P_n\}$ , tasked with maintaining a consistent view of the object’s orientation.
- **State Space ( $S$ ):** The state of the system is the orientation of the object. This can be uniquely represented by an element of the special orthogonal group in 3 dimensions,  $SO(3)$ . Each element of  $SO(3)$  is a  $3 \times 3$  orthogonal matrix with determinant  $+1$ , representing a rotation in Euclidean space. So, the state space  $S$  is the manifold  $SO(3)$ .
- **The Update Group ( $A$ ):** Clients issue commands to rotate the object. Each command corresponds to a specific rotation, which is an element of  $SO(3)$ . The set of all possible update operations forms the group  $A = SO(3)$ . A state transition is effected by left-multiplying the current state matrix  $s \in S$  by an update matrix  $a \in A$ :  $s_{new} = a \cdot s_{old}$ . The group  $A = SO(3)$  is famously non-abelian. For instance, a 90-degree rotation around the x-axis,  $R_x(\pi/2)$ , followed by a 90-degree rotation around the y-axis,  $R_y(\pi/2)$ , yields a different final orientation than performing the operations in the reverse order:  $R_y(\pi/2) \cdot R_x(\pi/2) \neq R_x(\pi/2) \cdot R_y(\pi/2)$ .
- **The Protocol and Execution Space ( $X$ ):** The system operates asynchronously. Processes receive update commands from clients and broadcast them to their peers. Due to network latency, the order in which these broadcasts are received and applied is not guaranteed to be the same for all processes. The set of all possible timed executions, accounting for all possible message orderings and interleavings, constitutes the protocol’s execution space, which we model as a topological space  $X$  (specifically, a simplicial or cubical complex). The vertices of this complex represent local states  $(P_i, k)$ , where  $P_i$  has processed  $k$  events, and higher-dimensional simplices represent states of shared knowledge resulting from communication.

**2. An Asynchronous Execution and the Emergence of Path-Dependence** Let’s analyze a simple yet critical execution involving two processes,  $P_1$  and  $P_2$ , and two concurrent client requests. Assume the initial state of the object, known to all processes, is the identity orientation,  $s_0 = I$ .

1. **Client A** sends an update command  $op_x = R_x(\pi/2)$  (rotate 90° about the x-axis). The message is routed to  $P_1$ .

2. **Client B** sends an update command  $\text{op\_y} = R_y(\pi/2)$  (rotate  $90^\circ$  about the y-axis). The message is routed to  $P_2$ .

Due to the concurrent nature of these requests and the asynchrony of the network, the following sequence of events occurs:

- **Path for  $P_1$ :**
  - $P_1$  receives  $\text{op\_x}$  from Client A.
  - $P_1$  applies the update:  $s_{1,\text{interim}} = \text{op}_x \cdot s_0 = R_x(\pi/2)$ .
  - $P_1$  broadcasts  $\text{op\_x}$  to other processes, including  $P_2$ .
  - Sometime later,  $P_1$  receives the broadcast of  $\text{op\_y}$  from  $P_2$ .
  - $P_1$  applies this second update:  $s_1 = \text{op}_y \cdot s_{1,\text{interim}} = R_y(\pi/2) \cdot R_x(\pi/2) \cdot s_0$ .
- **Path for  $P_2$ :**
  - $P_2$  receives  $\text{op\_y}$  from Client B.
  - $P_2$  applies the update:  $s_{2,\text{interim}} = \text{op}_y \cdot s_0 = R_y(\pi/2)$ .
  - $P_2$  broadcasts  $\text{op\_y}$  to other processes, including  $P_1$ .
  - Sometime later,  $P_2$  receives the broadcast of  $\text{op\_x}$  from  $P_1$ .
  - $P_2$  applies this second update:  $s_2 = \text{op}_x \cdot s_{2,\text{interim}} = R_x(\pi/2) \cdot R_y(\pi/2) \cdot s_0$ .

At the end of this execution sequence, both processes believe they have processed the complete set of operations,  $\{\text{op}_x, \text{op}_y\}$ , but their local states have diverged:

$$s_1 = R_y(\pi/2) \cdot R_x(\pi/2) \neq R_x(\pi/2) \cdot R_y(\pi/2) = s_2$$

This is a quintessential example of **path-dependent failure**. The final state depends on the path taken through the execution space.  $P_1$  followed the path “apply  $\text{op}_x$ , then apply  $\text{op}_y$ ,” while  $P_2$  followed the path “apply  $\text{op}_y$ , then apply  $\text{op}_x$ .” In the topological model  $X$ , these two execution histories correspond to two distinct directed paths between the initial event (no operations applied) and the final event (both operations applied).

**3. Holonomy as the Measure of Disagreement** The discrepancy between the final states  $s_1$  and  $s_2$  is not random; it is a specific element of the update group  $A = SO(3)$ . This discrepancy can be quantified as the transformation required to map one state onto the other. Let us define this transformation as  $h$ :

$$s_1 = h \cdot s_2 \quad h = s_1 \cdot s_2^{-1} = (R_y R_x s_0) \cdot (R_x R_y s_0)^{-1} = R_y R_x s_0^{-1} R_y^{-1} R_x^{-1} = R_y R_x R_y^{-1} R_x^{-1}$$

This element  $h$  is the group commutator of  $R_y$  and  $R_x$ , often denoted  $[R_y, R_x]$ . Since  $SO(3)$  is non-abelian, this commutator is not the identity element  $I$ .

This commutator has a profound geometric and topological meaning. The two different execution paths taken by  $P_1$  and  $P_2$  form a loop in the execution space  $X$ . This loop can be visualized in a “state-time” diagram as a diamond or square shape, a fundamental feature of asynchronous systems sometimes called a

“concurrency square.” The element  $h$  is precisely the **holonomy** of the system’s state bundle when transported around this loop. It measures the “twist” or “error” accumulated by traversing a cycle in the space of possible computations. The fact that  $h \neq I$  signifies that the system’s state space has a non-trivial curvature, induced by the non-commutativity of the update group.

**4. The Global State as a Non-Trivial  $A$ -Torsor** Now, let’s expand our view from two processes to the entire system  $\{P_1, \dots, P_n\}$ . In a complex asynchronous execution with many non-commuting operations, each process  $P_i$  will have observed a unique permutation of the operations, resulting in a unique local state  $s_i$ . The collection of all local states,  $\mathcal{S} = \{s_1, s_2, \dots, s_n\}$ , constitutes the global “state” of the system.

Consensus fails because there is no single state  $s \in S$  that can be identified with  $\mathcal{S}$ . However, this collection of states is not unstructured. It forms an  **$A$ -torsor**.

- **The Torsor Property:** For any two local states  $s_i$  and  $s_j$  in the collection  $\mathcal{S}$ , there exists a unique transformation  $a_{ij} \in A$  that relates them:  $s_j = a_{ij} \cdot s_i$ . This transformation  $a_{ij}$  is the composite of operations that  $P_j$  has applied in a different order than  $P_i$ . The group  $A = SO(3)$  acts freely and transitively on the set  $\mathcal{S}$  of observed states.
- **The Cohomological Interpretation:** This system of transformations  $\{a_{ij}\}$  defines a **1-cocycle**. Consider three processes  $P_i, P_j, P_k$ . The transformations must satisfy the cocycle condition:  $a_{ik} = a_{jk} \cdot a_{ij}$ . This simply means that transforming from state  $s_i$  to  $s_k$  is the same as transforming from  $s_i$  to  $s_j$  and then from  $s_j$  to  $s_k$ .

The inability to find a global consensus state  $s$  is equivalent to this cocycle being **non-trivial** (i.e., not a coboundary). If consensus were possible, there would exist a global state  $s$  such that for every process  $P_i$ , its local state  $s_i$  could be “corrected” back to  $s$  by some transformation  $b_i \in A$ . That is,  $s = b_i \cdot s_i$  for all  $i$ . If such a set of “correction factors”  $\{b_i\}$  existed, the cocycle  $\{a_{ij}\}$  would be a coboundary, as we could write:  $a_{ij} = s_j \cdot s_i^{-1} = (b_j^{-1} \cdot s) \cdot (b_i^{-1} \cdot s)^{-1} = b_j^{-1} \cdot s \cdot s^{-1} \cdot b_i = b_j^{-1} \cdot b_i$ . This is the definition of a 1-coboundary.

In our case study, the holonomy  $h = [R_y, R_x] \neq I$  around a single execution loop proves that at least one such cocycle value is non-trivial, preventing the existence of such a set  $\{b_i\}$ . Therefore, the set of states  $\mathcal{S}$  represents a non-trivial  $A$ -torsor, an element  $[a] \in H^1(X, A)$  where  $[a] \neq 1$ . The consensus failure is not a bug in any single process’s logic but a global, topological obstruction inherent to the system’s design.

**5. Implications for System Design** This case study illuminates why achieving consistency in the face of non-commutative operations is so challenging and why different families of distributed protocols make the design choices they do.



- **Total Ordering (e.g., Paxos, Raft):** Classical consensus algorithms like Paxos and Raft fundamentally solve this problem by eliminating path-dependence. They use complex protocols to agree on a single, global, totally-ordered log of operations. This forces every process to execute the *same* sequence of updates, effectively collapsing the execution space  $X$  into a single line. All paths are identical, so no loops with non-trivial holonomy can form. In our language, they ensure that  $H^1(X, A)$  is trivial by forcing  $X$  to be contractible. The price is significant overhead in coordination and reduced concurrency.
- **Commutativity (e.g., CRDTs):** Conflict-Free Replicated Data Types (CRDTs) take the opposite approach. Instead of restricting the topology of  $X$ , they restrict the algebra of  $A$ . CRDTs are designed such that all update operations are commutative. If  $A$  were an abelian group, then the commutator  $[op_y, op_x]$  would always be the identity, the holonomy would always be trivial, and path-dependence would vanish. All  $A$ -torsors would be trivial, and simple state-based merging would guarantee convergence.

This case study demonstrates that the framework of non-abelian cohomology provides the precise mathematical language to articulate the deep connection between the algebraic properties of state updates (commutativity) and the topological properties of the execution space (path-dependence). It reframes consensus failure from a mere programming error into a fundamental obstruction, measurable by holonomy and classifiable by cohomology, offering a powerful lens through which to analyze, understand, and design the next generation of distributed systems.

## Part 4: Cohomological Obstructions to Distributed Consensus: Torsors and Gerbes

### Chapter 4.1: The Torsor of Disagreement: $H^1(X, A)$ as the Formal Measure of Consensus Impossibility

The Torsor of Disagreement:  $H^1(X, A)$  as the Formal Measure of Consensus Impossibility

In the preceding chapters, we established a topological model for distributed executions—the protocol complex  $X$ —and an algebraic model for state transformations using a (potentially non-abelian) group  $A$ . We now arrive at the central synthesis of these two perspectives. This chapter demonstrates that the first non-abelian cohomology set, denoted  $H^1(X, A)$ , serves as a precise mathematical formalism for measuring the impossibility of achieving global consensus. The non-trivial elements of this set correspond to robust, topologically-enshrined states of disagreement, which we term *Torsors of Disagreement*.

The core thesis is as follows: **Consensus is achievable if and only if the first cohomology  $H^1(X, A)$  is trivial.** A non-trivial  $H^1(X, A)$  signifies the existence of an intrinsic, global obstruction that prevents local views from being

reconciled into a single, consistent global state. The elements of this set do not merely indicate failure; they classify the distinct, fundamental *modes* of failure.

**From Local Data to Global Obstruction: The 1-Cocycle of Disagreement** Imagine a distributed system where each process  $i$  maintains a local state  $s_i$ . We assume these states belong to a set  $S$  on which our group of discrepancies,  $A$ , acts. In an ideal world, consensus means finding a global state  $s$  such that  $s_i = s$  for all  $i$ . In reality, processes only have partial, local information.

Consider a small part of the protocol complex  $X$ , represented by a set of communicating processes. For any two processes,  $i$  and  $j$ , that can exchange information (represented by a 1-simplex, or edge,  $(i, j)$  in  $X$ ), they can compare their local states. The discrepancy between their views can be captured by a group element  $g_{ij} \in A$ , such that if  $s_i$  is the state from  $i$ 's perspective, then  $j$  sees it as  $g_{ij} \cdot s_i$ . This  $g_{ij}$  is the “gauge transformation” or “local translation” required to align  $i$ 's view with  $j$ 's. For this to be a consistent model, moving from  $i$ 's view to  $i$ 's view should require no transformation, so  $g_{ii} = e$  (the identity in  $A$ ). Furthermore, the transformation from  $j$  to  $i$  should be the inverse of the transformation from  $i$  to  $j$ , so  $g_{ji} = g_{ij}^{-1}$ .

This assignment of group elements  $g_{ij}$  to the edges of our protocol complex  $X$  constitutes a **1-cochain**. It represents the raw data of pairwise disagreements across the system.

The critical question is whether this collection of pairwise disagreements is internally consistent. Consider a 2-simplex (a triangle)  $(i, j, k)$  in  $X$ , representing a situation where processes  $i$ ,  $j$ , and  $k$  can all mutually communicate. For the system's local views to be consistent, the direct path from  $i$  to  $k$  must be equivalent to the path that goes via  $j$ . The transformation from  $i$ 's view to  $k$ 's view is  $g_{ik}$ . The transformation via  $j$  is more complex due to the non-abelian nature of  $A$ . If  $A$  acts on itself by conjugation, the composed transformation is  $g_{ij} \cdot g_{jk}$ . The condition for local consistency on the simplex  $(i, j, k)$  is therefore:

$$g_{ik} = g_{ij} \cdot g_{jk}$$

This is the **1-cocycle condition**. A 1-cochain  $\{g_{ij}\}$  that satisfies this condition for all 2-simplices in  $X$  is called a **1-cocycle**. A cocycle represents a state of disagreement that is *locally consistent*. There are no contradictions within any small, fully connected neighborhood of processes. The disagreement, if it exists, must be a global phenomenon.

**The Possibility of Consensus: Trivial Cocycles** If a 1-cocycle  $\{g_{ij}\}$  represents a state of disagreement, when can this disagreement be resolved to achieve consensus? It can be resolved if the disagreement is merely a matter of perspective—an artifact of each process using a different local “coordinate system” to define its state.

Formally, this means we can find a “correction factor” for each process  $i$ , represented by a group element  $h_i \in A$ , such that applying this correction ( $s_i' = h_i^{-1} \cdot s_i$ ) results in a unified global state. If such a set of corrections  $\{h_i\}$  (a **0-cochain**) exists, the observed disagreement  $g_{ij}$  can be explained entirely by these local coordinate choices:

$$g_{ij} = h_i \cdot h_j^{-1}$$

A 1-cocycle  $\{g_{ij}\}$  that can be expressed in this form is called a **1-coboundary**. A coboundary represents a *trivial* or *resolvable* disagreement. The existence of the 0-cochain  $\{h_i\}$  is tantamount to the existence of a global section—a single consistent state. By “re-gauging” each local state  $s_i$  with  $h_i$ , the system converges.

The set of all 1-cocycles is denoted  $Z^1(X, A)$ . The set of all 1-coboundaries is denoted  $B^1(X, A)$ . The first cohomology **set**,  $H^1(X, A)$ , is defined as the quotient of cocycles by an equivalence relation derived from coboundaries:

$$H^1(X, A) = Z^1(X, A) / \sim$$

where two cocycles  $g$  and  $g'$  are equivalent if there exists a 0-cochain  $h$  such that  $g'_{ij} = h_i^{-1} \cdot g_{ij} \cdot h_j$  for all  $(i, j)$ . When  $A$  is abelian, this simplifies to the familiar quotient group  $Z^1/B^1$ .

The trivial element of  $H^1(X, A)$  is the class of all coboundaries. If  $H^1(X, A)$  contains only this one element, it means that any locally consistent disagreement (any 1-cocycle) is necessarily a coboundary, and thus resolvable. Therefore:

**Consensus is possible**  $H^1(X, A) = \{\}$ , *where  $\{\}$  is the trivial element.*

**The Geometric Interpretation: The Torsor of Disagreement** The algebraic formalism of cocycles and coboundaries has a profound and intuitive geometric counterpart: the **A-torsor**. An  $A$ -torsor is the geometric embodiment of unresolvable disagreement.

**Definition:** An **A-torsor** (or principal  $A$ -bundle) over the protocol complex  $X$  is a space  $P$  equipped with a projection map  $p: P \rightarrow X$  and a right action of the group  $A$  on  $P$ . This structure must satisfy two conditions: 1. The action is free and transitive on the fibers: For any point  $x \in X$ , the group  $A$  acts on the fiber  $P_x = p^{-1}(x)$  such that for any two elements  $y_1, y_2 \in P_x$ , there exists a unique  $a \in A$  with  $y_2 = y_1 \cdot a$ . 2.  $P$  is locally trivial: For any point  $x \in X$ , there is an open neighborhood  $U$  such that the portion of the torsor over  $U$ ,  $p^{-1}(U)$ , is isomorphic to the trivial product  $U \times A$ .

Let’s unpack this in the context of distributed systems. \*  $X$  is the base space of possible system configurations. \* The fiber  $P_x$  over a configuration  $x$  represents the set of all possible valid, internally consistent global state assignments that could correspond to that configuration. \* The group  $A$  acts on this fiber, transforming one valid global state assignment into another. The fact that the action is *free and transitive* means that from any given valid state, all other

valid states are reachable via a unique transformation from  $A$ , and there is no preferred “zero” or “identity” state within the fiber itself. \* The *local triviality* condition captures the essence of local consistency. Over a small enough patch of the execution space, the situation looks simple. We can pick an arbitrary reference state (a local section) and describe all other states in relation to it, just like a product  $U \times A$ .

The fundamental obstruction to consensus is the potential for  $P$  to be **globally non-trivial**. While it looks like  $X \times A$  in small pieces, its global structure might be “twisted.” A classic analogy is the Möbius strip, which is a torsor for the group  $\mathbb{Z}_2 = \{+1, -1\}$  over the circle  $S^1$ . Locally, it’s just a strip (an interval  $\times$  ), but globally it has a twist and no continuous, non-zero global section.

A **global section** of the torsor  $P$  is a continuous map  $s: X \rightarrow P$  such that  $p \circ s = \text{id}_X$ . Such a section would, for every configuration  $x \in X$ , pick out a specific state  $s(x) \in P_x$  in a globally consistent way. **A global section is precisely a global consensus state.**

The central theorem connecting the algebra and geometry is:

**Theorem:** The set of isomorphism classes of  $A$ -torsors over  $X$  is in a canonical one-to-one correspondence with the first cohomology set  $H^1(X, A)$ .

Under this correspondence: \* The **trivial torsor**,  $P = X \times A$ , corresponds to the trivial element in  $H^1(X, A)$ . This torsor always admits global sections (e.g.,  $s(x) = (x, e)$ ). \* A **non-trivial torsor** corresponds to a non-trivial element in  $H^1(X, A)$ . A non-trivial torsor, by definition, has no global section.

This provides the ultimate interpretation. The impossibility of consensus is not just a computational failure; it is a topological property of the system. The system’s state space forms a “Torsor of Disagreement”—a twisted bundle where the fibers of possible states are woven together in such a way that no single, continuous choice of state can be made across the entire execution space. The non-trivial cocycle  $\{g_{ij}\}$  serves as the transition functions that define the “twist” in this bundle.

**Holonomy and the Fundamental Group** For path-connected spaces  $X$ , there is a deep connection between  $H^1(X, A)$  and the fundamental group  $\pi_1(X, x)$ . Specifically, there is a mapping from  $H^1(X, A)$  to the set of conjugacy classes of group homomorphisms  $\text{Hom}(\pi_1(X), A)$ .

In our context, a loop in the protocol complex  $X$  represents a sequence of operations and communications that returns the system to its initial configuration. As we traverse this loop, the local transformations  $\{g_{ij}\}$  accumulate. The total transformation upon returning to the start point is the **holonomy** of the loop.

- If for every loop in  $X$ , the holonomy is the identity element  $e \in A$ , then the cocycle is trivial, and consensus is possible.

- If there exists a loop in  $X$  with a non-trivial holonomy  $a \in A$ ,  $a \neq e$ , this signifies a non-trivial element in  $H^1(X, A)$ . This loop represents a cycle of operations that introduces an irresolvable discrepancy. No matter how local states are re-calibrated, traversing this execution path will always result in a state that differs from the starting point by the transformation  $a$ . This is the algebraic signature of a topological obstruction embodied by the **Torsor of Disagreement**. Famous impossibility results, like the FLP impossibility, can be re-cast as statements about the existence of execution loops with non-trivial holonomy in an appropriately defined protocol complex.

In summary,  $H^1(X, A)$  provides a complete and formal classification of consensus impossibility. It elevates the analysis from specific protocol behaviors to the intrinsic topological and algebraic properties of the problem itself. A non-trivial cohomology class is not merely an indicator of failure; it is a mathematical object—the Torsor of Disagreement—that gives geometric form to the very structure of inescapable conflict within a distributed system. The next chapter will explore the implications of higher cohomology,  $H^2(X, A)$ , which governs obstructions at an even more abstract level: the failure to even agree on a consistent protocol.

## Chapter 4.2: Interpreting Non-Abelian Cocycles as Local State Mismatches

### Interpreting Non-Abelian Cocycles as Local State Mismatches

In the preceding chapter, we established that the impossibility of achieving global consensus in a distributed system with state group  $A$  over a protocol complex  $X$  is formally captured by the non-triviality of the first non-abelian cohomology set,  $H^1(X, A)$ . Each element of this set corresponds to an equivalence class of  $A$ -torsors, which we termed “torsors of disagreement.” A torsor represents a global state of irreconcilable disagreement, a system where local consistency is maintained, yet no global state section exists.

While this provides a powerful classification of global failure modes, it remains an abstract characterization. To bridge the gap between this high-level algebraic invariant and the concrete, operational failures within a protocol, we must “unwind” the torsor and examine its constituent data. This data is the **non-abelian 1-cocycle**. This chapter will demonstrate that a 1-cocycle can be directly interpreted as a web of *local state mismatches*—the precise, quantifiable discrepancies in perspective between communicating processes that give rise to the global obstruction.

**From Abstract Torsors to Concrete Cocycles** To understand the system’s state, we adopt the perspective of Čech cohomology. Let the protocol complex  $X$  be covered by a collection of open sets  $\{U_i\}_{i \in I}$ . Each  $U_i$  represents a “local view,” for instance, the subset of the execution space accessible to a

process  $i$ , or a small neighborhood in the protocol state space where a consistent partial state can be defined.

A local state assignment, or *local section*, is a function  $s_i : U_i \rightarrow A$  that assigns a state from the group  $A$  to each point in the view  $U_i$ . Consensus is achieved if we can find a consistent set of local sections  $\{s_i\}$  that arise from a single *global section*  $s : X \rightarrow A$ , meaning  $s_i = s|_{U_i}$  for all  $i$ .

When consensus fails, no such global section  $s$  exists. However, the system is not necessarily chaotic. On the intersections of local views,  $U_{ij} = U_i \cap U_j$ , two processes  $i$  and  $j$  can compare their states. In a non-trivial torsor, their local sections  $s_i$  and  $s_j$  will not agree. Instead, they are related by a **transition function**, which is an element  $g_{ij} \in \mathcal{A}(U_{ij})$ , the group of  $A$ -valued functions on the intersection. The relationship is given by:

$$s_i(x) = g_{ij}(x) \cdot s_j(x) \text{ for all } x \in U_{ij}$$

Here, the group multiplication  $\cdot$  represents the action of the transformation  $g_{ij}$  on the state  $s_j$ . This equation is fundamental: it states that the view of process  $i$  ( $s_i$ ) can be reconciled with the view of process  $j$  ( $s_j$ ) over their shared domain  $U_{ij}$  by applying a specific, well-defined transformation  $g_{ij}$ . This transformation  $g_{ij}$  is the **local state mismatch**. It is the concrete, algebraic measure of disagreement between views  $i$  and  $j$ .

The collection of all such pairwise mismatches,  $\{g_{ij}\}$ , is not arbitrary. For the overall system state to be locally consistent, these mismatches must satisfy a consistency condition on triple intersections  $U_{ijk} = U_i \cap U_j \cap U_k$ . Consider the relationship between  $s_i$ ,  $s_j$ , and  $s_k$ :

- $s_i = g_{ij} \cdot s_j$
- $s_j = g_{jk} \cdot s_k$

Substituting the second into the first gives:  $s_i = g_{ij} \cdot (g_{jk} \cdot s_k) = (g_{ij}g_{jk}) \cdot s_k$

However, we also have a direct relationship between  $s_i$  and  $s_k$  via the mismatch  $g_{ik}$ :  $s_i = g_{ik} \cdot s_k$

For these to be consistent, we must have:  $g_{ik}(x) = g_{ij}(x) \cdot g_{jk}(x)$  for all  $x \in U_{ijk}$

This is the **1-cocycle condition**. It asserts that the mismatch between  $i$  and  $k$  can be consistently computed by composing the mismatch from  $i$  to  $j$  with the mismatch from  $j$  to  $k$ . A collection of local mismatches  $\{g_{ij}\}$  satisfying this condition is called a **non-abelian 1-cocycle**. It is the raw data that defines an  $\mathcal{A}$ -torsor.

**Coboundaries: Distinguishing Apparent Mismatch from Intrinsic Disagreement** The existence of a non-zero cocycle  $\{g_{ij}\}$  does not, by itself, guarantee a fundamental obstruction. We must distinguish between mismatches that are merely artifacts of local convention and those that represent

an intrinsic, topological “twist” in the system’s state. This is the distinction between a **coboundary** and a non-trivial cocycle.

A 1-cocycle  $\{g_{ij}\}$  is a **1-coboundary** if there exists a set of local transformations  $\{h_i \in \mathcal{A}(U_i)\}$  such that the mismatch  $g_{ij}$  can be factored through them:

$$g_{ij} = h_i \cdot h_j^{-1}$$

**Interpretation:** A coboundary represents an *apparent* disagreement. The mismatch  $g_{ij}$  between processes  $i$  and  $j$  can be entirely explained by assuming that each process  $k$  has independently chosen a local “coordinate system” or “gauge,” represented by the transformation  $h_k$ . The disagreement  $g_{ij}$  is not a property of the underlying state, but simply the transformation required to translate from system  $j$ ’s coordinates to system  $i$ ’s.

In this case, consensus is achievable. We can define a global state by picking any local section, say  $s_j$ , and “undoing” its local gauge choice. Let’s define a candidate global section  $s$  by the formula  $s(x) = h_i(x)^{-1} \cdot s_i(x)$  for  $x \in U_i$ . We must check that this definition is consistent on overlaps. For  $x \in U_{ij}$ :

$$\begin{aligned} h_j(x)^{-1} \cdot s_j(x) &= h_j(x)^{-1} \cdot (g_{ij}(x)^{-1} \cdot s_i(x)) = h_j(x)^{-1} \cdot ((h_i(x)h_j(x))^{-1})^{-1} \cdot s_i(x) \\ &= h_j(x)^{-1} \cdot (h_j(x)h_i(x)^{-1} \cdot s_i(x)) = (h_j(x)^{-1}h_j(x)) \cdot h_i(x)^{-1} \cdot s_i(x) = h_i(x)^{-1} \cdot s_i(x) \end{aligned}$$

The definition is consistent, and a global state section exists. The torsor is trivial.

Conversely, if the cocycle  $\{g_{ij}\}$  is **not** a coboundary, no such collection of local gauge choices  $\{h_i\}$  exists. The mismatches are intrinsic to the system’s state fabric. They cannot be explained away by any local change of coordinates. This is a genuine, topological obstruction to consensus, corresponding to a non-trivial element in  $H^1(X, \mathcal{A})$ .

**Cocycles, Path-Dependence, and Holonomy** The cocycle condition  $g_{ik} = g_{ij}g_{jk}$  provides a powerful link to the concept of path-dependence and holonomy. In a simplicial model of the protocol complex, where processes are vertices and communication channels are edges, the cocycle  $\{g_{ij}\}$  assigns a group element to each directed edge  $(i, j)$ . The cocycle condition ensures that for any 2-simplex (triangle)  $(i, j, k)$ , the transformation along the path  $i \rightarrow j \rightarrow k$  is identical to the transformation along the direct path  $i \rightarrow k$ .

This local consistency allows us to calculate the net effect of traversing any path in the complex. Consider a loop  $\gamma$  in the protocol complex, representing a cycle of interactions, e.g.,  $p_1 \rightarrow p_2 \rightarrow \dots \rightarrow p_n \rightarrow p_1$ . If we start with a hypothetical state  $s$  at  $p_1$  and “transport” it along this loop, applying the successive mismatch transformations, the final state  $s'$  upon returning to  $p_1$  will be:

$$s' = (g_{p_1 p_n} \cdots g_{p_3 p_2} g_{p_2 p_1}) \cdot s$$

The total transformation  $g_\gamma = g_{p_1 p_n} \cdots g_{p_2 p_1}$  is the **holonomy** of the loop  $\gamma$ .

- If the cocycle  $\{g_{ij}\}$  is a coboundary ( $g_{ij} = h_i h_j^{-1}$ ), the holonomy of any loop is trivial (the identity element  $e \in A$ ):  $g_\gamma = (h_{p_1} h_{p_n}^{-1}) \dots (h_{p_3} h_{p_2}^{-1})(h_{p_2} h_{p_1}^{-1}) = h_{p_1} (h_{p_n}^{-1} \dots h_{p_2}^{-1} h_{p_2}) h_{p_1}^{-1} = h_{p_1} e h_{p_1}^{-1} = e$ . In this case, the outcome is independent of the path taken. No memory of the path is encoded in the final state.
- If the cocycle is **not** a coboundary, there will exist at least one loop  $\gamma$  for which the holonomy  $g_\gamma \neq e$ . This means that a process, or the system as a whole, undergoing a cycle of interactions can return to its initial configuration only to find that its state has been transformed. This is the essence of path-dependence. The non-trivial cocycle is the local data that generates this global, path-dependent behavior.

**Example: The  $SO(3)$  Mismatch Torsor** Imagine a system of autonomous drones  $\{P1, P2, P3\}$  attempting to agree on a common spatial orientation. The state space is the group of 3D rotations,  $A = SO(3)$ . Due to sensor noise, asynchronous communication, and imperfect calibration, they can only establish *relative* orientations.

1. **Local Mismatches:**

- P1 and P2 communicate and determine their relative orientation is  $g_{12} \in SO(3)$ . This is a concrete rotation matrix.
- P2 and P3 determine their relative orientation is  $g_{23}$ .
- P3 and P1 determine their relative orientation is  $g_{31}$ .

The collection  $\{g_{12}, g_{21}, g_{23}, g_{32}, g_{13}, g_{31}\}$  (where  $g_{ji} = g_{ij}^{-1}$ ) forms a 1-cocycle over the complete graph  $K_3$  representing the system.

2. **Holonomy as Obstruction:** Can they all align to a master orientation (e.g., North-Up-East)? This requires finding absolute orientations  $\{s_1, s_2, s_3\} \subset SO(3)$  such that  $s_1 = g_{12}s_2$ ,  $s_2 = g_{23}s_3$ , and  $s_3 = g_{31}s_1$ . Substituting these relations gives:  $s_1 = g_{12}s_2 = g_{12}(g_{23}s_3) = g_{12}g_{23}(g_{31}s_1) = (g_{12}g_{23}g_{31})s_1$

A solution exists if and only if the holonomy around the triangle  $(1, 2, 3)$  is trivial:  $H = g_{12}g_{23}g_{31} = I$

3. **Interpreting Non-Triviality:** If, due to measurement error accumulation,  $H \neq I$ , no set of absolute orientations  $\{s_i\}$  can satisfy the pairwise constraints. The system is in a state of intrinsic geometric frustration. The cocycle  $\{g_{ij}\}$  is non-trivial. It defines a non-trivial  $SO(3)$ -torsor over the complex. The non-identity holonomy  $H$  is the tangible proof of this non-triviality. The system cannot agree on an orientation because the very web of relative orientation measurements is geometrically inconsistent.



## Conclusion

By deconstructing the abstract notion of a torsor into its constituent 1-cocycle, we find a direct and powerful interpretation within the operational reality of distributed systems. A non-abelian 1-cocycle  $\{g_{ij}\}$  is precisely the web of local, pairwise state mismatches. The cocycle condition enforces local consistency on this web of disagreement, while the distinction between cocycles and coboundaries separates intrinsic, topological obstructions from mere differences in local convention.

The values  $g_{ij}$  are not abstract symbols; they are concrete transformations—rotations, permutations, or other non-commutative operations—that a process must apply to reconcile its view with a neighbor's. The failure to find local “gauge transformations”  $\{h_i\}$  to trivialize this cocycle is synonymous with the existence of non-trivial holonomy, the hallmark of path-dependent computation.

This interpretation grounds the algebraic machinery of cohomology in the tangible world of measurement, comparison, and reconciliation failure. It provides the crucial link between the global classification of impossibility ( $H^1(X, \mathcal{A})$ ) and the local dynamics of disagreement. Having understood how 1-cocycles capture mismatches between *states*, we are naturally led to a more profound question: what happens when there is an obstruction to even defining a consistent web of mismatches? This leads us to the realm of second cohomology,  $H^2(X, \mathcal{A})$ , and the theory of gerbes, which we will explore in the next chapter.

### Chapter 4.3: The Structure of Disagreement: How the Non-Abelian Group $A$ Encodes Failure Modes

The Structure of Disagreement: How the Non-Abelian Group  $A$  Encodes Failure Modes

In the preceding chapters, we established the central thesis that the impossibility of achieving global consensus in a distributed system can be precisely formulated as the non-triviality of the first non-abelian cohomology set,  $H^1(X, \mathcal{A})$ . An element  $[g] \in H^1(X, \mathcal{A})$ , representing an isomorphism class of  $\mathcal{A}$ -torsors over the protocol complex  $X$ , corresponds to a state of “consistent disagreement”—a global configuration where local views are consistent with their neighbors but cannot be reconciled into a single, global state. This establishes a powerful correspondence: consensus failure is a topological obstruction.

However, this conclusion, while profound, only tells us *if* consensus is possible. It does not elaborate on the *nature* of the failure. Distributed systems do not fail in a monolithic way; they exhibit a rich taxonomy of failure modes, from simple state drift and split-brain scenarios to insidious Byzantine behaviors. The true power of the cohomological framework lies in its ability to not only detect but also classify these failures. The key to this classification is not the topology of  $X$  alone, but the intricate algebraic structure of the non-abelian group  $A$  of state transformations. This chapter will dissect the group  $A$ , demonstrating how its

internal structure—its center, its conjugacy classes, its commutator subgroup—provides a veritable “genome” for the failure modes of the system.

### The Commutator Subgroup and the Genesis of Path-Dependence

The quintessential feature of a non-abelian group  $A$  is the existence of elements  $a, b \in A$  such that their commutator,  $[a, b] = aba^{-1}b^{-1}$ , is not the identity element  $e$ . This simple algebraic fact is the seed of all path-dependent failures in a distributed system.

Consider a process that observes the system by following a path in the protocol complex  $X$ . The state it perceives is transformed by the operations corresponding to the edges of the path. If two execution paths,  $\gamma_1$  and  $\gamma_2$ , share the same start and end points but are not homotopic, they represent two different causal histories for arriving at the same point in the computation. In an abelian world where all operations commute, the net transformation would be identical regardless of the path taken. The system’s state would be a function of the *set* of operations performed, not their *order*.

In the non-abelian case, the outcome depends crucially on the path. The discrepancy between the two paths is measured by the holonomy of the loop  $\gamma_1 \circ \gamma_2^{-1}$ . The holonomy is an element of  $A$  obtained by composing the cocycle values along the loop. If the loop is generated by a sequence of operations and their inverses, such as `apply(a)`, `apply(b)`, `apply(a-1)`, `apply(b-1)`, the resulting holonomy is precisely the commutator  $[a, b]$ .

**Interpretation:** The non-triviality of the commutator subgroup  $[A, A]$  directly corresponds to the existence of execution loops that induce a net change in state. A process traversing such a loop returns to its starting point having altered its local state, not because of new information, but because the fabric of the state space has been “twisted” by the non-commutative nature of the updates.

- **System-Level Manifestation:** This is the algebraic root of the **split-brain** problem in partitioned systems. Imagine a database with two operations, `op_A` and `op_B`, that do not commute. During a network partition, one partition receives `op_A` then `op_B`, while the other receives `op_B` then `op_A`. Upon healing, the system is faced with two states, `op_B(op_A(S))` and `op_A(op_B(S))`, which are irreconcilably different. The obstruction to merging these states is the non-identity commutator  $[op_B, op_A]$ . The system is trapped in a non-trivial  $A$ -torsor where the failure mode is characterized by the commutator subgroup of  $A$ .

**The Center  $Z(A)$ : Islands of Abelian Predictability** The center of a group,  $Z(A) = \{z \in A \mid \forall a \in A, za = az\}$ , consists of all elements that commute with every other element in  $A$ . It forms an abelian subgroup, representing a domain of path-independent behavior embedded within the larger non-abelian structure.

What is the significance of a cocycle  $g_{ij}$  whose values lie entirely within  $Z(A)$ ?

Recall the 1-cocycle condition for a sheaf of groups  $\mathcal{A}$  with a group action (typically conjugation by local transition functions, which we suppress for clarity here but is implicit in the sheaf formalism):  $g_{ik} = g_{ij} \cdot (i \cdot g_{jk})$

If  $g_{jk} \in Z(A)$ , the action of any process  $i$  on it is trivial, i.e.,  $i \cdot g_{jk} = g_{jk}$ . The condition simplifies to the abelian form:  $g_{ik} = g_{ij} \cdot g_{jk}$

This implies that any disagreement encoded by cocycles valued in the center is path-independent. The holonomy around any loop will be the identity. The failure to achieve consensus is not due to the order of operations but to a persistent, global “offset” or “mismatch.”

**Interpretation:** Failures described by  $H^1(X, Z(A))$  represent a simpler class of disagreement. The system state is still fractured, but the disagreement is stable and predictable.

- **System-Level Manifestation:** Consider a distributed system where each node maintains a version vector, but due to a configuration error, each node has a different, fixed offset from the “true” version. Node 1 is always ahead by +2, Node 2 by +5. There is no global consensus on the version, so  $H^1(X, \mathbb{Z})$  is non-trivial (here,  $A = \mathbb{Z}$  acts as a model for the versioning logic). This is a “benign” disagreement; the relative difference between any two nodes (+3 between Node 1 and 2) is constant and globally agreed upon. The failure is an element of  $H^1(X, Z(A))$ , representing a global misalignment without the chaotic dynamics of path-dependence.

**Conjugacy Classes: The Observable Spectra of Failure** A local observer in a distributed system, say process  $i$ , can only measure the state of another process  $j$  relative to its own local reference frame. In our formalism, this choice of a local reference state corresponds to choosing a trivialization of the  $A$ -torsor over the open set  $U_i$  in the protocol complex. If we change this reference frame (i.e., choose a different trivialization), the cocycle  $g_{ij}$  that represents the mismatch between  $i$  and  $j$  transforms. Specifically, a change of trivialization by elements  $a_k \in A(U_k)$  transforms the cocycle  $g$  into a cohomologous cocycle  $g'$  via:  $g'_{ij} = a_i^{-1} \cdot g_{ij} \cdot (i \cdot a_j)$

If the group action is conjugation, this becomes  $g'_{ij} = a_i^{-1} \cdot g_{ij} \cdot a_j$ . This means that while the specific element  $g_{ij}$  is observer-dependent, its **conjugacy class** is an invariant. All possible measurements of the mismatch between regions  $U_i$  and  $U_j$  will yield elements within the same conjugacy class in  $A$ .

**Interpretation:** The set of conjugacy classes of  $A$  provides a catalog of the fundamentally distinct, observable types of disagreement that can arise in the system.

- **The Identity Class  $\{e\}$ :** Represents agreement or a trivial torsor. Any apparent mismatch can be eliminated by a suitable choice of reference frames. Consensus is possible.

- **Singleton Classes (Elements of  $Z(A)$ ):** As discussed above, these correspond to path-independent, stable disagreements. The observed mismatch is the same regardless of the observer’s state.
- **Large Conjugacy Classes:** These represent complex, observer-dependent failures. The nature of the discrepancy measured between two processes depends critically on the state of the measuring process. These are characteristic of highly non-abelian groups and correspond to the most volatile and unpredictable failure modes.

This perspective offers a powerful way to classify Byzantine faults. A Byzantine process can be seen as one whose state transformations correspond to elements from a large conjugacy class with a small centralizer. Its behavior appears inconsistent and chaotic to most other processes because very few operations “commute” with its state. Conversely, a more structured, “symmetric” failure (like a whole partition being off by a fixed rotation in quaternion-based state space) would correspond to an element with a larger centralizer.

**Group Extensions and Hierarchies of Obstruction** The analysis can be deepened by considering the subgroup structure of  $A$ . A short exact sequence of groups,  $1 \rightarrow N \rightarrow A \rightarrow Q \rightarrow 1$ , where  $N$  is a normal subgroup of  $A$  and  $Q \cong A/N$  is the quotient group, provides a way to decompose failures into a hierarchy. This sequence induces a long exact sequence in cohomology:

$$\dots \rightarrow H^1(X, N) \xrightarrow{\alpha} H^1(X, A) \xrightarrow{\beta} H^1(X, Q) \xrightarrow{\delta} H^2(X, N) \rightarrow \dots$$

This sequence is a diagnostic tool for pinpointing the source of disagreement:

1. **Map  $\alpha : H^1(X, N) \rightarrow H^1(X, A)$ :** This map takes failures of the “internal” type  $N$  and embeds them into the full spectrum of  $A$ -failures. If  $N = Z(A)$ , these are the path-independent failures.
2. **Map  $\beta : H^1(X, A) \rightarrow H^1(X, Q)$ :** This map takes a complex  $A$ -failure and projects it onto a “coarser”  $Q$ -failure. If an element  $[g] \in H^1(X, A)$  maps to the trivial element in  $H^1(X, Q)$ , it means the failure, when viewed through the coarse lens of the quotient group  $Q$ , disappears. The disagreement is therefore “of type  $N$ ”. The problem lies entirely within the algebraic structure of the subgroup  $N$ . For example, if  $A$  is the group of affine transformations  $x \mapsto ax + b$  and  $N$  is the subgroup of translations ( $a = 1$ ), then a failure that becomes trivial in the quotient group  $Q$  (the linear maps  $x \mapsto ax$ ) is purely a translational mismatch, with no rotational or scaling component.
3. **The Connecting Homomorphism  $\delta : H^1(X, Q) \rightarrow H^2(X, N)$ :** This is the most subtle and powerful part of the sequence. It addresses the question: given a “coarse” consensus failure of type  $Q$ , can we “lift” it or “refine” it to a full-fledged  $A$ -failure? The map  $\delta$  provides the answer. A  $Q$ -torsor (an element of  $H^1(X, Q)$ ) can be lifted to an  $A$ -torsor if and only if its image under  $\delta$ , which is an element of  $H^2(X, N)$ , is trivial.

**Interpretation:**  $\delta$  represents an obstruction of a higher order. It signifies that even if a simplified version of the problem (in the quotient group  $Q$ ) appears to have a certain type of disagreement, the internal structure of  $N$  may introduce a more complex, second-order obstruction that prevents a consistent global state. This obstruction is not a torsor but a **gerbe**, the geometric object classified by  $H^2$ . This reveals that the problem of resolving consensus is hierarchical: solving the “coarse” part of the disagreement might be obstructed by a more fundamental inconsistency at a finer level, an inconsistency not of states, but of the protocols for reconciling states themselves.

### Conclusion: The Failure Genome

The non-abelian group  $A$  is far more than a label for the set of states. It is the repository of the system’s potential pathologies. By examining its algebraic features, we can create a detailed taxonomy of consensus failures:

- **Commutators** encode the potential for **path-dependent failures** and **split-brain scenarios**.
- The **center**  $Z(A)$  encodes **path-independent, stable disagreements**.
- **Conjugacy classes** classify the **observable types of failure**, distinguishing stable, symmetric faults from chaotic, Byzantine-like ones.
- **Subgroup and quotient structures** reveal a **hierarchy of failures**, allowing us to determine if a disagreement is of a specific “type” and to identify higher-order obstructions (gerbes) to resolving simpler, coarse-grained problems.

In this light, designing a robust distributed system is analogous to genetic engineering. The choice of state representation and the set of allowed operations define the group  $A$ . Choosing operations that commute (an abelian  $A$ ) eliminates path-dependence entirely. If non-commutativity is unavoidable, designing  $A$  to have a large center or a simple commutator structure can constrain the system to more manageable failure modes. The cohomological framework, powered by the algebraic structure of  $A$ , thus transforms from a purely descriptive tool into a predictive and prescriptive one, providing a powerful mathematical language to reason about the very structure of disagreement.

### Chapter 4.4: Beyond State Disagreement: Gerbes and $H^2$ as Obstructions to Protocol Coherence

Beyond State Disagreement: Gerbes and  $H^2$  as Obstructions to Protocol Coherence

In the preceding chapters, we have developed a powerful correspondence between the failure to achieve consensus in a distributed system and the non-triviality of the first non-abelian cohomology group,  $H^1(X, \mathcal{A})$ . The topological space  $X$  represents the space of all possible protocol executions, and the sheaf of non-abelian groups  $\mathcal{A}$  represents the local symmetries of the state space. A non-trivial class in  $H^1(X, \mathcal{A})$  corresponds to an  $\mathcal{A}$ -torsor over  $X$ —a “twisted”

version of the state space that admits no global section. Such a section would be a globally agreed-upon state, and its absence is the formal definition of consensus impossibility. This “torsor of disagreement” provides a complete, geometric picture of how local views fail to cohere into a consistent global state.

This framework, however, presumes a fundamental level of consistency. The torsor model, while describing disagreement, implies that the *rules of disagreement* are themselves globally coherent. The 1-cocycle  $g_{ij}$  that relates the state  $s_j$  in process  $j$ ’s view to the state  $s_i$  in process  $i$ ’s view ( $s_i = g_{ij} \cdot s_j$ ) must obey the consistency condition  $g_{ik} = g_{ij} \cdot g_{jk}$  on any triple intersection of local views. This condition ensures that the web of pairwise disagreements is self-consistent.

But what if this underlying consistency breaks down? What if the very protocols for reconciling local views are themselves incoherent? This points to a deeper, more subtle class of failures—not an obstruction to a global *state*, but an obstruction to a global *protocol*. To analyze such phenomena, we must ascend the cohomological ladder from  $H^1$  to  $H^2$ , and from the geometry of torsors to the more intricate world of gerbes.

### The Geometry of Higher Obstructions: Gerbes

Just as a torsor is an object that is locally isomorphic to a group  $\mathcal{A}$  but may lack a global identity element, a **gerbe** is a higher categorical object that is locally equivalent to a torsor. It can be intuitively understood as a “stack of torsors.” The local data of a gerbe over a cover  $\{U_i\}$  of the execution space  $X$  consists of: 1. On each patch  $U_i$ , an  $\mathcal{A}$ -torsor  $T_i$ . 2. On each intersection  $U_{ij} = U_i \cap U_j$ , an isomorphism of torsors  $\phi_{ij} : T_j|_{U_{ij}} \rightarrow T_i|_{U_{ij}}$ .

If we were to stop here, these isomorphisms could be used to glue the local torsors  $T_i$  into a single global torsor. The obstruction to this gluing process is precisely what defines a gerbe. The gluing data  $\phi_{ij}$  must itself satisfy a consistency condition on triple intersections  $U_{ijk} = U_i \cap U_j \cap U_k$ . Specifically, the composition of isomorphisms  $\phi_{ij} \circ \phi_{jk}$  going from  $T_k$  to  $T_i$  via  $T_j$  must be compared to the direct isomorphism  $\phi_{ik}$ .

In the case of a trivial gerbe (which corresponds to a global torsor), this composition is exact:  $\phi_{ik} = \phi_{ij} \circ \phi_{jk}$ . However, for a non-trivial gerbe, this equality fails. The failure is not arbitrary; it is captured by an automorphism of the torsor  $T_k$ . Since automorphisms of an  $\mathcal{A}$ -torsor are given by the action of the group  $\mathcal{A}$  itself, this failure is measured by an element of  $\mathcal{A}$ . This “failure of gluing the gluing data” is the geometric essence of a gerbe and the hallmark of a non-trivial class in  $H^2(X, \mathcal{A})$ .

This situation is analogous to the construction of vector bundles in differential geometry. A vector bundle is defined by local product structures  $U_i \times \mathbb{R}^n$  and transition functions  $g_{ij} : U_{ij} \rightarrow GL(n, \mathbb{R})$  that satisfy the 1-cocycle condition  $g_{ik} = g_{ij}g_{jk}$ . This is an  $H^1$  phenomenon. A gerbe is analogous to having local vector bundles  $E_i$  on each  $U_i$ , isomorphisms  $\phi_{ij} : E_j \rightarrow E_i$  on overlaps, where

the composition of these isomorphisms around a triple overlap fails to be the identity, returning a non-trivial bundle automorphism.

### The Algebraic Signature: Non-Abelian 2-Cocycles

The geometric picture of a gerbe is mirrored by the algebraic structure of non-abelian 2-cocycles. Let us formalize the breakdown of the 1-cocycle condition.

Recall that a 1-cocycle  $\{g_{ij} \in \mathcal{A}(U_{ij})\}$  represents an  $\mathcal{A}$ -torsor and satisfies the relation  $g_{ij}g_{jk}g_{ki} = e$  on every triple overlap  $U_{ijk}$ . This allows for the consistent definition of a global object from local data.

Now, consider a situation where we can define local transition functions  $g_{ij}$ , but they fail the 1-cocycle condition. The failure itself becomes the object of study. We define a 2-cochain  $h_{ijk}$  by the relation:

$$h_{ijk} = g_{ij} \cdot g_{jk} \cdot g_{ki}$$

This element  $h_{ijk} \in \mathcal{A}(U_{ijk})$  measures the “error” or “curvature” in the attempt to patch local data. If all  $h_{ijk}$  are the identity element  $e$ , we recover the 1-cocycle condition, and the obstruction lies in  $H^1$ . If  $h_{ijk}$  is non-trivial, it represents a potential  $H^2$  obstruction.

For  $h_{ijk}$  to define a valid cohomology class, it must itself satisfy a cocycle condition. By considering a quadruple overlap  $U_{ijkl}$ , one can show that the  $g_{ij}$  relations imply a condition on the  $h_{ijk}$  elements. This is the **2-cocycle condition**:

$$g_{ij}h_{jkl}g_{ij}^{-1} \cdot h_{ikl}^{-1} \cdot h_{ijl} \cdot h_{ijk}^{-1} = e$$

This expression simplifies considerably if the coefficients  $h_{ijk}$  take values in the center of the group,  $Z(\mathcal{A})$ . If  $h_{ijk} \in Z(\mathcal{A})$ , the condition becomes the standard abelian 2-cocycle condition:

$$h_{jkl} - h_{ikl} + h_{ijl} - h_{ijk} = 0 \quad (\text{written additively})$$

$$h_{jkl} \cdot h_{ikl}^{-1} \cdot h_{ijl} \cdot h_{ijk}^{-1} = e \quad (\text{written multiplicatively})$$

A non-trivial solution to this equation that cannot be written as the boundary of a 1-cochain (i.e.,  $h_{ijk} \neq f_{jk}f_{ik}^{-1}f_{ij}$  for some  $f_{ij}$ ) defines a non-zero class in  $H^2(X, Z(\mathcal{A}))$ . This class is the algebraic manifestation of a non-trivial gerbe.

### Interpreting $H^2$ in Distributed Systems

Translating this abstract framework into the language of distributed consensus reveals its profound implications. While  $H^1$  describes a failure to agree on a state,  $H^2$  describes a failure to agree on the *protocol for managing disagreement*.

- **Local Protocols of Reconciliation ( $T_i$ ):** Imagine a large, partitioned system (e.g., across multiple geographic data centers). Within a single

partition  $U_i$  (a subset of processes and their execution paths), the system may run a protocol that successfully establishes a coherent model of disagreement—a local torsor  $T_i$ . For example, all nodes within the partition might agree on a consistent set of transformations that relate their local states, even if they cannot agree on one state.

- **Inconsistent Protocol Isomorphisms** ( $g_{ij}$ ): When two partitions  $U_i$  and  $U_j$  interact, they must translate between their local disagreement models. This translation is the isomorphism of torsors, represented by the 1-cochain  $g_{ij}$ . It acts as a “meta-protocol” for reconciling the reconciliation protocols. For instance,  $g_{ij}$  could be a set of rules for mapping the state transformations valid in partition  $j$  to those valid in partition  $i$ .
- **The Gerbe as a Meta-Obstruction** ( $h_{ijk} \neq e$ ): The  $H^2$  obstruction emerges when we consider three interacting partitions,  $U_i, U_j, U_k$ . A non-trivial 2-cocycle  $h_{ijk}$  means that the meta-protocols do not compose correctly. Translating the disagreement model from  $k \rightarrow j$  (via  $g_{jk}$ ) and then from  $j \rightarrow i$  (via  $g_{ij}$ ) yields a different result than translating directly from  $k \rightarrow i$  (via  $g_{ik}$ ). The discrepancy is a fundamental, non-removable transformation  $h_{ijk}$ .

This signifies a catastrophic failure of **protocol coherence**. There is no globally consistent way to manage local disagreements. Any attempt to define a system-wide set of rules for how different processes should interpret each other’s states is doomed to fail with path-dependent inconsistencies. This is an obstruction to **meta-consensus**: the system cannot even agree on a common framework for disagreeing. Such a failure mode is far more subtle than a simple fork or split-brain scenario. It could manifest as a system that appears locally healthy within different segments, but whose global state integration is fundamentally impossible, leading to inexplicable data corruption or logical inconsistencies that cannot be traced to a single faulty component or message.

### Coefficient Structures: From Groups to 2-Groups

The rigorous definition of non-abelian  $H^2$  requires a richer algebraic structure for the coefficients than a simple group. The appropriate object is a **2-group**, which can be modeled algebraically by a **crossed module**. A crossed module is a pair of groups  $(H, G)$  equipped with an action of  $G$  on  $H$  and a group homomorphism  $\partial : H \rightarrow G$  that satisfies two axioms: 1.  $\partial(g \cdot h) = g\partial(h)g^{-1}$  for  $g \in G, h \in H$ . 2.  $(\partial h_1) \cdot h_2 = h_1 h_2 h_1^{-1}$  for  $h_1, h_2 \in H$ .

In the context of gerbes,  $G$  can be thought of as the group of automorphisms of an object (our original  $\mathcal{A}$ ), and  $H$  as the group of isomorphisms between those automorphisms. The map  $\partial$  connects these two levels.

For distributed systems, this implies that a complete model of higher-order failures requires us to specify not only the group  $\mathcal{A}$  of state transformations but also a second group describing the “transformations between transformations.”



This algebraic step-up reflects the computational complexity of the failure mode. An  $H^1$  failure is about states and their transformations ( $\mathcal{A}$ ). An  $H^2$  failure is about the protocols for managing those transformations, requiring an algebraic object that captures the relationships *between* the transformations themselves. This could model, for example, Byzantine failures where an adversary does not simply alter a state value but alters the interpretation of a state-update operation, thereby corrupting the logic of the protocol itself in a way that is inconsistent across different observer groups.

### Conclusion: Gerbes, Meta-Consensus, and Protocol Composition

The introduction of second non-abelian cohomology and gerbes elevates our topological framework from a theory of state disagreement to a theory of protocol incoherence.  $H^1(X, \mathcal{A})$  classifies the obstructions to achieving a global state, materialized as the torsor of disagreement. It formalizes the impossibility of consensus.  $H^2(X, \mathcal{A})$  classifies the obstructions to achieving a globally consistent protocol for managing local disagreements, materialized as a gerbe. It formalizes the impossibility of meta-consensus.

This higher-order obstruction has direct relevance for the design and analysis of complex, large-scale distributed systems. It provides a formal language to describe failures in **protocol composition**, where protocols that are individually correct and fault-tolerant may yield a globally incoherent system when combined. The non-trivial 2-cocycle  $h_{ijk}$  represents the fundamental incompatibility born from their composition.

Furthermore, the gerbe framework offers a new lens for classifying Byzantine faults. Beyond simple data corruption, it can model sophisticated attacks that manipulate the semantic interpretation of the protocol rules, creating systemic logical paradoxes. By identifying the conditions that lead to non-trivial  $H^2$  classes, we can develop principles for designing protocols that are robust not only against state-level divergence but also against these deeper, coherence-level attacks. The path from  $H^1$  to  $H^2$  is a journey from the algebra of disagreement to the algebra of a disagreement about disagreement itself—a crucial step in understanding the most subtle and challenging failure modes in distributed computing.

### Chapter 4.5: A Cohomological Classification of Faults: From Benign Asynchrony to Byzantine Torsors

A Cohomological Classification of Faults: From Benign Asynchrony to Byzantine Torsors

In the preceding chapters, we have established the central thesis that the impossibility of distributed consensus can be framed as a cohomological obstruction. Specifically, the non-triviality of the first non-abelian cohomology set,  $H^1(X, \mathcal{A})$ , where  $X$  is the protocol complex and  $\mathcal{A}$  is a sheaf of (generally non-abelian) groups representing local state transformations, formalizes the absence

of a global agreement. This chapter builds upon this foundation to propose a novel classification of fault models in distributed computing. We argue that the traditional taxonomy of faults—crash-stop, omission, Byzantine—while descriptively useful, can be subsumed and refined by a more fundamental, algebraic classification. The nature of a fault, we contend, is encoded in the algebraic structure of the coefficient group  $\mathcal{A}$  and the properties of the resulting cocycles and torsors. This perspective allows us to distinguish between “benign” disagreements arising from asynchrony and “malicious” disagreements arising from Byzantine behavior, not by their operational cause, but by the mathematical structure of the resulting inconsistency.

**The Baseline: Benign Asynchrony and Abelian Disagreement** Let us begin with the simplest non-trivial scenario: a set of correct processes attempting to reach consensus in a purely asynchronous system, as in the classic FLP impossibility setting. The only “fault” is the unbounded message delay. Suppose the processes aim to agree on a single integer value from a proposed set. The state updates are simple additions or assignments, which are fundamentally commutative operations. For example, if a process’s state is a version number, applying update  $\mathbf{u}_1$  then  $\mathbf{u}_2$  results in the same state as applying  $\mathbf{u}_2$  then  $\mathbf{u}_1$ .

In this context, the appropriate coefficient group  $A$  is **abelian**. For instance, if processes are trying to agree on a version vector, the differences between their local versions can be modeled by the group of integers,  $\mathbb{Z}$ . A 1-cocycle  $\phi \in Z^1(X, \mathbb{Z})$  is a function that assigns an integer  $\phi_{ij}$  to each directed edge  $(i, j)$  in the protocol complex, representing the “version drift” or “time lag” observed by process  $i$  relative to process  $j$ . The cocycle condition, for an abelian group, simplifies to:

$$\phi_{ik} = \phi_{ij} + \phi_{jk}$$

for any 2-simplex  $(i, j, k) \in X$ . This equation expresses a simple, consistent transitivity of differences. The difference between  $i$  and  $k$  is the sum of the difference between  $i$  and  $j$  and the difference between  $j$  and  $k$ .

Consensus is possible if and only if this cocycle is a coboundary; that is, if there exists a global state function  $s : V(X) \rightarrow \mathbb{Z}$  (a 0-cochain) such that  $\phi_{ij} = s_j - s_i$ . In this case,  $H^1(X, \mathbb{Z}) = 0$ . The FLP impossibility result demonstrates that for a typical asynchronous protocol complex, its topology is non-trivial, leading to  $H^1(X, \mathbb{Z}) \neq 0$ .

The key insight is that the resulting disagreement, while preventing consensus, is **benign** and **quantitative**. The obstruction is a non-trivial  $\mathbb{Z}$ -torsor, which can be visualized as a principal bundle over  $X$  with fiber  $\mathbb{Z}$ . Geometrically, this is akin to a “spiral staircase” or helicoid. While no single floor (global state) can be chosen consistently, the *relative height* between any two points is well-defined and path-independent. The disagreement between any two processes is a simple numerical offset, and all processes agree on the magnitude of this offset. The structure of the disagreement is constrained by the commutativity of  $\mathbb{Z}$ .

**Crash-Stop and Omission Faults: Puncturing the Topology** Next, consider crash-stop faults, where a process ceases all activity permanently. From a cohomological viewpoint, a crash fault can be modeled in two primary ways:

1. **Topological Modification:** The most direct model is to state that a crashed process  $p_c$  and any communication involving it are absent from the execution. This means that any simplex  $\sigma = (p_0, \dots, p_k)$  in the protocol complex  $X$  where one of the vertices is  $p_c$  is never formed after the crash. The resulting execution complex  $X_{\text{faulty}}$  is a subcomplex of the potential full complex. The consensus problem is then reformulated on this punctured space. The impossibility of consensus is then determined by the cohomology of this modified space,  $H^1(X_{\text{faulty}}, A)$ . A crash may or may not resolve the impossibility; for example, if a crash disconnects the complex in a way that simplifies its topology, it could paradoxically make consensus easier for the remaining participants.
2. **Algebraic Trivialization:** Alternatively, we can keep the complex  $X$  intact but modify the coefficient sheaf  $\mathcal{A}$ . A crashed process can be considered to be in a fixed, absorbent “crashed” state. The transformations associated with it become trivial. The cocycle values  $g_{ij}$  on any edge  $(i, j)$  involving the crashed process  $p_c$  would be the identity element  $e \in A$ . This effectively “flattens” the torsor in the dimensions corresponding to the crashed process, but the non-trivial cohomological structure may persist among the remaining active processes.

Omission faults, where messages are intermittently dropped, can be seen as a dynamic version of this. The topology of the protocol complex flickers, with certain edges and simplices appearing and disappearing over time. An obstruction to consensus would mean that a non-trivial cocycle exists on a persistent topological “hole” that is not eliminated by the transient communication failures. In both crash and omission faults, the fundamental nature of state transformations remains unchanged; the group  $A$  can still be abelian if the underlying operations are commutative. The fault is primarily topological (affecting  $X$ ) rather than algebraic (affecting  $A$ ).

**The Byzantine Torsor: The Algebra of Malice** Byzantine faults represent the most challenging failure model, where a faulty process can behave arbitrarily and maliciously, sending conflicting information to different peers. This is precisely where the transition from abelian to **non-abelian** cohomology becomes essential. The essence of a Byzantine fault is not just providing incorrect data, but creating a web of inconsistencies that destroys the very notion of a coherent, shared reality.

Let the group  $A$  represent not just a value, but the set of allowed symmetries or transformations of the entire state space. For instance, if processes are trying to agree on an ordered list of three candidates  $\{C_1, C_2, C_3\}$ , the state transformations could be modeled by the permutation group  $A = S_3$ . This group is

famously non-abelian.

A 1-cocycle in the non-abelian setting is a function  $g$  that assigns a group element  $g_{ij} \in A$  to each edge  $(i, j)$  of the protocol complex, subject to the twisted cocycle condition on each 2-simplex  $(i, j, k)$ :

$$g_{ik} = g_{ij} \cdot j(g_{jk})$$

Here,  $j(g_{jk})$  denotes the action of being in process  $j$ 's frame of reference on the transformation  $g_{jk}$ . For simplicity, if the action is trivial (i.e., the group structure is globally agreed upon), this becomes  $g_{ik} = g_{ij}g_{jk}$ .

Now, consider how a Byzantine process  $b$  generates a non-trivial cocycle. - Suppose processes  $p_1$  and  $p_2$  are correct, and  $b$  is Byzantine. -  $b$  tells  $p_1$  that the correct ordering is  $(C_1, C_2, C_3)$ . -  $b$  tells  $p_2$  that the correct ordering is  $(C_2, C_1, C_3)$ . -  $p_1$  and  $p_2$  then communicate with each other.

Let's trace the transformations. From  $p_1$ 's perspective, to align its view with  $b$ 's stated view, it needs no transformation, so  $g_{p_1,b} = e$  (the identity permutation). From  $p_2$ 's perspective, to align its (initially identical) view with what  $b$  told it, it must apply the transposition  $\tau = (C_1 C_2)$ . So,  $g_{p_2,b} = \tau$ . When  $p_1$  and  $p_2$  communicate, they discover this discrepancy. The transformation required to reconcile  $p_1$ 's view with  $p_2$ 's view, as mediated through  $b$ , is path-dependent.

- Path 1:  $p_1 \rightarrow p_2$  (direct communication). Let's assume they find no direct discrepancy, so  $g_{p_1,p_2} = e$ .
- Path 2:  $p_1 \rightarrow b \rightarrow p_2$ . The composite transformation is  $g_{p_1,b} \cdot g_{b,p_2}$ . Note that  $g_{b,p_2} = g_{p_2,b}^{-1} = \tau^{-1} = \tau$ . So the transformation is  $e \cdot \tau = \tau$ .

The disagreement between the two paths ( $e$  vs.  $\tau$ ) indicates that the local data cannot be consistently patched together. This mismatch,  $\delta g(p_1, p_2, b) = g_{p_1,p_2} \cdot g_{p_2,b} \cdot g_{b,p_1} \neq e$ , signifies that  $g$  is not a coboundary. The set of all such locally consistent but globally irreconcilable views forms a non-trivial **A-torsor**.

This "Byzantine torsor" is fundamentally different from the benign abelian torsor of asynchrony. In an  $A$ -torsor for non-abelian  $A$ : 1. **There is no global section:** No consensus is possible. 2. **There is no canonical difference:** The "disagreement" between process  $i$ 's state and process  $j$ 's state is not a single element of  $A$ . It is an entire coset of a stabilizer subgroup. Worse, the transformation depends on the path of comparison. The discrepancy between  $i$  and  $j$  measured via  $k_1$  ( $g_{ik_1}g_{k_1j}$ ) may be a different group element from the discrepancy measured via  $k_2$  ( $g_{ik_2}g_{k_2j}$ ) due to non-commutativity.

This path-dependent, qualitative ambiguity is the mathematical signature of Byzantine failure. The lie is not a simple error; it is a structural sabotage of the system's geometry of information.

**A Cohomological Fault Classification** We can now summarize this classification in a hierarchical table:

Fault Type	Nature of Disagreement	Coefficient Group $A$	Cohomology	Geometric Interpretation
<b>Benign Asynchrony</b>	Quantitative, Path-Independent	<b>Abelian</b> (e.g., $\mathbb{Z}, \mathbb{R}$ )	$H^1(X, A) \neq 0$	A principal bundle (e.g., a helicoid). Disagreement is a consistent, global “twist”.
<b>Crash/Omission</b>	Topological Absence	Any group (Abelian or not)	$H^1(X_{\text{faulty}}, A) \neq 0$	Disagreement stems from holes in the communication graph/protocol complex itself.
<b>Byzantine</b>	Qualitative, Path-Dependent	<b>Non-Abelian</b> (e.g., $S_n, GL(n, F)$ )	Non-trivial class in $H^1(X, A)$	A non-trivial $A$ -torsor. The geometry of state space is twisted and path-dependent.

This framework demonstrates that what we call a Byzantine fault is not merely a different behavior but corresponds to a fundamental shift in the algebraic structure of state inconsistency—from an abelian group to a non-abelian one.

**Higher Obstructions: Failures in Fault Detection** This classification points toward higher-order phenomena. If  $H^1(X, A)$  measures the obstruction to consensus on a *state*, then the second cohomology group,  $H^2(X, A)$ , can be interpreted as an obstruction to a “meta-consensus”: a consensus on the *protocol itself*, or on the nature of the disagreement. A non-trivial element in  $H^2(X, A)$ , known as a gerbe, represents an obstruction to patching local torsors together.

In the context of faults, this could manifest as an impossibility of achieving a consistent, global fault diagnosis. Imagine a fault-detection protocol running on top of the base protocol. Different cliques of processes might successfully use an  $H^1$ -level test to identify a local set of Byzantine actors. However, their conclusions might be fundamentally irreconcilable at a global scale. For instance, clique  $C_1$  concludes node  $b_1$  is faulty based on evidence  $E_1$ , and clique  $C_2$  concludes  $b_2$  is faulty based on evidence  $E_2$ . A gerbe-like obstruction would

emerge if, upon combining their information, they find that  $E_1$  and  $E_2$  are mutually contradictory, leading to a paradox where no globally consistent set of faulty nodes can be identified. This represents a failure in the logic of the fault-detection protocol itself, an obstruction to forming a coherent “torsor of disagreements.”

In conclusion, the language of non-abelian cohomology provides a powerful, unifying lens through which to view distributed faults. It elevates the discussion from a catalog of behaviors to a structural analysis of disagreement. By associating benign asynchrony with abelian cohomology and malicious Byzantine behavior with non-abelian torsors, we not only gain a deeper understanding of why these systems fail but also acquire a precise mathematical toolkit to measure and classify the very fabric of their inconsistency.

## Part 5: Applications in Protocol Analysis and Fault-Tolerance

### Chapter 5.1: Cohomological Analysis of Paxos: Proving Consensus through a Trivial Disagreement Torsor

Cohomological Analysis of Paxos: Proving Consensus through a Trivial Disagreement Torsor

The preceding chapters have established a powerful, if abstract, framework: the impossibility of consensus in a distributed system can be precisely formulated as the existence of a non-trivial class in the first nonabelian cohomology set,  $H^1(X, A)$ . This class corresponds to a “disagreement torsor”—a global state of ambiguity that cannot be resolved into a single, consistent choice because of the topological structure of the protocol’s execution space,  $X$ . In this framework, a protocol succeeds not by chance, but by actively implementing mechanisms that guarantee the triviality of this cohomology group.

This chapter applies this analytical lens to one of the most foundational protocols in fault-tolerant distributed computing: Paxos. We will demonstrate that the safety of Paxos is equivalent to the statement that its procedural rules—specifically, its two-phase structure, totally ordered proposal numbers, and quorum intersection requirement—are precisely the algebraic machinery required to ensure that any potential 1-cocycle of disagreement is, in fact, a 1-coboundary. In doing so, we prove that for any execution of Paxos, the disagreement torsor is trivial, which is the cohomological statement of consensus.

**Establishing the Cohomological Setting for Paxos** To begin our analysis, we must first translate the components of the Paxos protocol into the language of our algebraic-topological model.

**1. The Protocol Complex,  $X$ :** The space  $X$  is the simplicial complex (or more accurately, the directed space) of all possible partial executions of the Paxos protocol. \* **Vertices (0-simplices):** A vertex represents the local state

of a single process (Proposer, Acceptor, or Learner) at a given logical time. For an Acceptor  $p$ , a state can be denoted by a tuple  $\mathbf{s}_p = (\text{max\_prepare}, \text{accepted\_n}, \text{accepted\_v})$ , representing the highest proposal number it has promised to, and the proposal number and value it has accepted, if any. For a Proposer, the state includes its current proposal number  $n$  and proposed value  $v$ .

\* **Edges (1-simplices):** An edge  $(\mathbf{s}_i, \mathbf{s}_j)$  represents a possible transition between two global states  $\mathbf{s}_i$  and  $\mathbf{s}_j$ , typically caused by a single event like a message send/receive or an internal state change of one process.

\* **Higher-dimensional Simplices:** A  $k$ -simplex  $(\mathbf{s}_0, \dots, \mathbf{s}_k)$  represents a set of  $k+1$  global states that are mutually consistent, meaning they could coexist in a single global view of the system. The topology of  $X$  is defined by the asynchronous nature of the network. The absence of a global clock and the possibility of message delays or process failures create a complex, non-trivial topology with potential “holes” or ambiguous paths.

**2. The Coefficient Group of Disagreement,  $A$ :** The group  $A$  models the nature of the disagreement we wish to measure. In the context of single-value consensus, the fundamental disagreement is over the chosen value. Let  $\mathcal{V}$  be the set of all possible values that can be proposed. A disagreement can be seen as a transformation from one proposed value to another.

\* **Simple Case (Binary Consensus):** If  $\mathcal{V} = \{0, 1\}$ , the disagreement can be modeled by the group  $A = \mathbb{Z}_2 = \{id, flip\}$ , where  $flip$  represents the transformation from 0 to 1 and vice-versa.

\* **General Case:** For a larger set  $\mathcal{V}$ , we could let  $A$  be the symmetric group  $Sym(\mathcal{V})$ , where each element represents a permutation of the possible outcomes. An element  $g \in A$  mapping  $v_1 \rightarrow v_2$  represents the “disagreement” between a process that believes the outcome is  $v_1$  and one that believes it is  $v_2$ . The non-abelian nature of  $Sym(\mathcal{V})$  for  $|\mathcal{V}| > 2$  correctly captures the fact that resolving a disagreement between  $v_1$  and  $v_2$ , and then between  $v_2$  and  $v_3$ , is not necessarily the same as resolving it in a different order.

With this setup, the statement “consensus is not achieved” is equivalent to the existence of a non-trivial 1-cocycle  $g \in Z^1(X, A)$ , meaning there is a persistent, path-dependent ambiguity in the final value that cannot be resolved by assigning a single value to each process.

**Paxos Mechanics as a Coboundary Operator** The genius of the Paxos algorithm lies in how its two phases actively construct a 0-cochain that trivializes any potential 1-cocycle of disagreement. A 1-cocycle  $g$  is a coboundary if there exists a 0-cochain  $h : \text{Vertices}(X) \rightarrow A$  such that for any edge  $(u, v)$ ,  $g_{uv} = h_u^{-1} \cdot h_v$ . The function  $h$  represents a global, consistent assignment of “state,” and its existence proves that any local disagreement  $g_{uv}$  is merely an artifact of observing this global state from different local perspectives,  $u$  and  $v$ . Paxos constructs this  $h$ .

Let’s interpret the phases of Paxos cohomologically.

### Phase 1: Prepare/Promise - Probing the Topology

A Proposer initiates this phase by selecting a proposal number  $n$  (guaranteed to be unique and higher than any it has used before) and sending a **prepare**( $n$ ) message to a set of Acceptors. \* **Cohomological Action:** The Proposer is attempting to define a consistent “gauge” across a region of the protocol complex. The proposal number  $n$  acts as a coordinate. \* An Acceptor  $p$  receives **prepare**( $n$ ). If  $n$  is greater than any **max\_prepare** it has seen, it updates its state (**max\_prepare** =  $n$ ) and responds with **promise**( $n$ , **accepted\_n**, **accepted\_v**). \* **Cohomological Interpretation:** The **promise** is a crucial piece of local data. The Acceptor is reporting its state relative to a previous potential choice (**accepted\_n**, **accepted\_v**). By promising not to accept any proposal numbered less than  $n$ , the Acceptor is topologically “pruning” future execution paths. It forbids transitions that would violate the total ordering of proposals, effectively simplifying the local topology of  $X$ .

### Phase 2: Propose/Accept - Enforcing the Coboundary Condition

If the Proposer receives **promise** messages from a quorum of Acceptors, it proceeds to Phase 2. The core of Paxos’s safety lies in the rule for choosing the value to propose. \* **The Value Selection Rule:** The Proposer examines all **promise** responses. If any of them contained a previously accepted value, the Proposer *must* choose the value  $v$  associated with the highest proposal number **accepted\_n** among all responses. Otherwise, it is free to choose any value. \* **Cohomological Function:** This rule is the explicit mechanism for constructing the coboundary. Let’s imagine a potential 1-cocycle representing a disagreement. Suppose a Proposer  $P_1$  with proposal  $(n_1, v_1)$  has managed to get it accepted by a quorum  $Q_1$ . Now, a second Proposer  $P_2$  starts a new proposal with number  $n_2 > n_1$ . \* For  $P_2$  to succeed, it must get promises from a quorum  $Q_2$ . Due to the **quorum intersection property**,  $Q_1 \cap Q_2 \neq \emptyset$ . \* At least one Acceptor  $p \in Q_1 \cap Q_2$  that accepted  $(n_1, v_1)$  will respond to  $P_2$ ’s **prepare**( $n_2$ ) request with **promise**( $n_2$ , **n\_1**, **v\_1**). \*  $P_2$  is now algorithmically constrained. It sees the prior choice  $(n_1, v_1)$ . Assuming no other Acceptor reports a higher proposal number,  $P_2$  is *forced* to abandon its own intended value (if different) and instead propose  $(n_2, v_1)$ . \* This act of forcing  $P_2$  to adopt  $v_1$  is what “repairs” the potential disagreement. A potential cycle in the state space where one path leads to  $v_1$  and another to  $v_2$  is collapsed. The algorithm forces the path for  $n_2$  to “learn from” the path for  $n_1$  and conform to its value.

Let  $h$  be the (eventual) globally chosen value, say  $v^*$ . The 0-cochain  $h$  assigns to each process state  $s_p$  the transformation in  $A$  required to get from its local view of the value to  $v^*$ . The Paxos rule ensures that for any two communicating processes with states  $u$  and  $v$ , the local disagreement  $g_{uv}$  between them is exactly the difference reconcilable by mapping both to  $v^*$ , i.e.,  $g_{uv} = h_u^{-1} \cdot h_v$ . Any apparent disagreement is resolved by referencing a higher-order, globally consistent fact: the value associated with the highest-numbered proposal seen so far.



**Quorum Intersection: The Topological Engine of Triviality** From a geometric perspective, the quorum intersection requirement is what prevents the formation of stable “holes” in the protocol complex  $X$  that would support a non-trivial cohomology class. \* A disagreement torsor, represented by a non-trivial element of  $H^1(X, A)$ , can be visualized as a “twist” in the state space. To navigate a loop around a hole in  $X$  and return to your starting point, your view of the global state has been altered by a non-identity element of  $A$ . This holonomy signifies an irresolvable, path-dependent ambiguity. \* Quorums ensure that no two such loops can exist independently. Any two potential “consensus groups” (the quorums) must overlap. This overlap provides a topological “bridge” across which information must flow. \* The information that flows is precisely the constraint from the value selection rule: the history of the highest-numbered accepted proposal. This information propagates across the quorum intersection, stitching the space together and preventing any twisted, inconsistent cycles from stabilizing. It guarantees that any local view of disagreement can be connected to a globally consistent history, thereby collapsing the hole and trivializing the cocycle.

**Conclusion: Paxos and the Trivial Disagreement Torsor** By translating the Paxos protocol into the language of nonabelian cohomology, we arrive at a deeper understanding of its safety property. The traditional proof of Paxos safety is an inductive argument over proposal numbers. The cohomological proof is a statement about the global structure of all possible executions.

The analysis can be summarized as follows: 1. **The Problem:** Consensus is obstructed if the protocol execution space  $X$  and the disagreement group  $A$  admit a non-trivial first cohomology set,  $H^1(X, A) \neq \{[1]\}$ . This signifies the existence of a disagreement torsor. 2. **The Paxos Solution:** Paxos implements a set of rules that act as a dynamic algorithm for proving the triviality of  $H^1(X, A)$  for any given execution. 3. **The Mechanism:** The totally ordered proposal numbers provide a global coordinate system. The quorum intersection property ensures there are no disjoint regions in the state space where independent, contradictory decisions can be made. Finally, the value-selection rule acts as a concrete coboundary operator, using information propagated across quorum intersections to resolve any potential 1-cocycle of disagreement into a 1-coboundary.

Therefore, the safety of Paxos is synonymous with the statement that it guarantees the disagreement torsor is trivial. The protocol does not simply hope for consensus; it actively manipulates the topology and algebra of its own execution space to make disagreement cohomologically impossible. This perspective elevates the analysis of consensus protocols from a procedural checklist to a fundamental inquiry into the geometric nature of distributed computation itself.

## Chapter 5.2: Raft’s Leader Election as Symmetry Breaking in the Protocol Groupoid

### Raft’s Leader Election as Symmetry Breaking in the Protocol Groupoid

The preceding analysis of Paxos revealed a system whose very design ensures the triviality of the disagreement torsor,  $H^1(X, A) = 0$ . Paxos achieves this guarantee through a monolithic and abstract process of quorums and proposals, effectively preventing any non-trivial cocycle from forming. While powerful, this abstraction can be opaque. We now turn to the Raft consensus algorithm, a protocol explicitly designed for understandability. Raft achieves the same goal—consensus—but does so by decoupling the problem into distinct, manageable subproblems: leader election, log replication, and safety.

This chapter posits that Raft’s primary innovation, from a geometric perspective, is its explicit mechanism for **symmetry breaking**. We will model the leader election process as an operator that transforms the underlying *protocol groupoid* from a highly symmetric state, where consensus is obstructed, to a starkly asymmetric one, where consensus becomes trivial. In this view, leader election is not merely a prerequisite for consensus; it is the dynamical process that actively reconfigures the system’s topology and algebra to guarantee that the first cohomology group, which measures disagreement, vanishes.

**The Symmetric Protocol Groupoid: The Pre-Leader State** To understand the role of leader election, we must first characterize the system in its absence. Before a stable leader emerges, the Raft cluster exists in a state of high symmetry. This state can be described by a **protocol groupoid**, which we denote  $\Pi_{\text{sym}}(X)$ , where  $X$  is the protocol complex of possible executions.

- **Objects:** The objects of  $\Pi_{\text{sym}}(X)$  are the global states of the system,  $= (s_1, s_2, \dots, s_N)$ , where each process  $p_i$  is in a state  $s_i$  belonging to  $\{\text{Follower}, \text{Candidate}\}$ . Associated with each process is its current **term** number and its log.
- **Morphisms:** The morphisms are the sequences of allowed computational events: message sends/receives, and critically, internal timeouts that trigger a process to transition from Follower to Candidate.

The defining characteristic of  $\Pi_{\text{sym}}(X)$  is its symmetry. In this state, any non-faulty node is, *a priori*, equivalent to any other. Each node possesses the potential to initiate an election and become the leader. This democratic potential can be formalized by considering the action of the permutation group  $S_N$  on the set of nodes. While the nodes are not identical (they have different IDs), the *rules of the protocol* treat them symmetrically with respect to their right to bid for leadership.

This symmetry has profound algebraic consequences. Let  $A$  be the group of transformations on the system’s replicated state machine (i.e., log entries). If any node can potentially propose an update, the effective group of operations

on the global state is complex. It is not merely  $A$ , but a structure that must also encode *who* is performing the action. This creates a fertile ground for non-commuting operations. For instance, if Candidate  $i$  in term  $T$  and Candidate  $j$  in term  $T'$  (where  $T'$  is initiated after  $T$  but before  $T$  is globally known) both attempt to establish leadership, their actions conflict. Different interleavings of their **RequestVote** messages can lead to divergent global states—specifically, states of “split vote” where no leader is chosen.

These split-vote scenarios are the concrete manifestation of cohomological obstruction in the symmetric groupoid  $\Pi_{\text{sym}}(X)$ . A cycle of failed elections, where the system transitions through various candidate states but never converges on a leader, represents a loop in  $\Pi_{\text{sym}}(X)$  with non-trivial holonomy. The system state fails to advance in a consistent manner because the symmetry has not been broken; there is no globally agreed-upon “preferred” direction. The set of possible future views of the system (e.g., “ $i$  is leader,” “ $j$  is leader,” “no leader”) forms a non-trivial fiber over the current state, indicating the potential for a disagreement torsor. Consensus on the next state is impossible because the rules allow for multiple, conflicting, yet equally valid, future paths.

**The Election Mechanism: A Symmetry-Breaking Operator** Raft’s leader election algorithm is precisely the mechanism designed to destroy the symmetry of  $\Pi_{\text{sym}}(X)$ . It is a dynamic process that prunes the groupoid of possible executions, collapsing the set of potential leaders to a singleton and thereby rendering the system asymmetric. This process can be viewed as an operator  $E$  that maps the system from the symmetric groupoid to an asymmetric one,  $E: \Pi_{\text{sym}}(X) \rightarrow \Pi_{\text{asym}}(X)$ .

The key instruments of this symmetry breaking are:

1. **Terms (term):** The monotonically increasing **term** number acts as a logical clock, foliating the execution space. When a node initiates an election for a new term  $T+1$ , it immediately invalidates any lingering information or leadership claims from term  $T$ . This is a powerful tool for cutting off cycles and preventing outdated state from creating paradoxes. Any message bearing an older term is rejected, effectively severing huge subgraphs of outdated execution paths from the protocol groupoid.
2. **Unilateral Voting:** When a Follower receives a **RequestVote** RPC from a valid Candidate, it makes a local, definitive choice. By casting its vote for Candidate  $c$ , it relinquishes its own potential to become a leader (or to vote for another candidate) within that same term. This is a fundamental act of local symmetry breaking. The Follower projects the  $N$ -dimensional space of potential leaders onto a single choice.
3. **Majority Quorum:** The global symmetry is definitively broken when a Candidate receives votes from a majority of nodes ( $\text{floor}(N/2) + 1$ ). At this moment, the set of potential leaders for the current term collapses from many to one. The system transitions into a new configuration where

one node, the Leader, is uniquely privileged. This event constitutes an irreversible transition in the protocol groupoid, moving the system into the  $\Pi_{\text{asym}}(X)$  configuration.

The election process is thus a procedure for selecting a preferred basis vector in the space of possible system evolutions. Before the election, the system can evolve along paths initiated by any node. After the election, only paths sanctioned by the single, unique leader are valid.

### The Asymmetric Groupoid and Trivial Log-Replication Cohomology

Once a leader  $L$  is elected for term  $T$ , the protocol groupoid  $\Pi_{\text{asym}}(X)$  has a radically different and simpler structure.

- **Asymmetry:** The system is now fundamentally asymmetric. The leader  $L$  is the sole source of log entries. All other nodes become passive Followers whose only role is to replicate the leader’s log and acknowledge receipt. The symmetry group acting on the nodes’ roles collapses from  $S_N$  to the trivial group.
- **Linearization of Operations:** All state machine updates are channeled through the leader. Client requests are redirected to  $L$ , which determines the order in which operations will be appended to the log. This imposes a total ordering on all state transitions. The non-commutativity that plagued the symmetric state—arising from concurrent, conflicting proposals from different nodes—is completely eliminated. The acting group on the replicated log is simply  $A$ , the group of state updates, and its application is controlled by a single agent.

The cohomological consequence of this broken symmetry is profound. We are no longer concerned with consensus on the global state of the *entire* protocol (which includes who is leader), but with the subproblem of log agreement. Let  $X_L$  be the protocol complex of executions under the stable leadership of  $L$ . The task is to determine if disagreement on the log is possible, i.e., to compute  $H^1(X_L, A)$ .

In this asymmetric configuration, any two valid execution paths  $\gamma_1$  and  $\gamma_2$  from one consistent log state to another  $\gamma'$  must ultimately correspond to the same sequence of committed entries, as dictated by  $L$ . The leader acts as a global coordinator, ensuring that the “holonomy” of any loop is trivial. If a follower receives a set of updates via path  $\gamma_1$  and another set via  $\gamma_2$  (perhaps due to message reordering), Raft’s safety mechanisms (e.g., consistency checks in **AppendEntries** RPCs) ensure that these paths are either reconciled to the leader’s single history or one path is rejected. There is no ambiguity. Any 1-cocycle  $g(\gamma)$  measuring the discrepancy along a path  $\gamma$  must be cohomologous to the trivial cocycle.

Therefore,  $H^1(X_L, A) = 0$ . Consensus on the log is guaranteed *precisely because the leader election phase has already reconfigured the system into one where*

*disagreement is topologically impossible.*

### System Dynamics: Flipping Between Symmetry and Asymmetry

The lifecycle of a Raft cluster can be viewed as a dynamic oscillation between the symmetric and asymmetric groupoids.

- **Leader Failure:** When a leader crashes or becomes partitioned from a majority of the cluster, the asymmetric state  $\Pi_{\text{asym}}(\mathbf{X})$  is destroyed. The remaining nodes, upon timeout, will transition to the Candidate state. The system reverts to the symmetric, high-potential-energy state of  $\Pi_{\text{sym}}(\mathbf{X})$ . The symmetry is restored, and with it, the obstruction to consensus.
- **Split Vote:** The canonical failure mode of the election process is the “split vote,” where no candidate secures a majority. This is a state where the system becomes trapped in  $\Pi_{\text{sym}}(\mathbf{X})$ . Multiple nodes remain candidates, each renewing their term and re-soliciting votes. This liveness failure is the price of navigating the symmetric state; it is a direct visualization of the system failing to break the symmetry and collapse the state space. The randomized election timeouts in Raft are a probabilistic mechanism designed to eventually break this deadlock, allowing one node to request votes before others, thus breaking the temporal symmetry and initiating a successful election.

In conclusion, Raft’s architecture provides a beautiful, concrete example of consensus as a two-phase cohomological process. The first phase, leader election, is a frontal assault on the system’s inherent symmetry. It is a messy, probabilistic, and potentially non-terminating process whose sole goal is to transform the protocol groupoid  $\Pi_{\text{sym}}(\mathbf{X})$ , where disagreement ( $H^1 = 0$ ) is possible, into an asymmetric groupoid  $\Pi_{\text{asym}}(\mathbf{X})$ . The second phase, log replication, operates within this simplified, asymmetric world. Here, the centralized control exercised by the leader linearizes all operations, guaranteeing that the disagreement torsor is trivial ( $H^1(\mathbf{X}_L, \mathbf{A}) = 0$ ) and making consensus a deterministic outcome. Raft, by separating these two concerns, externalizes the complex task of symmetry breaking, leaving behind a core replication problem whose solution is, from a cohomological standpoint, trivial.

### Chapter 5.3: Modeling Byzantine Faults: The Non-Abelian Torsor of Malicious Views

#### Modeling Byzantine Faults: The Non-Abelian Torsor of Malicious Views

The preceding analyses of protocols like Paxos and Raft have demonstrated the power of our cohomological framework in systems subject to crash-faults and asynchrony. In those contexts, consensus failure arises from an inability to distinguish a crashed process from a slow one, leading to a global “drift” in perspective that can be modeled as a trivial or carefully managed torsor. We now turn our attention to the most challenging failure model in distributed

computing: Byzantine faults. A Byzantine-faulty process is not merely silent or slow; it is an active adversary. It can deviate from the protocol in arbitrary ways, send conflicting information to different peers, feign failure, and collude with other malicious actors to sabotage consensus.

This active, intentional injection of inconsistency elevates the problem from managing passive uncertainty to combating structured deception. Simple models of state disagreement, sufficient for crash-faults, are inadequate here. The very nature of a Byzantine lie—presenting one face to process  $P_i$  and another to  $P_j$ —is inherently relational and non-commutative. The perceived state of the system becomes radically path-dependent, contingent on the order in which a process receives information from honest and malicious sources. This chapter will demonstrate that this complex landscape of deception is perfectly captured by the concept of a non-abelian  $A$ -torsor, where the non-abelian group  $A$  represents the “algebra of lies” and the non-triviality of the first cohomology group  $H^1(X, A)$  serves as a formal measure of the system’s susceptibility to Byzantine sabotage.

**The Group of Lies: A Non-Abelian Structure for Deception** To model Byzantine behavior, we must first algebrize the act of lying. In a correct system, a state transition from a local view  $s_i$  to  $s'_i$  is governed by a deterministic function of the current state and a received message  $m$ :  $s'_i = f(s_i, m)$ . We can abstract this as the action of an element from a group of valid transformations,  $G$ , on the state space  $S$ . An honest process  $P_j$  sending an update to  $P_i$  induces a transformation  $g_{\{ji\}} \in G$  on  $P_i$ ’s state.

A Byzantine process  $P_b$  shatters this well-defined structure. It can send message  $m_i$  to process  $P_i$  and a conflicting message  $m_j$  to  $P_j$ . This act of *equivocation* cannot be modeled by a single transformation  $g \in G$ . Instead,  $P_b$ ’s action must be described by a mapping from its observers to a set of transformations.

We define the **group of possible transformations**,  $A$ , as the algebraic structure encompassing all possible perceived state updates, both honest and malicious. This group  $A$  must contain the group of correct updates  $G$  as a subgroup, but it is substantially larger. For instance, if the state is a vector in  $\mathbb{R}^n$ ,  $G$  might be the group of translations, while  $A$  could be the full affine group  $\text{Aff}(n, \mathbb{R})$ , including rotations, scaling, and shears that represent malicious data corruption.

The crucial property of  $A$  in the Byzantine context is its **non-abelian nature**. Consider two Byzantine processes,  $B_1$  and  $B_2$ , and an honest process  $P$ .  $B_1$  sends a message that induces transformation  $a_1 \in A$ , and  $B_2$  sends a message inducing  $a_2 \in A$ . If  $P$  receives from  $B_1$  then  $B_2$ , its state becomes  $a_2 \cdot a_1 \cdot s_0$ . If it receives in the opposite order, its state is  $a_1 \cdot a_2 \cdot s_0$ . Since  $a_1 a_2 \neq a_2 a_1$  in general, the order of receiving lies fundamentally alters the outcome. This path-dependence

is the hallmark of non-abelian systems and the primary tool of the Byzantine adversary. The group  $A$  is thus the “group of lies,” whose non-commutativity algebraically encodes the path-dependent confusion sown by malicious actors.

**The Byzantine Torsor:  $H^1(X, A)$  as the Space of Irreconcilable Worldviews** With the protocol execution modeled as a simplicial complex  $X$  and the algebra of state transformations by a non-abelian group  $A$ , we can now formalize the impact of Byzantine faults. As established previously, the first cohomology set  $H^1(X, \mathcal{A})$  (where  $\mathcal{A}$  is the sheaf of  $A$ -valued functions) classifies  $A$ -torsors over  $X$ . An  $A$ -torsor  $E$  is a space that locally resembles  $A$  but lacks a global identity element, or “origin.” Its existence signifies an inability to establish a global, consistent frame of reference.

In the Byzantine setting, this abstract definition gains a potent, concrete interpretation. A 1-cocycle is a map  $g$  that assigns a transformation  $g_{ij}$  in  $A$  to each directed 1-simplex  $(i, j)$  in the protocol complex  $X$  (representing a communication event from a process configuration  $i$  to  $j$ ), subject to the cocycle condition on 2-simplices  $(i, j, k)$ :  $g_{ik} = g_{ij} \cdot g_{jk}$  (Where  $g_{jk}$  may be acted upon by  $g_{ij}$  depending on the specific formulation, reflecting the non-abelian nature of  $A$ ).

An honest system, even with asynchrony, strives to maintain this condition. The transformations compose consistently. A Byzantine process, however, is a factory for violating it. Consider processes  $P_i, P_j, P_k$  and a Byzantine node  $P_b$ . 1.  $P_b$  tells  $P_i$ : “The value is  $v_1$ .” 2.  $P_b$  tells  $P_j$ : “The value is  $v_2$ .” 3.  $P_i$  forwards its view (based on  $v_1$ ) to  $P_k$ . This corresponds to a cocycle element  $g_{ik}$ . 4.  $P_j$  forwards its view (based on  $v_2$ ) to  $P_k$ . This corresponds to  $g_{jk}$ .

Now,  $P_k$  has received conflicting information about the system’s history. The “lie” from  $P_b$  has created a discrepancy. If we trace a loop in the state space of  $P_k$ ’s knowledge—for instance,  $P_k$ ’s state before hearing from anyone, its state after hearing from  $P_i$ , and its state after also hearing from  $P_j$ —the composition of transformations does not return to the identity. This non-trivial holonomy is the signature of a 1-cocycle that is not a coboundary.

The **Byzantine torsor** is the  $A$ -torsor  $E$  classified by this non-trivial cocycle  $[g]$  in  $H^1(X, A)$ . It represents the space of all possible, locally self-consistent, but mutually irreconcilable “worldviews” that can exist within the system. Each fiber of the torsor  $E_\sigma$  over a simplex  $\sigma$  is the set of these coherent-but-wrong perspectives available to the processes in  $\sigma$ . Consensus requires finding a global section of this torsor—a single, globally agreed-upon worldview. The non-triviality of the Byzantine torsor is precisely the obstruction to finding such a section. **Consensus is impossible if and only if  $H^1(X, A)$  is non-trivial.**

**Classifying Byzantine Strategies via Cohomology** This framework allows us to move beyond simply labeling behavior as “Byzantine” and toward a

formal classification of adversarial strategies based on the algebraic structure of the cocycles they generate.

- **Equivocation (e.g., Double-Spending):** This is the fundamental Byzantine act. A node  $P_b$  sends conflicting messages  $m_1$  and  $m_2$  concerning the same event. This directly engineers a “split” in the cocycle  $g$ . For two honest observers  $P_i$  and  $P_j$ , the transformations  $g_{\{bi\}}$  and  $g_{\{bj\}}$  associated with receiving information from  $P_b$  are different elements of  $A$ . This creates a local “tear” in the fabric of causality, ensuring that any path in the protocol complex that attempts to reconcile the views of  $P_i$  and  $P_j$  will accumulate a non-trivial holonomy.
- **Relaying and Framing Attacks:** A Byzantine node  $P_b$  can intercept a message from an honest node  $P_i$  to  $P_j$  and either relay a modified message or relay it selectively. This attack manipulates the cocycle elements corresponding to paths that pass through  $P_b$ . By corrupting the information flow,  $P_b$  can effectively “frame”  $P_i$  as faulty or logically partition the network, creating disjoint regions with inconsistent views that cannot be reconciled.
- **Coordinated Collusion:** A cabal of  $f$  Byzantine nodes  $\{B_1, \dots, B_f\}$  can orchestrate their lies. Algebraically, this means the cocycle elements they generate,  $g_{\{B_k, P_j\}}$ , are not random. They are carefully chosen to reinforce a single, coherent, but false narrative. Their goal is to construct a malicious cocycle  $g_{\{\text{malicious}\}}$  that is robustly non-trivial. They can ensure that the holonomy around many different loops is consistently non-identity, making the disagreement pervasive and difficult to detect or resolve. This coordinated action aims to make the malicious cocycle  $g$  non-cohomologous to the trivial cocycle.

**Fault Tolerance as a Mechanism for Torsor Trivialization** Byzantine Fault Tolerant (BFT) protocols, such as PBFT, can be understood as powerful geometric and algebraic mechanisms designed to systematically **trivialize the Byzantine torsor**. They impose rigid constraints on the protocol’s execution, effectively forcing any potentially malicious cocycle to become a coboundary, rendering the lies globally irrelevant.

1. **Authenticated Channels (Digital Signatures):** Signatures impose a fundamental constraint on the group of lies  $A$ . A Byzantine node  $P_b$  can no longer generate an arbitrary transformation and claim it originated from an honest node  $P_i$ . It can only forward  $P_i$ ’s correctly signed message or create a new message signed with its own key. This prunes  $A$  significantly, eliminating the possibility of “framing” attacks and limiting the adversary’s ability to arbitrarily manipulate cocycle elements associated with honest communication paths.
2. **Quorum Certificates ( $n > 3f$ ):** The requirement in many BFT protocols that a process must wait to receive  $2f+1$  messages (where  $f$  is the



maximum number of Byzantine nodes) before accepting a state transition is a profound topological constraint. In our model, this is a **quorum condition** on the advancement of the protocol. A process will only traverse an edge in the protocol complex  $X$  if it can construct a “certificate” from a specific, large sub-complex of  $2f+1$  reporting nodes.

This condition acts as a **cohomological filter**. By definition, any set of  $2f+1$  nodes must contain at least  $f+1$  honest nodes. Therefore, any valid certificate is guaranteed to be endorsed by a majority of honest processes. A Byzantine cabal of  $f$  nodes can lie, creating malicious cocycle elements. However, their lies will be embedded in a communication structure that also contains at least  $f+1$  consistent, honest reports.

When an honest node evaluates its state, it is effectively integrating information along various paths in  $X$ . The quorum mechanism forces it to only consider paths that are “wide” enough to contain this honest supermajority. When computing the holonomy around any loop constructed from these valid, quorum-certified steps, the contributions from the  $f$  malicious nodes are effectively “cancelled out” or “outvoted” by the  $f+1$  honest nodes. The lies are still present locally, but the protocol’s structure ensures that their global effect integrates to zero. The malicious cocycle  $g_{\{\text{malicious}\}}$  is shown to be a coboundary:  $g_{\{\text{malicious}\}} = \partial h$  for some 0-cochain  $h$ . The disagreement introduced by the adversary can be resolved by a local change of basis ( $h$ ), meaning a global consensus state exists. The condition  $n > 3f$  is the topological requirement that the connectivity of the honest sub-graph of the protocol complex is sufficiently high to “flatten” any twist or hole that  $f$  Byzantine nodes can introduce.

In conclusion, the cohomological framework provides a new and incisive lens through which to analyze Byzantine fault tolerance. By modeling malicious behavior as the action of a non-abelian group  $A$  and consensus failure as the non-triviality of an  $A$ -torsor over the protocol complex, we transform the problem from one of algorithmic cat-and-mouse to one of algebraic topology. The existence of Byzantine consensus protocols is a statement that  $H^1(X, A)$  can be forced to be trivial. The design of these protocols is an exercise in applied algebraic topology: structuring communication and authentication to guarantee that all malicious cocycles are, in fact, trivial coboundaries, thereby proving that a globally consistent state is not just a target, but a topological necessity.

#### Chapter 5.4: CRDTs as a Strategy for Trivializing Cohomology: Enforcing Commutativity to Guarantee Convergence

CRDTs as a Strategy for Trivializing Cohomology: Enforcing Commutativity to Guarantee Convergence

The preceding chapters have established a formal equivalence between the failure to achieve consensus in a distributed system and the existence of non-trivial non-abelian cohomology. Specifically, we have demonstrated that the set of stable, global states of disagreement corresponds to the first nonabelian cohomology set

$H^1(X, A)$ , where  $X$  is the protocol complex representing the system’s possible executions and  $A$  is the (generally non-abelian) group of state transformations. A non-trivial class in  $H^1(X, A)$  represents a “disagreement torsor”—a configuration of local states where pairwise views are consistent, but no single global state can be defined.

This obstruction,  $[g] \in H^1(X, A)$ , arises from the interplay of two factors: 1. **Topological Complexity:** The asynchrony and partial failure modes of the system imbue the protocol complex  $X$  with a rich topology, particularly non-trivial loops corresponding to different orderings of concurrent operations. The fundamental groupoid  $\text{Cat}_1(X)$  is complex. 2. **Algebraic Non-commutativity:** The state update operations form a non-abelian group  $A$ , meaning the order of application matters. This algebraic structure gives rise to non-trivial holonomy; traversing a loop in  $X$  can induce a net transformation in  $A$ , captured by a cocycle  $g$ .

Traditional consensus algorithms, such as Paxos and Raft, achieve consistency by fundamentally attacking the topological complexity. By electing a leader and enforcing a total order on operations, they effectively collapse the intricate protocol complex  $X$  into a topologically trivial, linear history  $X'$ . In this simplified space, the fundamental group is trivial ( $\pi_1(X') = \{e\}$ ), which forces  $H^1(X', A)$  to be trivial, irrespective of the non-commutativity of  $A$ . These protocols pay a high coordination cost to tame the topology.

This chapter explores a dual strategy, one that leaves the complex topology of asynchronous execution untouched but instead tames the algebra. This is the strategy embodied by Conflict-free Replicated Data Types (CRDTs). We will demonstrate that CRDTs are, from a cohomological perspective, a design pattern for constructing state spaces  $A$  with algebraic properties specifically chosen to render  $H^1(X, A)$  trivial for any execution topology  $X$ . They guarantee convergence not by preventing concurrency, but by making its effects algebraically irrelevant.

**The Algebraic Essence of CRDTs: The Join-Semilattice** In conventional terms, a CRDT is a data structure that can be replicated across multiple computers in a network, where replicas can be updated independently and concurrently without coordination between them, and it is always mathematically possible to resolve any inconsistencies that might result. This property is known as **strong eventual consistency**.

From the perspective of our algebraic framework, the guarantees of a CRDT are encoded in the algebraic structure of its state space and the **merge** function that combines states. The state of any replica can be viewed as an element of a set  $S$ , and the operations are functions that map states to new states. The core insight is that for any two states  $s_i, s_j \in S$ , representing the state at replicas  $i$  and  $j$ , there exists a **merge** or **join** operation, denoted  $s_i \vee s_j$ , which produces a new, reconciled state. This operation is endowed with three crucial properties:

1. **Commutativity:**  $s_i \vee s_j = s_j \vee s_i$ . The order in which two states are merged is irrelevant.
2. **Associativity:**  $(s_i \vee s_j) \vee s_k = s_i \vee (s_j \vee s_k)$ . The grouping of merge operations is irrelevant.
3. **Idempotency:**  $s_i \vee s_i = s_i$ . Merging a state with itself produces no change.

A set equipped with such a binary operation is known as a **(join-)semilattice**. The states of a CRDT form a semilattice, where the partial order is defined by  $s_i \leq s_j$  if and only if  $s_i \vee s_j = s_j$ . Convergence in a CRDT system is the process of replicas exchanging their states and merging them, causing their local states to move “up” the semilattice towards the least upper bound of all states.

Let us now reinterpret this within our cohomological framework. The set of states  $S$  is our space  $A$ . The **merge** operation corresponds to the “group” operation. The commutativity property is the most profound: it dictates that the state algebra  $A$  is abelian.

**Trivializing Cohomology by Enforcing Commutativity** The existence of a non-trivial disagreement torsor is predicated on the failure of the cocycle condition for a coboundary. A 1-cocycle is a map  $g : E(X) \rightarrow A$  assigning a transformation  $g_{ij}$  to each directed edge  $(i, j)$  in the execution complex, representing the transformation from process  $i$ ’s view to process  $j$ ’s. The cocycle condition on a 2-simplex  $(i, j, k)$  is  $g_{ij} \cdot (g_{jk})_{\sigma_i} = g_{ik}$ , where  $(g_{jk})_{\sigma_i}$  denotes the action of the path from the basepoint to  $i$  on the group element  $g_{jk}$ .

When the group  $A$  is non-abelian, this condition is difficult to satisfy globally. The path-dependence introduced by the group action and the non-commutativity of the multiplication means that local patches of consistency ( $g_{ij}$  functions) cannot be stitched together into a global state.

Now, consider the CRDT strategy. By designing the data type such that its state space and merge function form a commutative monoid (a semilattice), we enforce that the group  $A$  is abelian. In an abelian group, two critical simplifications occur:

1. The group action is typically trivial. The effect of an update does not depend on the path taken to the current state, only on the state itself. The cocycle condition simplifies to  $g_{ij} + g_{jk} = g_{ik}$  (using additive notation for the abelian group).
2. The operation is commutative. This directly attacks the root cause of path-dependent disagreement.

Let’s formalize this. In our model, a non-trivial class in  $H^1(X, A)$  corresponds to a holonomy representation  $\rho : \text{Cat}_1(X) \rightarrow A$  that is not a coboundary. The holonomy around a loop  $\gamma$  in the execution space  $X$  is the product of the group elements  $g_{ij}$  along the edges of the loop. If  $A$  is non-abelian, a loop corresponding to receiving operations  $(op_1, op_2)$  and then their inverses in a different order,

e.g.,  $(op_2^{-1}, op_1^{-1})$ , may have a non-trivial holonomy given by the commutator  $[op_1, op_2] = op_1 op_2 op_1^{-1} op_2^{-1}$ . This commutator represents the fundamental “disagreement” generated by the concurrency of  $op_1$  and  $op_2$ .

By enforcing commutativity, CRDTs ensure that for any two operations  $op_1, op_2 \in A$ , their commutator is the identity element:  $op_1 op_2 op_1^{-1} op_2^{-1} = e$ . This implies that the holonomy of any loop in the execution space  $X$  arising from reordered operations is trivial. The representation  $\rho : \text{Cat}_1(X) \rightarrow A$  maps all such commutators to the identity.

This effectively trivializes all path-dependent obstructions. The state resulting from a sequence of operations depends only on the *set* of operations applied, not their permutation. A cocycle  $g$  is now a coboundary, meaning it can be expressed as  $g_{ij} = s_j - s_i$  for some global state assignment  $\{s_k\}_{k \in V(X)}$ . The semilattice structure of the CRDT provides this global assignment explicitly. The “global state” can be conceived as the join (least upper bound) of all operations ever created in the system,  $S_{\text{global}} = \bigvee_{op \in \text{AllOps}} op$ . The local state at any replica  $i$ ,  $S_i$ , is simply the join of the subset of operations it has observed. The existence of a least upper bound for any two states  $S_i$  and  $S_j$ —namely  $S_i \vee S_j$ —means there is a canonical, obstruction-free path to reconciliation.

The  $A$ -torsor over  $X$  is trivialized because the semilattice provides a canonical section. For any open cover  $\{U_i\}$  of  $X$  (representing the local views of processes), and a collection of local states  $\{s_i \in A(U_i)\}$ , the section is given by mapping this collection to their join,  $s = \bigvee s_i$ . There is no “twist” that prevents this gluing. The disagreement torsor collapses into a trivial structure, guaranteeing that all replicas will eventually converge to a consistent state.

**A Tale of Two Strategies: Taming Topology vs. Taming Algebra** The cohomological framework provides a beautifully clear distinction between the strategy of classical consensus algorithms and that of CRDTs.

- **Paxos/Raft: Taming the Topology**
  - **Given:** A complex execution space  $X$  with non-trivial  $\text{Cat}_1(X)$  and a non-abelian state algebra  $A$ .
  - **Strategy:** Employ a coordination-heavy protocol (e.g., leader election, two-phase commit) to constrain executions to a single, totally-ordered log. This prunes the complex  $X$  down to a linear graph  $X'$ .
  - **Result:** The fundamental groupoid of  $X'$  is trivial. Therefore,  $H^1(X', A)$  is trivial for *any* group  $A$ . Consensus is achieved by eliminating the topological possibility of disagreement.
  - **Analogy:** Forcing all cars to take one specific, pre-defined road to a destination. There can be no confusion about the route.
- **CRDTs: Taming the Algebra**
  - **Given:** A complex execution space  $X$  with non-trivial  $\text{Cat}_1(X)$  and a state algebra  $A$ .
  - **Strategy:** Engineer the data type itself such that its state algebra

$A$  is a commutative join-semilattice.

- **Result:** The holonomy representation  $\rho : \text{Cat}_1(X) \rightarrow A$  becomes trivial for any loop generated by concurrency. Therefore,  $H^1(X, A)$  is trivial for *any* asynchronous execution space  $X$ . Convergence is guaranteed by making the algebra of states immune to topological complexity.
- **Analogy:** Inventing teleportation. The path taken to the destination is irrelevant; the end result is always the same.

CRDTs trade expressive power in the state algebra for extreme resilience to topological complexity (asynchrony, network partitions, out-of-order delivery). They represent a fundamental shift from procedural solutions for consensus to structural, algebraic ones.

### Beyond State Convergence: Lingering Higher-Order Obstructions

Does the CRDT strategy of trivializing  $H^1(X, A)$  eliminate all cohomological obstructions? Not necessarily. While state convergence is guaranteed, more subtle inconsistencies, corresponding to higher cohomology groups like  $H^2(X, A)$ , may persist.

Recall that  $H^2(X, A)$  classifies gerbes and represents obstructions to the consistency of protocol composition. A non-trivial 2-cocycle represents a situation where local state patchings are consistent on pairs of overlaps (e.g., between views of processes  $i$  and  $j$ , and  $j$  and  $k$ ), but fail to cohere on triple overlaps (among  $i$ ,  $j$ , and  $k$ ).

In a system built with CRDTs, this would not manifest as a failure for a single CRDT's state to converge. Instead, it could emerge as a violation of a global invariant that spans multiple CRDTs or a CRDT and an external system.

**Example:** Consider a distributed banking application. A **GCounter** CRDT tracks an account balance, and a **GSet** CRDT tracks a list of authorized transaction approvers. \* The balance CRDT guarantees all nodes will converge on the correct balance.  $H^1(X, A_{\text{balance}})$  is trivial. \* The approver set CRDT guarantees all nodes will converge on the same set of approvers.  $H^1(X, A_{\text{approvers}})$  is trivial.

However, suppose there is a system-wide invariant: “A withdrawal can only be processed if the number of approvers in the set is greater than two.” This invariant links the state of the two CRDTs. A 2-cocycle could manifest as a scenario where: \* Process  $i$ 's view: {Balance: 100, Approvers: {A, B, C}}. Withdrawal is valid. \* Process  $j$ 's view: {Balance: 200, Approvers: {A, B}}. Withdrawal is invalid. \* Process  $k$ 's view: {Balance: 100, Approvers: {A, B}}. Withdrawal is invalid.

The pairwise consistency between  $(i, j)$  and  $(j, k)$  might seem resolvable, but the inconsistency across the triad  $(i, j, k)$  regarding the *validity of an action* persists. This is a gerbe-like obstruction. The local data (the CRDT states) are

individually consistent and will converge, but their relationship, which defines the correctness of the composed protocol, is flawed. The “state” is not just the CRDT value but the validity of the global application logic.

Therefore, while CRDTs provide a powerful, surgically precise method for trivializing the first cohomology group and thus solving the state convergence problem, they do not absolve designers from reasoning about higher-order invariants and the potential for compositional inconsistencies, which our framework identifies with the persistence of non-trivial higher cohomology. They solve for consensus, but not necessarily for all forms of distributed correctness.

### Chapter 5.5: Blockchain Forks as Manifestations of Non-Trivial Holonomy in the State Space

#### Blockchain Forks as Manifestations of Non-Trivial Holonomy in the State Space

The preceding chapters have demonstrated how protocols like Paxos and Raft achieve consensus by actively constraining the topology of their execution space or by enforcing a symmetry-breaking mechanism that selects a canonical path. Paxos, through its two-phase commit structure, ensures that any potential disagreement torsor remains trivial. Raft, via its leader election, collapses the protocol groupoid to a single authoritative history. Blockchains, particularly those employing Nakamoto-style consensus (e.g., Proof-of-Work), present a starkly different paradigm. Instead of preventing disagreement, they permit it to arise naturally and provide a probabilistic mechanism for its eventual resolution. This chapter argues that a blockchain *fork* is the precise, observable manifestation of non-trivial holonomy within the protocol’s state space, an embodiment of a non-trivial class in the first non-abelian cohomology group  $H^1(X, A)$ .

#### The Blockchain State Space and the Non-Abelian Nature of Updates

To formalize this analysis, we must first define the components of our cohomological model in the context of a blockchain.

- **The State Space ( $\mathcal{S}$ ):** The global state of a blockchain is the ledger itself—a unique, ordered sequence of blocks. However, in a distributed system, no single participant has access to this platonic ideal. Instead, each node  $p_i$  possesses a *local state*  $s_i \in \mathcal{S}$ . This local state consists of the sequence of blocks that  $p_i$  has received, validated, and accepted as its current canonical chain, along with any known but unconfirmed (orphaned) blocks. The total state space  $\mathcal{S}$  is the set of all possible valid chains.
- **The Update Algebra ( $A$ ):** The state of a node evolves through the application of updates. The fundamental update operation is the reception and validation of a new block,  $B$ . Let us denote the operator that appends a valid block  $B$  to a chain  $C$  as  $T_B$ . Thus, if a node is in state  $s_C$  (representing chain  $C$ ), receiving block  $B$  transitions it to state  $s_{C.B} = T_B(s_C)$ .

The set of all such valid transformations can be conceptualized as our coefficient algebra,  $A$ .

The critical property of  $A$  is that it is fundamentally **non-abelian**. This non-commutativity is more subtle than simple matrix multiplication. It arises from the rules of the protocol itself. Consider a node whose canonical chain ends in block  $B_n$ . Suppose two distinct, valid successor blocks,  $B_{n+1,\alpha}$  and  $B_{n+1,\beta}$ , are mined and broadcast nearly simultaneously. A node receiving these blocks must make a choice based on its local policy (e.g., which one it saw first).

The operation  $T_{B_{n+1,\alpha}}$  transforms the state by appending  $B_{n+1,\alpha}$  to the canonical chain. Once this is done, the block  $B_{n+1,\beta}$  can no longer be appended at the same height; it is relegated to the status of an orphan. The application of  $T_{B_{n+1,\alpha}}$  changes the context and invalidates the precondition for applying  $T_{B_{n+1,\beta}}$  to the same parent. Therefore, the composition of updates is path-dependent:  $T_{B_{n+1,\beta}} \circ T_{B_{n+1,\alpha}}(s_{B_n}) \neq T_{B_{n+1,\alpha}} \circ T_{B_{n+1,\beta}}(s_{B_n})$ . The resulting states represent two distinct ledgers, one in which the transactions in  $B_{n+1,\alpha}$  are confirmed and those in  $B_{n+1,\beta}$  are not, and vice versa. This path-dependence, rooted in the non-abelian nature of state updates, is the algebraic seed of consensus failure.

**The Protocol Complex and Execution Cycles** We model the set of all possible asynchronous executions as a topological space, the protocol complex  $X$ . The vertices of  $X$  represent local states  $(p_i, s_i)$ , where node  $p_i$  holds local chain  $s_i$ . A simplex is formed by a set of vertices  $\{(p_1, s_1), \dots, (p_k, s_k)\}$  if those local states can coexist simultaneously within the rules of the protocol.

Asynchrony, in the form of unpredictable network latency, is the primary source of non-trivial topology in  $X$ . It allows for execution paths to diverge and form cycles. Consider the canonical fork scenario: 1. **Initial State:** All nodes  $\{p_i\}$  are in a consistent state, having all agreed on a chain  $C_n$  ending in block  $B_n$ . 2. **Divergence:** Two miners,  $M_\alpha$  and  $M_\beta$ , find valid blocks  $B_{n+1,\alpha}$  and  $B_{n+1,\beta}$  respectively, both building on  $B_n$ . 3. **Asynchronous Propagation:** \*  $M_\alpha$  broadcasts  $B_{n+1,\alpha}$ . Due to network topology, node  $p_1$  receives it first.  $p_1$  updates its state to  $s_\alpha = C_n \cdot B_{n+1,\alpha}$ . \*  $M_\beta$  broadcasts  $B_{n+1,\beta}$ . Node  $p_2$  receives this block first and updates its state to  $s_\beta = C_n \cdot B_{n+1,\beta}$ . 4. **Cross-Communication:** Subsequently,  $p_1$  receives  $B_{n+1,\beta}$  and  $p_2$  receives  $B_{n+1,\alpha}$ . At this point,  $p_1$ 's view is (canonical:  $s_\alpha$ , orphan:  $B_{n+1,\beta}$ ), while  $p_2$ 's view is (canonical:  $s_\beta$ , orphan:  $B_{n+1,\alpha}$ ).

This sequence of events traces a loop in the execution space  $X$ . We can visualize a path in  $X$  corresponding to  $p_1$ 's event history and another for  $p_2$ . These paths start at a common point (agreement on  $C_n$ ) and end at a point where they both know of both new blocks, but they have traversed different intermediate states. This loop structure corresponds to a non-trivial element in the fundamental groupoid (or group,  $\pi_1(X)$ ) of the protocol complex.

**The Fork as a Holonomy Element** As established in previous chapters, a loop  $\gamma$  in the base space  $X$  induces a holonomy transformation  $h_\gamma \in A$  on the fiber (the state space). The holonomy measures the discrepancy in the state after it has been parallel-transported around the loop.

Let's apply this to our fork scenario. Let the initial state of agreement be  $s_0$ , representing the chain  $C_n$ . \* **Path 1** ( $\gamma_1$ ): This path corresponds to the event sequence experienced by node  $p_1$ . The state is transformed by the update operator  $g_1 = T_{B_{n+1,\alpha}}$ . The final state for  $p_1$  is  $s_1 = g_1(s_0) = s_\alpha$ . \* **Path 2** ( $\gamma_2$ ): This path corresponds to the event sequence experienced by node  $p_2$ . The state is transformed by  $g_2 = T_{B_{n+1,\beta}}$ . The final state for  $p_2$  is  $s_2 = g_2(s_0) = s_\beta$ .

The loop in the execution space is  $\gamma = \gamma_1 \circ \gamma_2^{-1}$ . The holonomy of this loop is the transformation required to map the state at the end of path 1 to the state at the end of path 2, relative to their common starting point. This is given by:  $h_\gamma = g_1 g_2^{-1} \in A$

Since  $s_1 \neq s_2$ , the states are physically different, meaning  $g_1 \neq g_2$  and the holonomy element  $h_\gamma$  is not the identity. **This non-trivial holonomy element is the fork.** It is an algebraic object in  $A$  that precisely quantifies the difference between the two divergent views of the ledger. It represents the “reorganization” operation needed to transform one version of the chain into the other.

This entire situation is formally described by a non-trivial 1-cocycle. We can define a function  $f$  on the edges of our execution complex. For the edge representing  $p_1$  adopting  $B_{n+1,\alpha}$ , the value is  $g_1$ . For the edge representing  $p_2$  adopting  $B_{n+1,\beta}$ , the value is  $g_2$ . The transition function  $g_{12}$  that reconciles the view of  $p_1$  with  $p_2$  is precisely this non-identity holonomy element. The collection of these transition functions  $\{g_{ij}\}$  defines a cocycle that is not a coboundary.

The existence of such a cocycle means there is no global section—no single, globally agreed-upon chain. The set of possible states  $\{s_\alpha, s_\beta, \dots\}$  forms an  **$A$ -torsor of disagreement**. It is a space that locally looks like the group  $A$  (any chain can be transformed into another via a re-org operation) but has no globally defined “zero” element (the one true chain).

**Fork Resolution as Cohomology Trivialization** Unlike systems that are designed to be fault-intolerant to such ambiguities, Nakamoto consensus embraces them and provides a probabilistic resolution mechanism: the **longest-chain rule**.

Suppose a new block,  $B_{n+2}$ , is mined, and its parent is  $B_{n+1,\alpha}$ . The chain  $C'_\alpha = C_n \cdot B_{n+1,\alpha} \cdot B_{n+2}$  is now strictly longer than  $C_\beta = C_n \cdot B_{n+1,\beta}$ . When node  $p_2$  (currently on the  $C_\beta$  fork) receives  $B_{n+2}$ , its protocol rules compel it to perform a *reorganization*. It must abandon  $B_{n+1,\beta}$ , rewind its state to  $C_n$ , and then apply the updates for  $B_{n+1,\alpha}$  and  $B_{n+2}$ .

This resolution mechanism can be interpreted as a dynamic process that actively trivializes the non-trivial cocycle. The arrival of  $B_{n+2}$  provides new information



that acts as a coboundary. It allows for the definition of a global state assignment  $\sigma$  (a section) such that the local views can be reconciled. Before  $B_{n+2}$ , the transition  $g_{21}$  needed to map  $p_2$ 's view to  $p_1$ 's view was a non-trivial re-org. After  $p_2$  receives and processes the longer chain information, its state becomes identical to  $p_1$ 's new state. The transition function  $g_{21}$  becomes the identity element in  $A$ . The holonomy vanishes.

This process is probabilistic, not deterministic. A competing block  $B'_{n+2}$  could be found building on  $B_{n+1,\beta}$ , perpetuating the fork and keeping the holonomy non-trivial. The security of a blockchain transaction is therefore a measure of the probability that the cocycle defining the fork on which it resides has been, and will remain, trivialized.

**Conclusion: The Geometry of Probabilistic Consensus** Viewing blockchain forks through the lens of non-abelian cohomology provides a powerful explanatory framework that elevates the concept from a mere race condition to a fundamental geometric property.

- **Forks are not bugs; they are features of the geometry.** The asynchronous, decentralized nature of block propagation inevitably creates a protocol complex  $X$  with non-trivial loops ( $\pi_1(X) \neq \{e\}$ ).
- **The non-abelian algebra of state updates ( $A$ ) provides the “paint” for this topology.** The fact that choosing one block invalidates another at the same height ensures that traversing a loop in  $X$  results in a non-identity transformation in  $A$ .
- **A fork is the physical manifestation of a non-trivial holonomy element.** The algebraic object  $h_\gamma \in A$  is a precise representation of the reorganization required to reconcile the divergent chains. The set of competing chains at any moment constitutes an  $A$ -torsor, a direct witness to a non-trivial class in  $H^1(X, A)$ .
- **Consensus is an act of dynamic trivialization.** The longest-chain rule is not a mechanism for *preventing* disagreement but for *resolving* it. It is an algorithmic process that propagates information (coboundaries) to eventually render the disagreement cocycle trivial, collapsing the torsor to a single point.

This perspective recasts blockchain consensus as a system that does not fear the complex topology of asynchrony but instead leverages it, allowing temporary, localized disagreements (forks/holonomies) to exist while providing a powerful, albeit probabilistic, mechanism for achieving eventual, global consistency. It provides a formal language to distinguish blockchains from protocols like Paxos and Raft, not by their goals, but by their fundamentally different relationship with the algebra of disagreement.